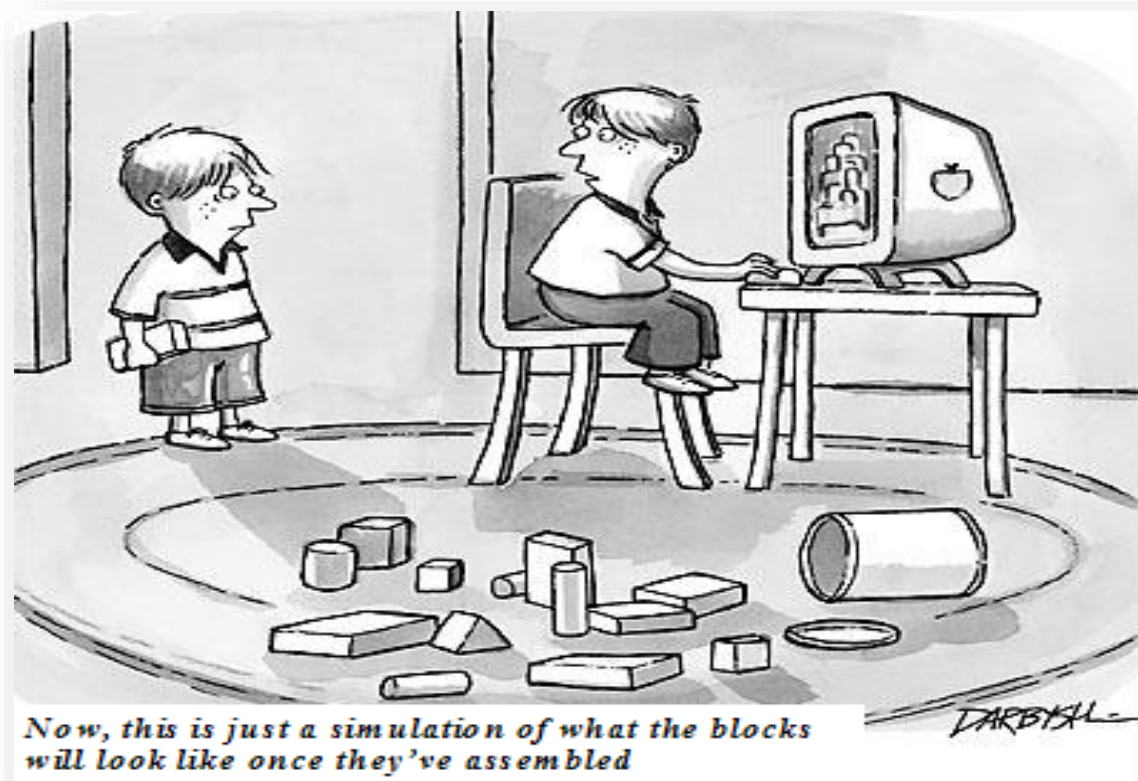


ECE569

Module 28



- Atomic Operations

A Common Arbitration Pattern

- Multiple customers booking air tickets, each
 - Brings up a flight seat map
 - Decides on a seat
 - Update the seat map, mark the seat as taken
- *A bad outcome*
 - Multiple passengers ended up booking the same seat

A Common Collaboration Pattern


- **Multiple bank tellers count the total amount of cash in the safe**
 - Each grab a pile and count
 - Have a central display of the running total
 - Whenever someone finishes counting a pile, add the subtotal of the pile to the running total
- **A bad outcome**
 - Some of the piles were not accounted for.

Histogram Generation

- A method for extracting notable features and patterns from large data sets
 - Feature extraction for object recognition in images
 - Fraud detection in credit card transactions
 - Unusual patterns (just like object recognition): Type and location of purchases
- Basic histograms - for each element in the data set, use the value to identify a “bin” to increment
 - With a cumulative distribution function you can quickly analyze the data

Text Histogram Example

- Define the **bins** as four-letter sections of the alphabet: a-d, e-h, i-l, n-p, ...
- For each character in an input string, **increment** the appropriate bin counter.
- In the phrase “**Programming Massively Parallel Processors**” the output histogram is shown below:



Serial Code

```
sequential_Histogram(char *data, int length, int *histo) {  
  
    for (int i = 0; i < length; i++) {  
  
        int alphabet_position = data[i] - 'a';  
  
        if (alphabet_position >= 0 && alphabet_position < 26) {  
  
            histo[alphabet_position/4]++  
        }  
    }  
}
```

Data Race in Parallel Thread Execution

- Old and New are per-thread register variables.
 - Question 1: If Mem[x] was initially 0, what would the value of Mem[x] be after threads 1 and 2 have completed?
 - Question 2: What does each thread get in their Old variable?
- Unfortunately, the answers may vary according to the relative execution timing between the two threads, which is referred to as a data race.

thread1: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow
New

thread2: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow
New

Data Race in Parallel Thread Execution

Time	Thread 1	Thread 2
1	(0) Old \leftarrow Mem[x]	
2	(1) New \leftarrow Old + 1	
3	(1) Mem[x] \leftarrow New	
4		(1) Old \leftarrow Mem[x]
5		(2) New \leftarrow Old + 1
6		(2) Mem[x] \leftarrow New

Time	Thread 1	Thread 2
1	(0) Old \leftarrow Mem[x]	
2	(1) New \leftarrow Old + 1	
3		(0) Old \leftarrow Mem[x]
4	(1) Mem[x] \leftarrow New	
5		(1) New \leftarrow Old + 1
6		(1) Mem[x] \leftarrow New

Purpose of Atomic Operation – Correct Outcome

thread1: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

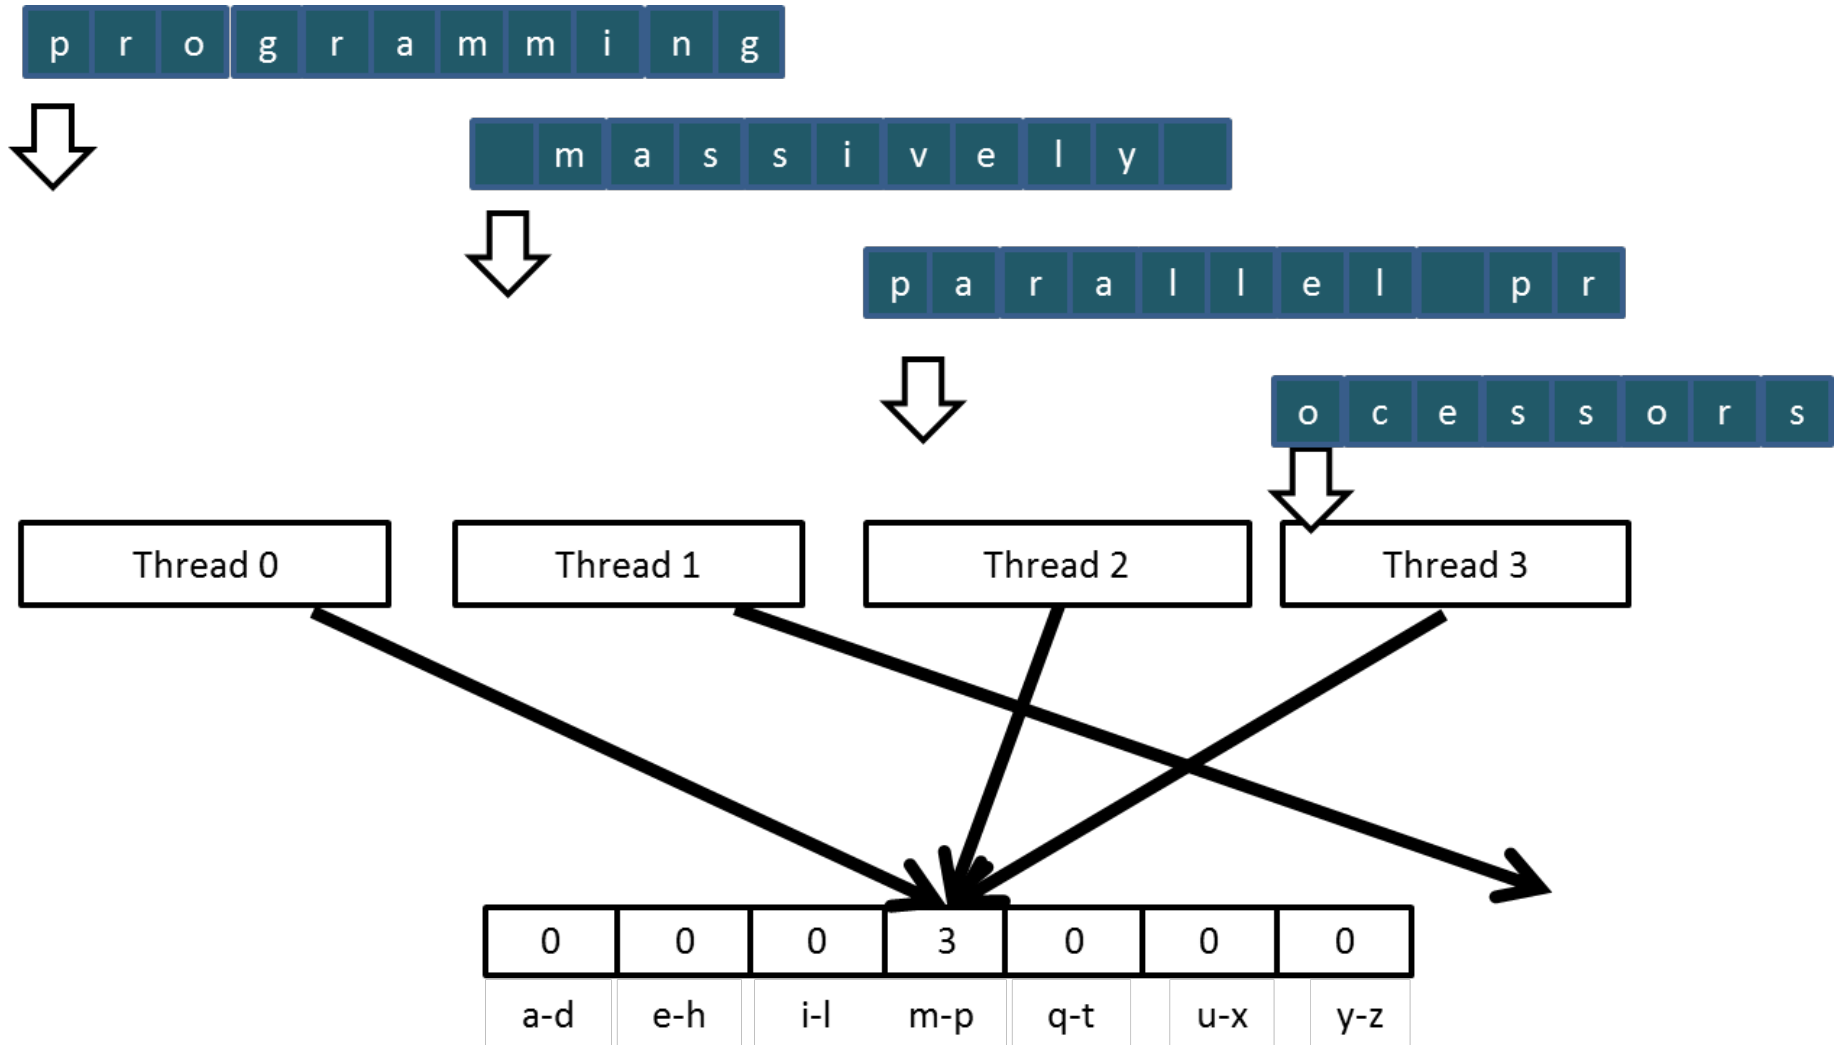
thread2: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

Or

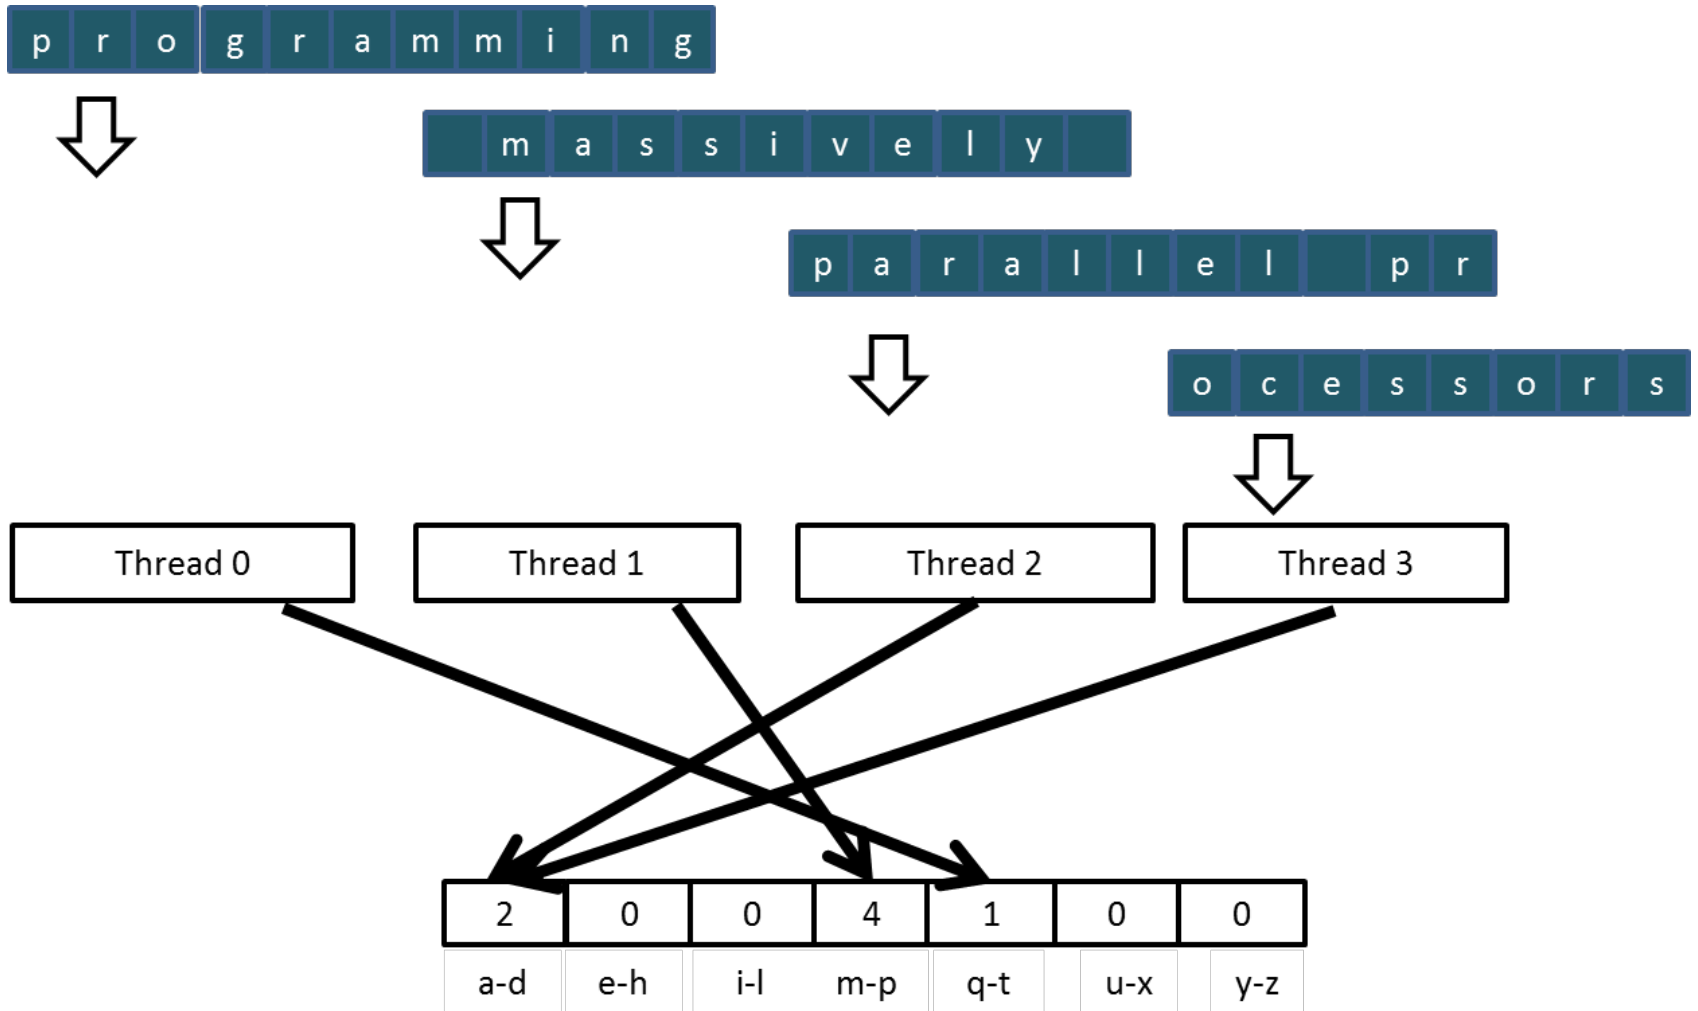
thread1: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

thread2: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

Histogram Generation Example – Iteration 1



Histogram Generation Example – Iteration 2



Atomic Operations

- Atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)
 - Read CUDA C programming Guide for details
- Atomic Add
 - `int atomicAdd(int* address, int val);`
 - reads 32-bit word from the location pointed to by address in global or shared memory,
 - computes (old + val),
 - stores the result back to memory at the same address.
 - returns old value
 - These operations are performed in one atomic transaction.

Atomic Add

- Unsigned 32-bit integer atomic add
 - `unsigned int atomicAdd(unsigned int* address, unsigned int val);`
- Unsigned 64-bit integer atomic add
 - `unsigned long long int atomicAdd(unsigned long long int* address, unsigned long long int val);`
- Single-precision floating-point atomic add (Compute capability 2.x+)
 - `float atomicAdd(float* address, float val);`
- Double-precision floating-point atomic add (Compute capability 6.x+)
 - `double atomicAdd(double* address, double val);`
- 16-bit floating-point atomic add (Compute capability 7.x+)
 - `__half atomicAdd(__half* address, __half val);`

Atomic Add Demo

- **Refer to D2L**
 - Content → Demo → 6.Atomic

```

#include <stdio.h>
#define NUM_THREADS 1000000
#define ARRAY_SIZE 10
#define BLOCK_WIDTH 1000
__global__ void increment_naive(int *g){
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    i = i % ARRAY_SIZE;
    g[i] = g[i] + 1; }

int main(int argc, char **argv){
    // declare and allocate host memory
    int h_array[ARRAY_SIZE];
    const int ARRAY_BYTES = ARRAY_SIZE * sizeof(int);
    // declare, allocate, and zero out GPU memory
    int * d_array;
    cudaMalloc((void **) &d_array, ARRAY_BYTES);
    cudaMemset((void *) d_array, 0, ARRAY_BYTES);

    increment_naive<<<NUM_THREADS/BLOCK_WIDTH, BLOCK_WIDTH>>>(d_array);

    // copy back the array of sums from GPU and print
    cudaMemcpy(h_array, d_array, ARRAY_BYTES, cudaMemcpyDeviceToHost);
    // What will be values in each element in h_array?
    // free GPU memory allocation and exit
    cudaFree(d_array);
    return 0;}

```

```

#include <stdio.h>
#define NUM_THREADS 1000000
#define ARRAY_SIZE 10
#define BLOCK_WIDTH 1000
__global__ void increment_atomic(int *g){
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    i = i % ARRAY_SIZE;
    atomicAdd(& g[i], 1);}

int main(int argc, char **argv){
    // declare and allocate host memory
    int h_array[ARRAY_SIZE];
    const int ARRAY_BYTES = ARRAY_SIZE * sizeof(int);
    // declare, allocate, and zero out GPU memory
    int * d_array;
    cudaMalloc((void **) &d_array, ARRAY_BYTES);
    cudaMemset((void *) d_array, 0, ARRAY_BYTES);

    increment_atomic<<<NUM_THREADS/BLOCK_WIDTH, BLOCK_WIDTH>>>(d_array);

    // copy back the array of sums from GPU and print
    cudaMemcpy(h_array, d_array, ARRAY_BYTES, cudaMemcpyDeviceToHost);
    // What will be values in each element in h_array?
    // free GPU memory allocation and exit
    cudaFree(d_array);
    return 0;}

```


Demo code with timer utility to experiment with

```
#include <stdio.h>
#include "gputimer.h"
#define NUM_THREADS 1000000
#define ARRAY_SIZE 10
#define BLOCK_WIDTH 1000

int main(int argc, char **argv)
{
    int runs;
    GpuTimer timer;
    printf("%d total threads in %d blocks writing into %d array elements\n",
        NUM_THREADS, NUM_THREADS / BLOCK_WIDTH, ARRAY_SIZE);
    // declare and allocate host memory
    int h_array[ARRAY_SIZE];
    const int ARRAY_BYTES = ARRAY_SIZE * sizeof(int);

    // declare, allocate, and zero out GPU memory
    int * d_array;
    cudaMalloc((void **) &d_array, ARRAY_BYTES);
    cudaMemset((void *) d_array, 0, ARRAY_BYTES);
    // launch the kernel - comment out one of these
    timer.Start();
    for (runs=0; runs<100; runs++){
        // increment_naive<<<NUM_THREADS/BLOCK_WIDTH, BLOCK_WIDTH>>>(d_array);
        increment_atomic<<<NUM_THREADS/BLOCK_WIDTH, BLOCK_WIDTH>>>(d_array);
    }
    timer.Stop();
    // copy back the array of sums from GPU and print
    cudaMemcpy(h_array, d_array, ARRAY_BYTES, cudaMemcpyDeviceToHost);
    print_array(h_array, ARRAY_SIZE);
    printf("For %d runs total time elapsed = %g ms with average per run = %g ms\n", runs, timer.Elapsed(), timer.Elapsed()/runs);
    // free GPU memory allocation and exit
    cudaFree(d_array);
    return 0;}
```

Atomic Add Demo

- **Refer to D2L**
 - Content → Demo → 6.Atomic
 - Comment out the “print_array” in the main
 - Run the increment_atomic kernel with the following configurations
 - 1 million threads, incrementing 10 elements
 - 1 million threads, incrementing 100 elements
 - 1 million threads, incrementing 500 elements
 - 1 million threads, incrementing 1000 elements
 - 1 million threads, incrementing 5000 elements
 - Does the execution time trend make sense?

Reading

- CUDA Programming Guide

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

- Section 5.4.2: control flow and predicates
- Section 5.4.3: synchronization
- Appendix B.5: `__threadfence()` and variants
- Appendix B.6: `__syncthreads()` and variants
- Appendix B.14: atomic functions
- Appendix B.18: warp voting

Next

- **Histogram**