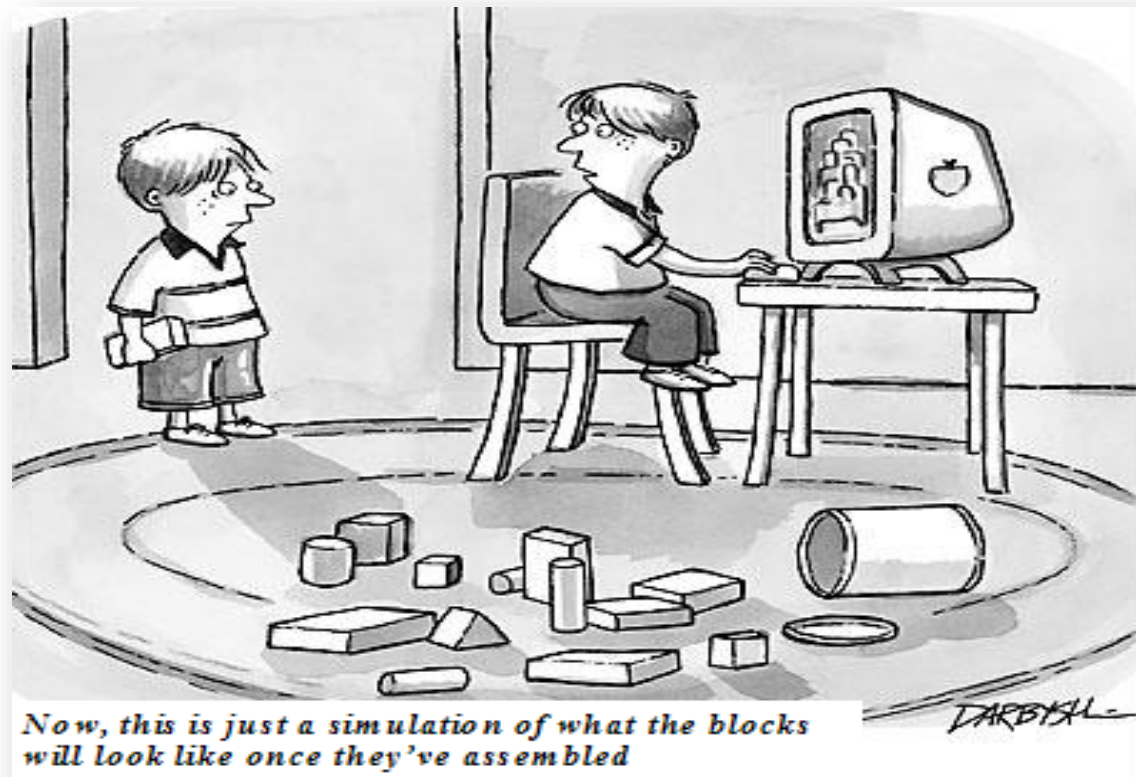


# ECE569

## Module 7

---

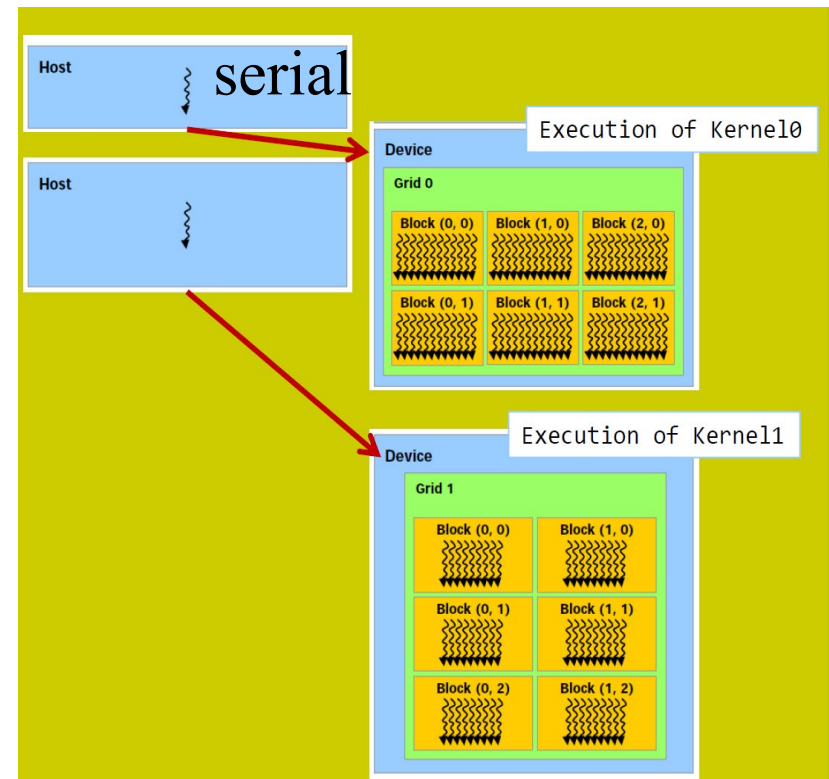


- Thread Organization (Grid, Block)

# CUDA Execution Model is Asynchronous

- **Heterogeneous host (CPU) + device (GPU) application C program**
  - Serial parts in host C code
  - Parallel parts in device SPMD kernel code

By default, execution on host doesn't wait for kernel to finish



# Quick check

---

- **What is GPU good at?**

- ☐ Launching a small number of threads efficiently
- ☐ Launching a large number of threads efficiently
- ☐ Running one thread very quickly
- ☐ Running one thread that does lots of work in parallel
- ☐ Running a large number of threads in parallel

# Vector Addition Host Code: Kernel Launch

---

```
void vecAdd(float *h_A, float *h_B, float *h_C, int n){
    int size = n * sizeof(float);
    float *d_A, *d_B, *d_C;

    cudaMalloc((void **) &d_A, size);
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_B, size);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &d_C, size);

    // Kernel invocation:
    // vecAddKernel <<<grid,block>>>(params)

    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

# Launching the kernel from host

---

- `kernel_routine<<<gridDim, blockDim>>>(args);`
  - `gridDim` is the number of instances of the kernel
    - (the “grid” size = number of blocks)
  - `blockDim` is the number of threads within each instance
    - (the “block” size)
  - `args` is a limited number of arguments, usually mainly pointers to arrays in graphics memory, and some constants which get copied by value

# Thread Blocks: Scalable Cooperation

---

- `kernel_routine<<<10, 256>>>(args);`
  - Number of blocks = 10, Number of threads per block = 256



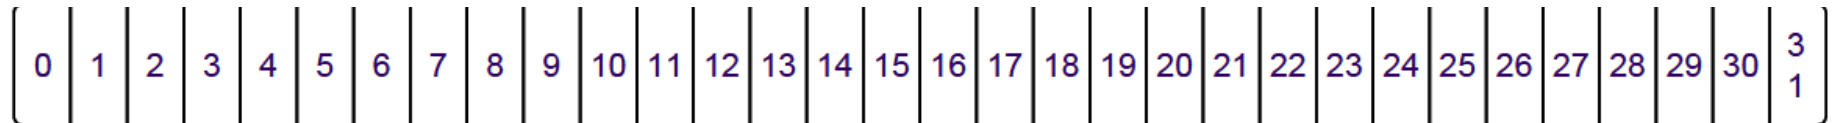
- **Thread position: Which block, which thread?**
  - Need a block identifier: `blockIdx.x`
  - Need a thread identifier: `threadIdx.x`
    - `.x` notation
      - In `x` direction (can be 3D: `x,y,z`)
  - Analogy
    - `blockIdx.x`: area code, `threadIdx.x`: local phone number
  - Mapping function
    - $Id = blockIdx.x * blockDim.x + threadIdx.x$

# Example Array Indexing

---

- Consider indexing into an array A, where one thread is accessing one element. Assume that you launch **8** threads per block and the array is 32 entries long. Which element of the A array will be processed by the thread **ld** of index 5 in block of index 2?

- ☐ 13
- ☐ 21
- ☐ 29
- ☐ 42



# Grid Size and Block Size

---

- **1D grid with 4 blocks, each with 64 threads:**
  - gridDim = 4
  - blockDim = 64
  - blockIdx ranges from 0 to 3
  - threadIdx ranges from 0 to 63



# Next

---

- **Writing kernel code**
  - Vector addition