ECE569
Module 11

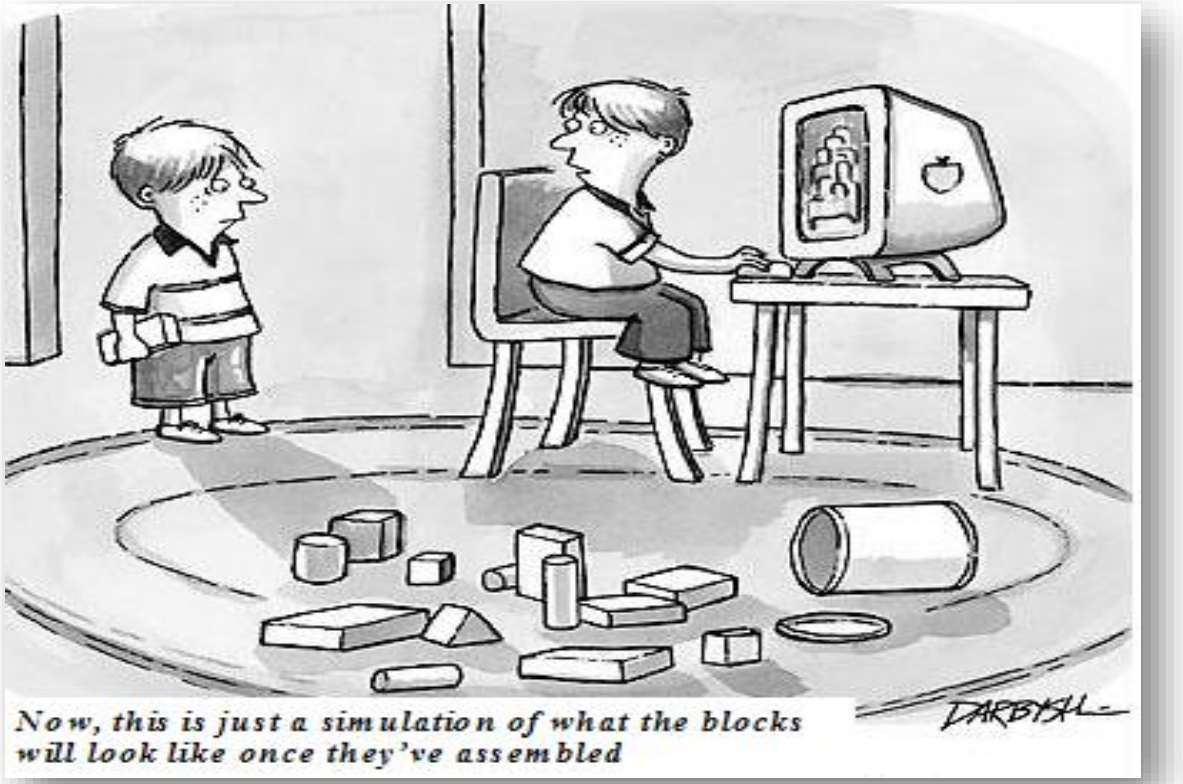

- Profiling

Developer Tools - Profilers

NSIGHT

NVVP

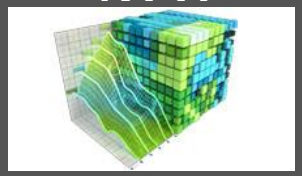

NVPROF

```
==20561== Profiling result:
Time(%)      Time     Calls       Avg       Min       Max  Name
 49.88%  866.69ms    504758  1.7170us  1.5840us  2.0160us  void th
int, thrust::detail::device_generate_functor<thrust::detail::fill_
 25.33%  440.05ms    252662  1.7410us  1.5360us  2.3680us  void th
t, thrust::detail::device_generate_functor<thrust::detail::fill_fu
 17.07%  296.60ms       200  1.4830ms  1.2040ms  1.7253ms  kerComp
  2.98%  51.819ms       200  259.09us  246.97us  264.83us  kerMake
  1.16%  20.173ms       501  40.265us     928ns  17.677ms  [CUDA m
  0.93%  16.190ms       200  80.991us  71.840us  90.751us  kerColl
  0.73%  12.636ms       400  31.589us  14.720us  50.432us  [CUDA m
  0.69%  12.075ms       200  60.376us  59.680us  62.304us  kerRema
  0.63%  10.993ms       200  54.963us  52.608us  58.208us  kerMake
  0.32%  5.5524ms       200  27.761us  22.559us  33.152us  [CUDA m
  0.12%  2.1342ms         1  2.1342ms  2.1342ms  2.1342ms  void th
```

NVIDIA Provided

TAU

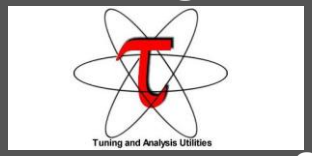

VampirTrace

3rd Party

https://developer.nvidia.com/performance-analysis-tools

NVPROF EXERCISE

- Command Line Profiler
  - Compute time in each kernel
  - Compute memory transfer time
  - Collect metrics and events
  - Support complex process hierarchy's
  - Collect profiles for NVIDIA Visual Profiler
  - No need to recompile
- **$ nvprof ./myadd**
  - View available metrics

    $ **nvprof --query-metrics**
  - View global load/store efficiency

    $ nvprof --metrics gld_efficiency,gst_efficiency ./myadd
  - Store a timeline to load in NVVP

    $ **nvprof –o profile.timeline ./myadd**
  - Store analysis metrics to load in NVVP

    $ **nvprof –o profile.metrics --analysis-metrics ./myadd**

NVPROF EXERCISE

- Either in interactive session
- or
- Use job submission script given in hw1 as a template and include nvprof as part of the “/usr/bin/time mpirun” command
  - /usr/bin/time mpirun -np 1 nvprof --log-file profile.txt ./myadd
  - /usr/bin/time mpirun -np 1 nvprof --metrics all -o nvvp_profile.nvvp ./myadd
  - /usr/bin/time mpirun -n 1 nvprof -o profile.timeline ./myadd
  - /usr/bin/time mpirun -n 1 nvprof –o profile.metrics --analysis-metrics ./myadd

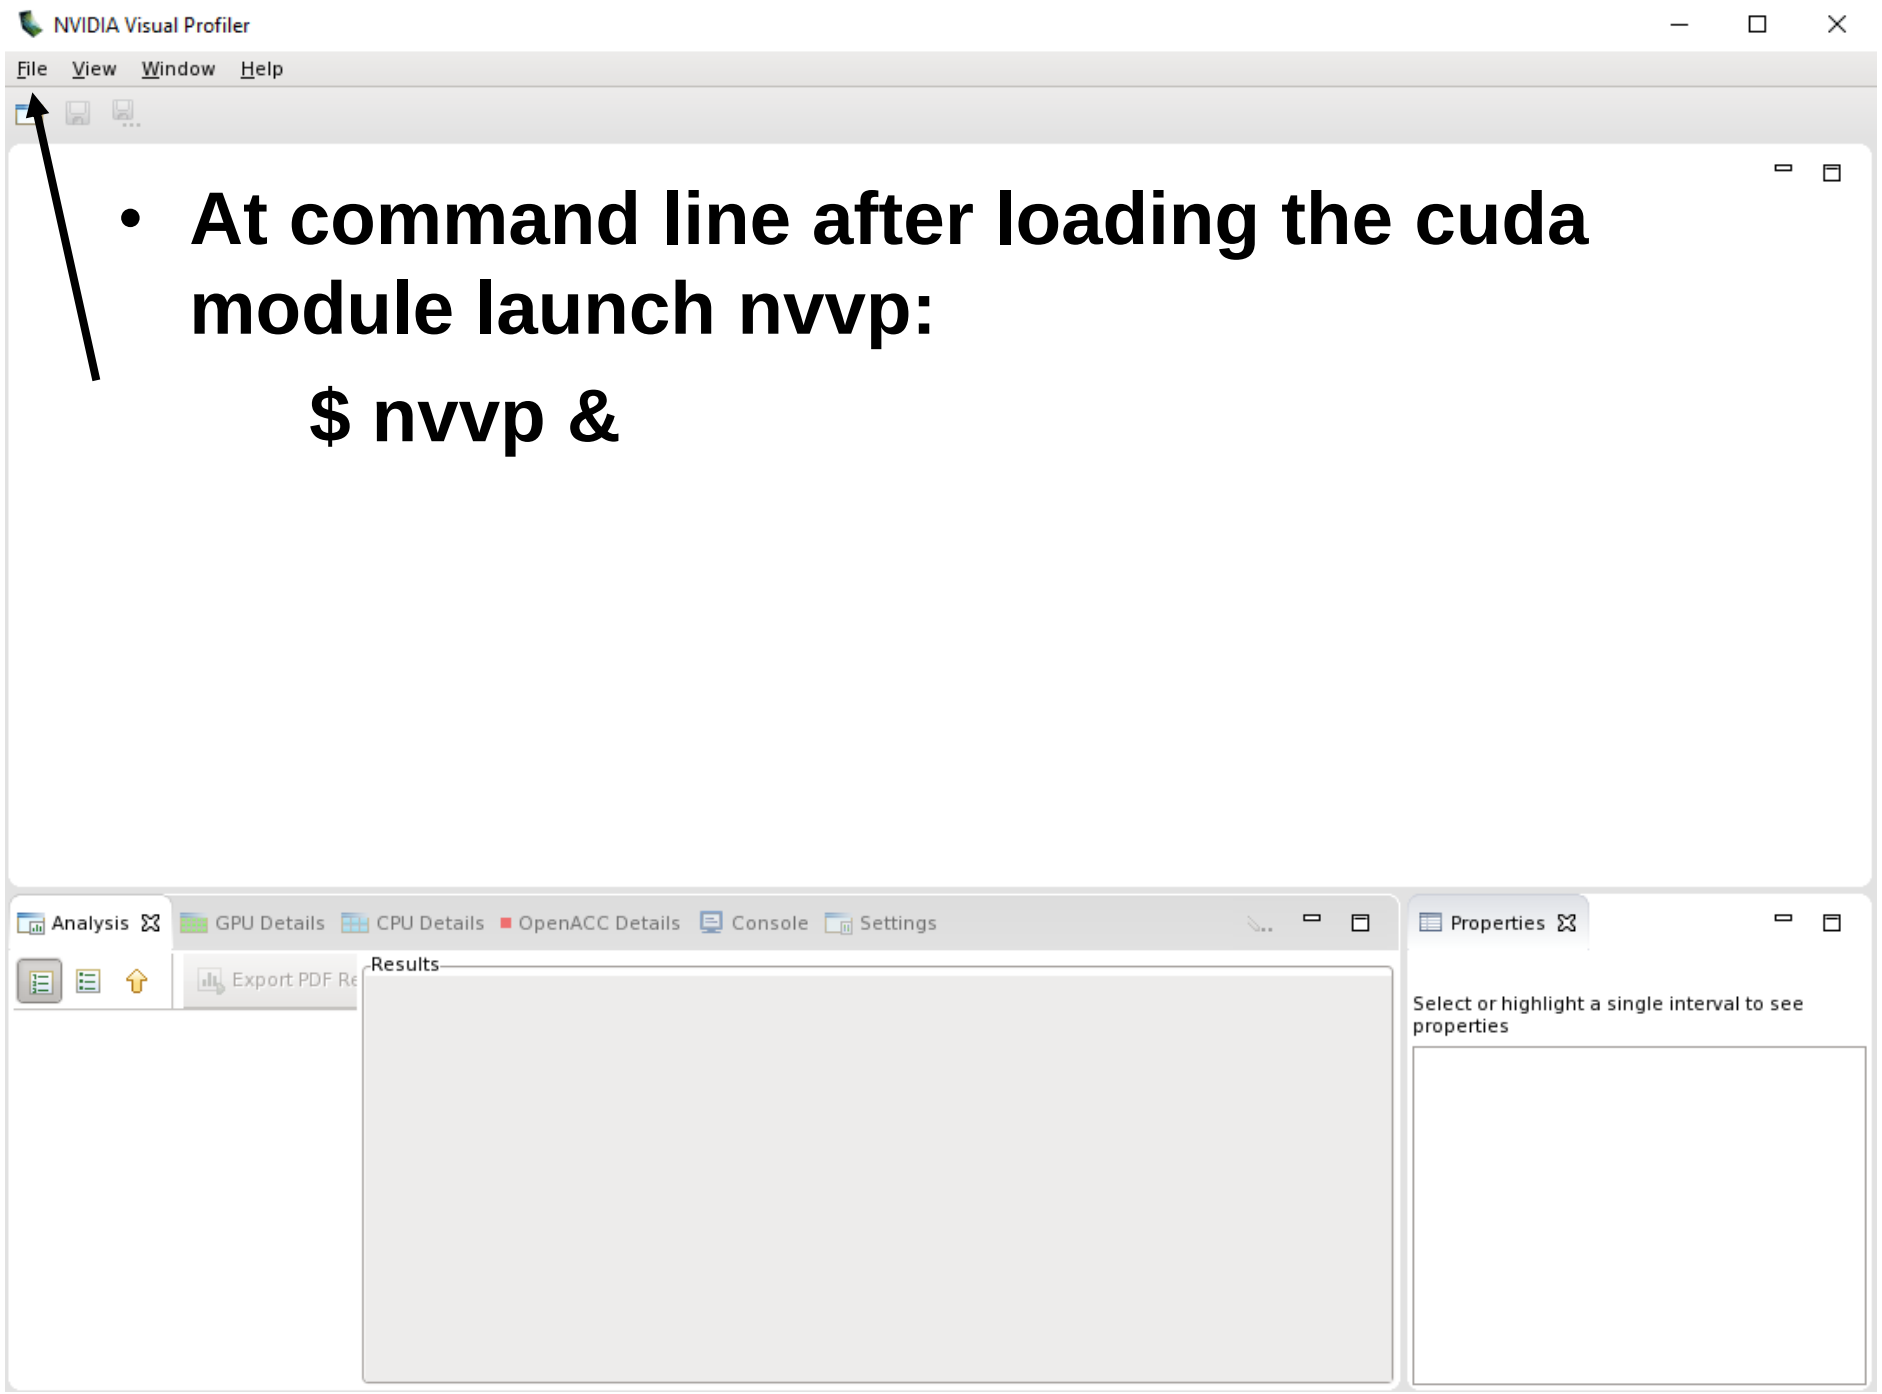

- **At command line after loading the cuda module launch nvvp:**

  $ nvvp &

Import nvprof into NVVP

- Launch nvvp
- Click
  - File/ Import/ Nvprof/ **Next**/ Single process/ Next / Browse

  (next to Timeline data file)
  - Select profile.timeline
- Add Metrics to timeline
  - Click on 2nd Browse (for Event/Metric data files)
  - Select profile.metrics

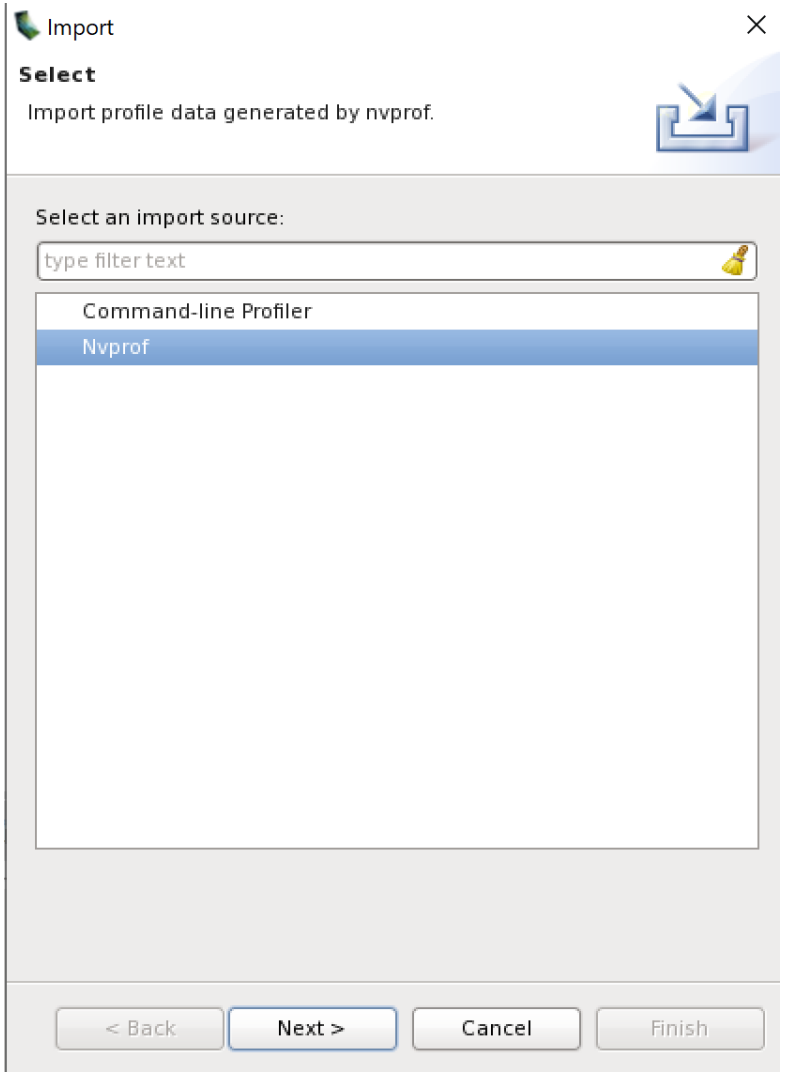

Load your profile files

- Select profile.timeline
- Click on 2nd Browse (for Event/Metric data files)
- Select profile.metrics

- Click Finish
- Explore Timeline
  - Control + mouse drag in timeline to zoom in
  - Control + mouse drag in measure bar (on top) to measure time

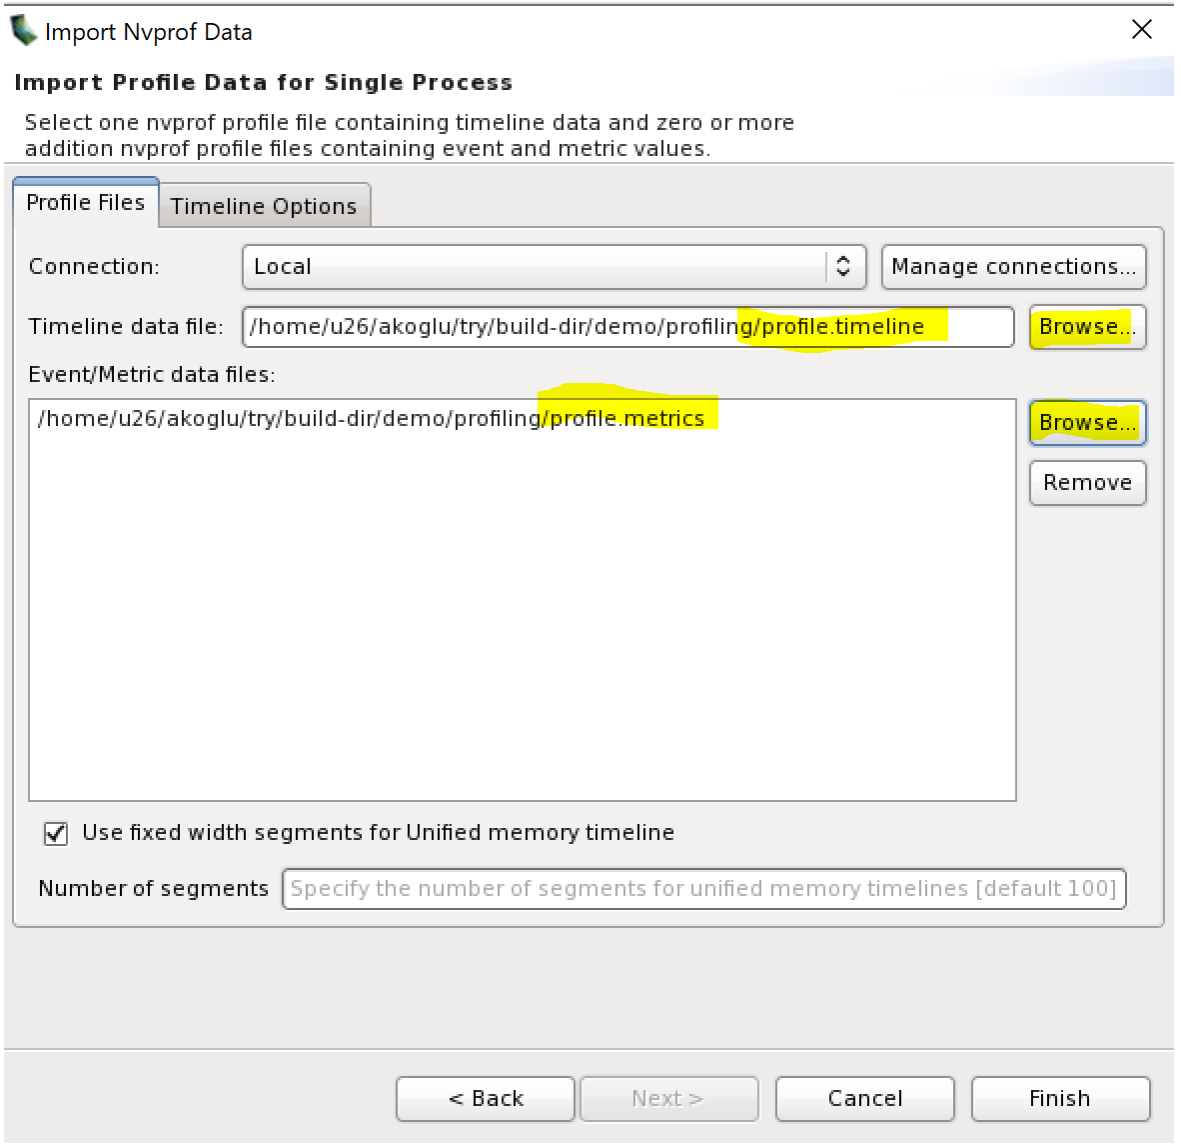

NVVP

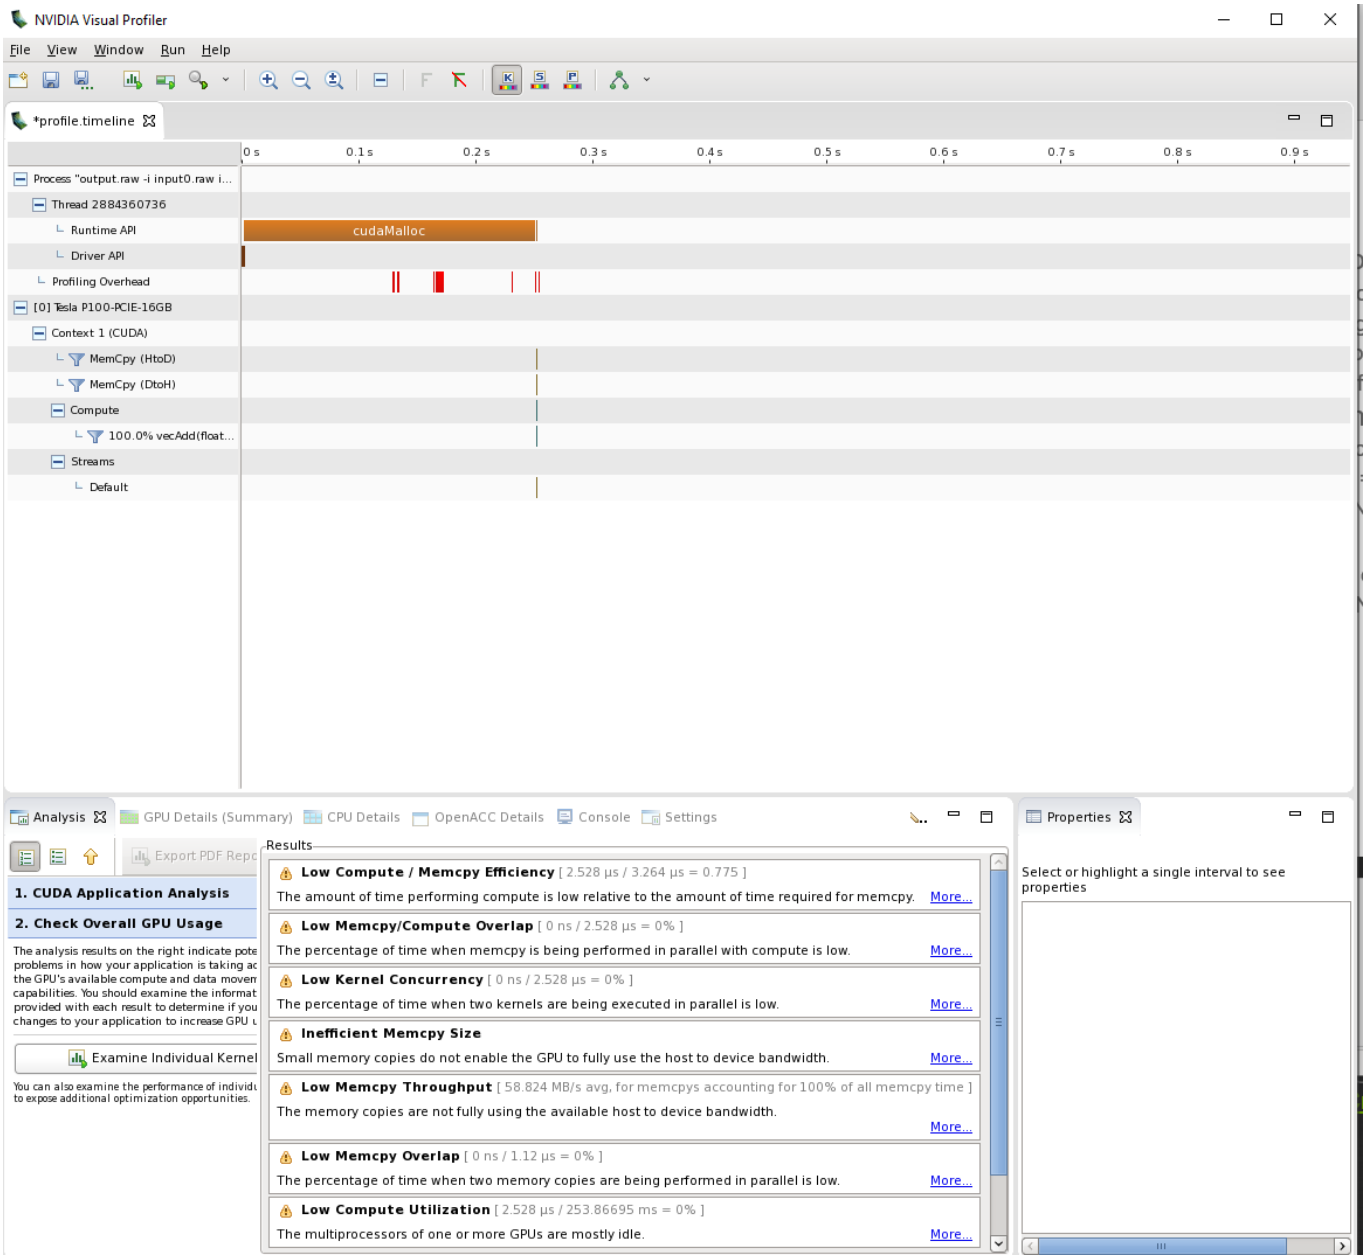
NVIDIA Visual Profiler
File View Window Run Help
*profile.timeline
0 s 0.1 s 0.2 s 0.3 s 0.4 s 0.5 s 0.6 s 0.7 s 0.8 s 0.9 s
Process "output.raw -i input0.raw i...
Thread 2884360736
Runtime API
cudaMalloc
Driver API
Profiling Overhead
[0] Tesla P100-PCIE-16GB
Context 1 (CUDA)
MemCpy (HtoD)
MemCpy (DtoH)
Compute
100.0% vecAdd(float...
Streams
Default
Analysis
GPU Details (Summary)
CPU Details
OpenACC Details
Console
Settings
Properties
Export PDF Repo
1. CUDA Application Analysis
2. Check Overall GPU Usage
The analysis results on the right indicate pote
problems in how your application is taking ac
the GPU's available compute and data mover
capabilities. You should examine the informat
provided with each result to determine if you
changes to your application to increase GPU u
Examine Individual Kernel
You can also examine the performance of individu
to expose additional optimization opportunities.
Results
Low Compute / Memcpy Efficiency [2.528 µs / 3.264 µs = 0.775]
The amount of time performing compute is low relative to the amount of time required for memcpy. More...
Low Memcpy/Compute Overlap [0 ns / 2.528 µs = 0%]
The percentage of time when memcpy is being performed in parallel with compute is low. More...
Low Kernel Concurrency [0 ns / 2.528 µs = 0%]
The percentage of time when two kernels are being executed in parallel is low. More...
Inefficient Memcpy Size
Small memory copies do not enable the GPU to fully use the host to device bandwidth. More...
Low Memcpy Throughput [58.824 MB/s avg, for memcpys accounting for 100% of all memcpy time]
The memory copies are not fully using the available host to device bandwidth.
More...
Low Memcpy Overlap [0 ns / 1.12 µs = 0%]
The percentage of time when two memory copies are being performed in parallel is low. More...
Low Compute Utilization [2.528 µs / 253.86695 ms = 0%]
The multiprocessors of one or more GPUs are mostly idle. More...
Select or highlight a single interval to see properties

Timing Analysis Window

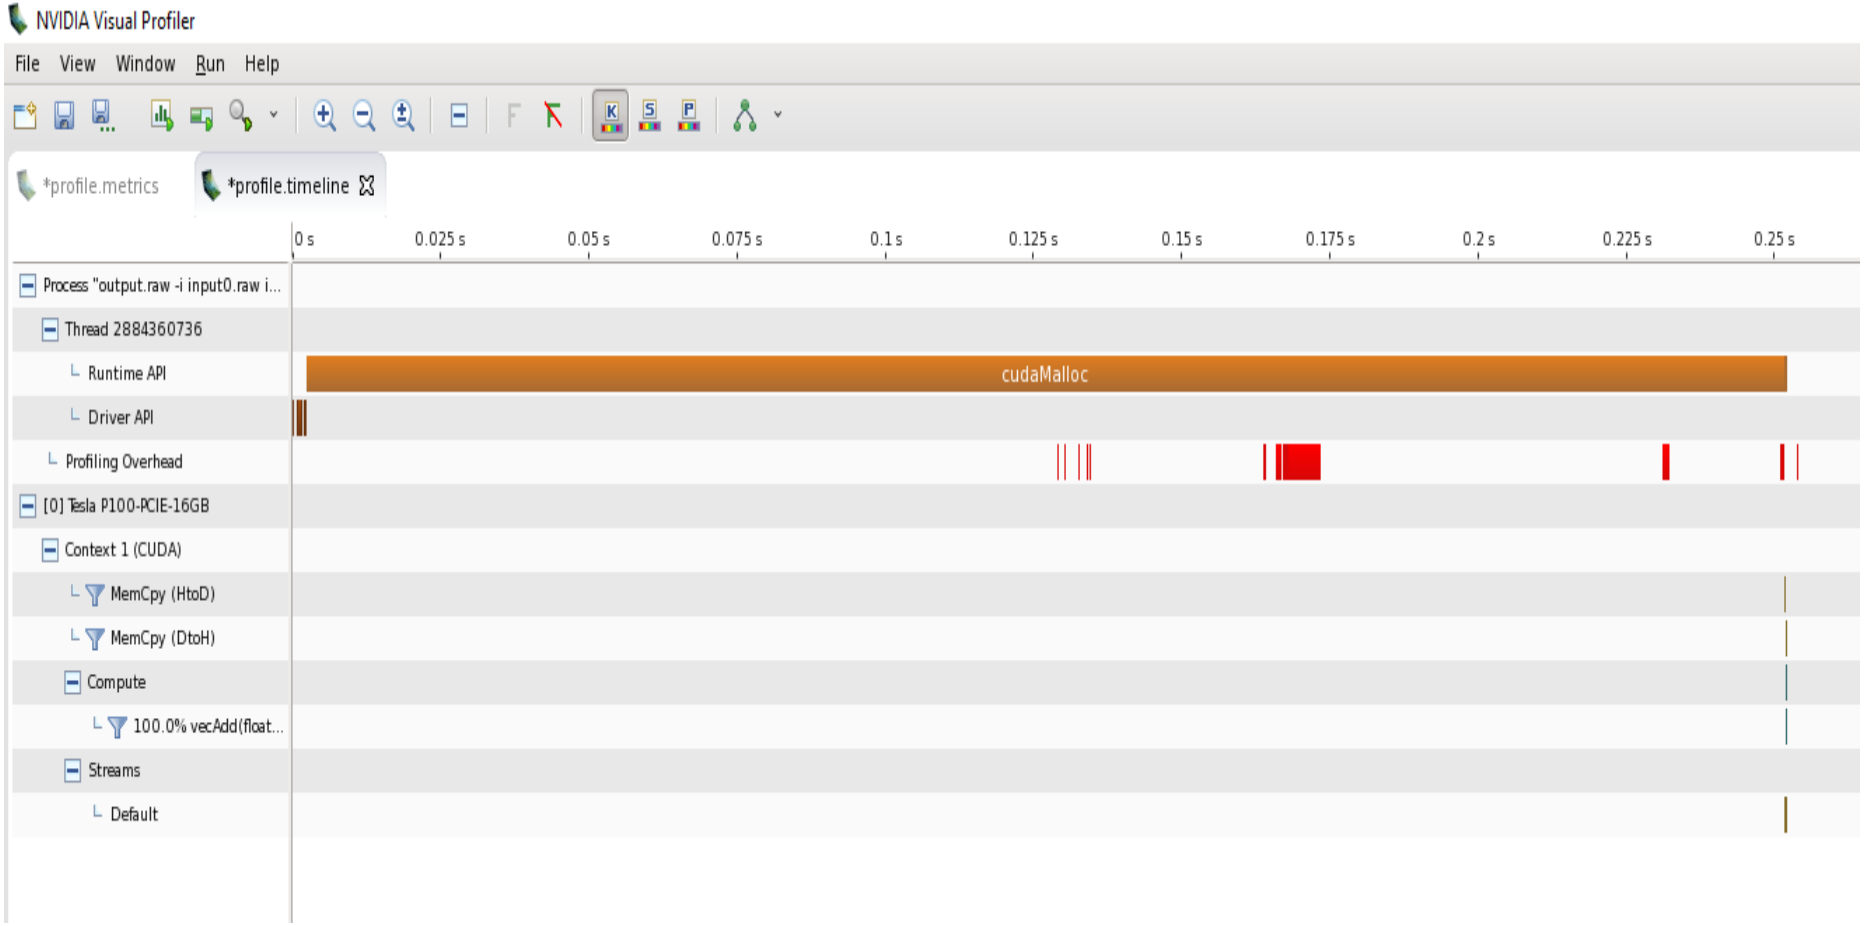

Guided Analysis Window

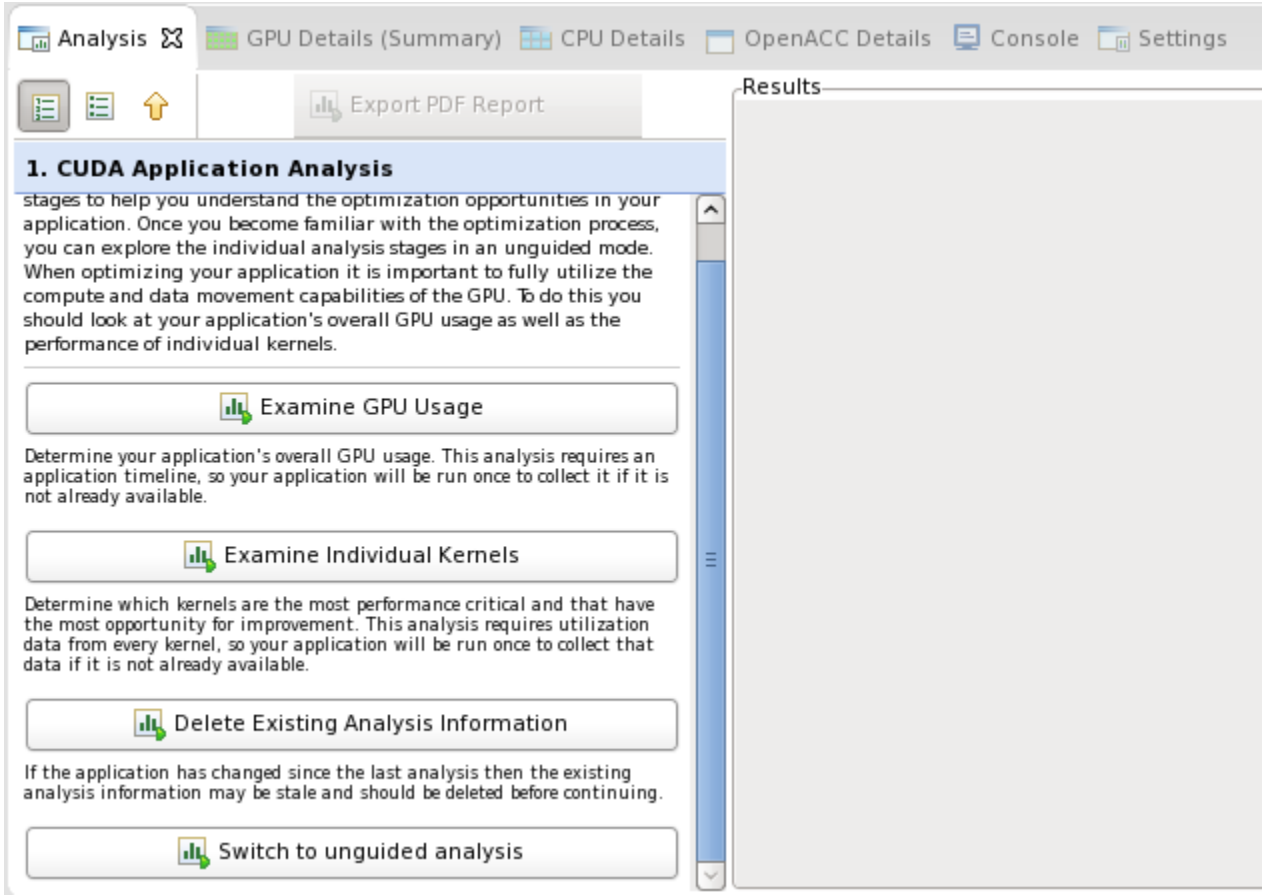

Guided Analysis – GPU Usage

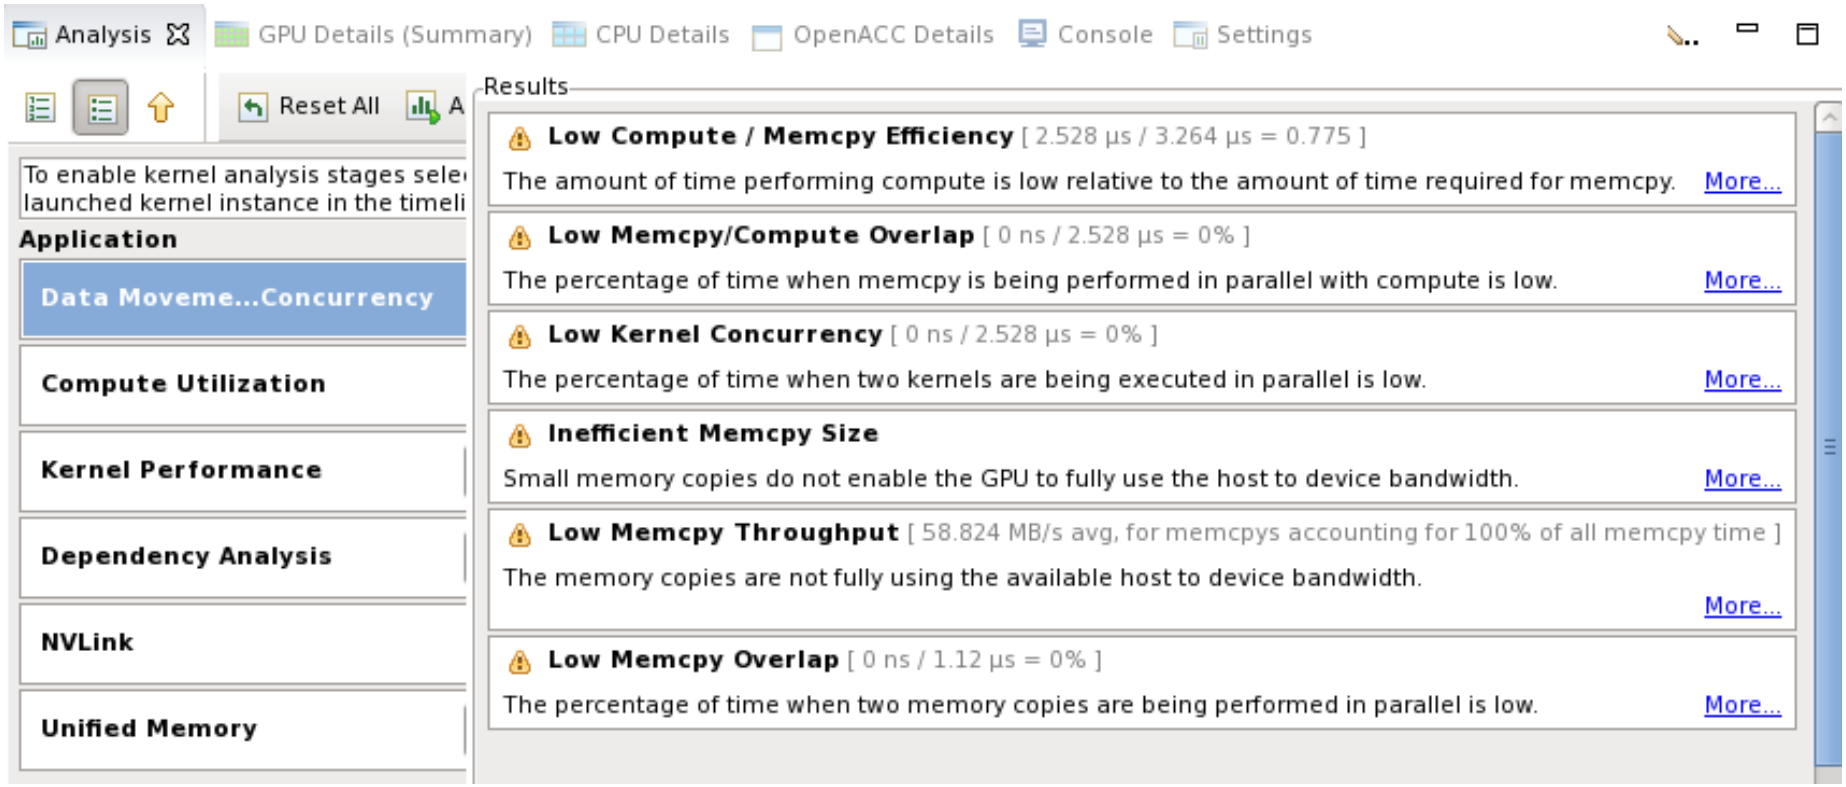

Unguided Analysis

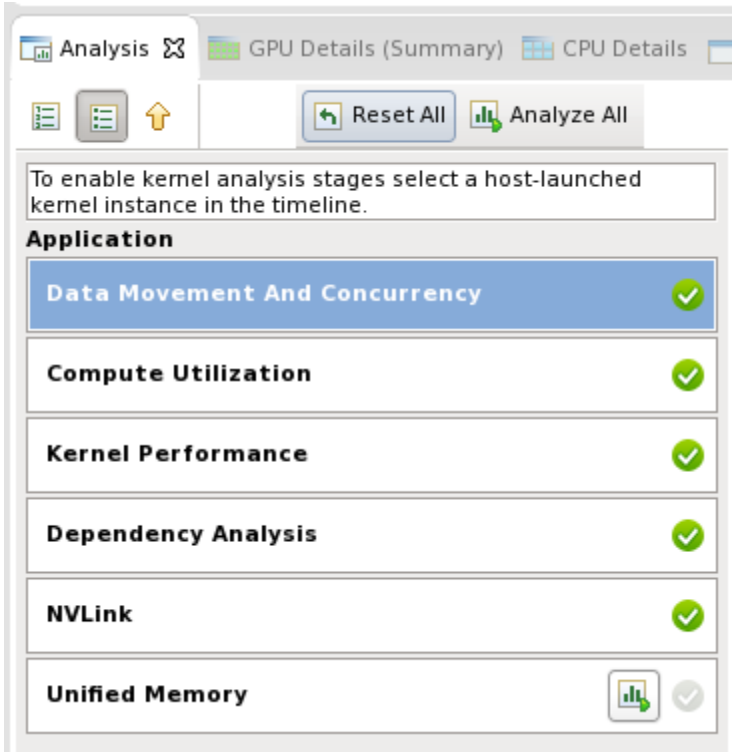

Instructions:

1. Click on a kernel
2. On Analysis tab click on the unguided analysis

2. Click Analyze All

   Explore metrics and properties

Unguided Analysis

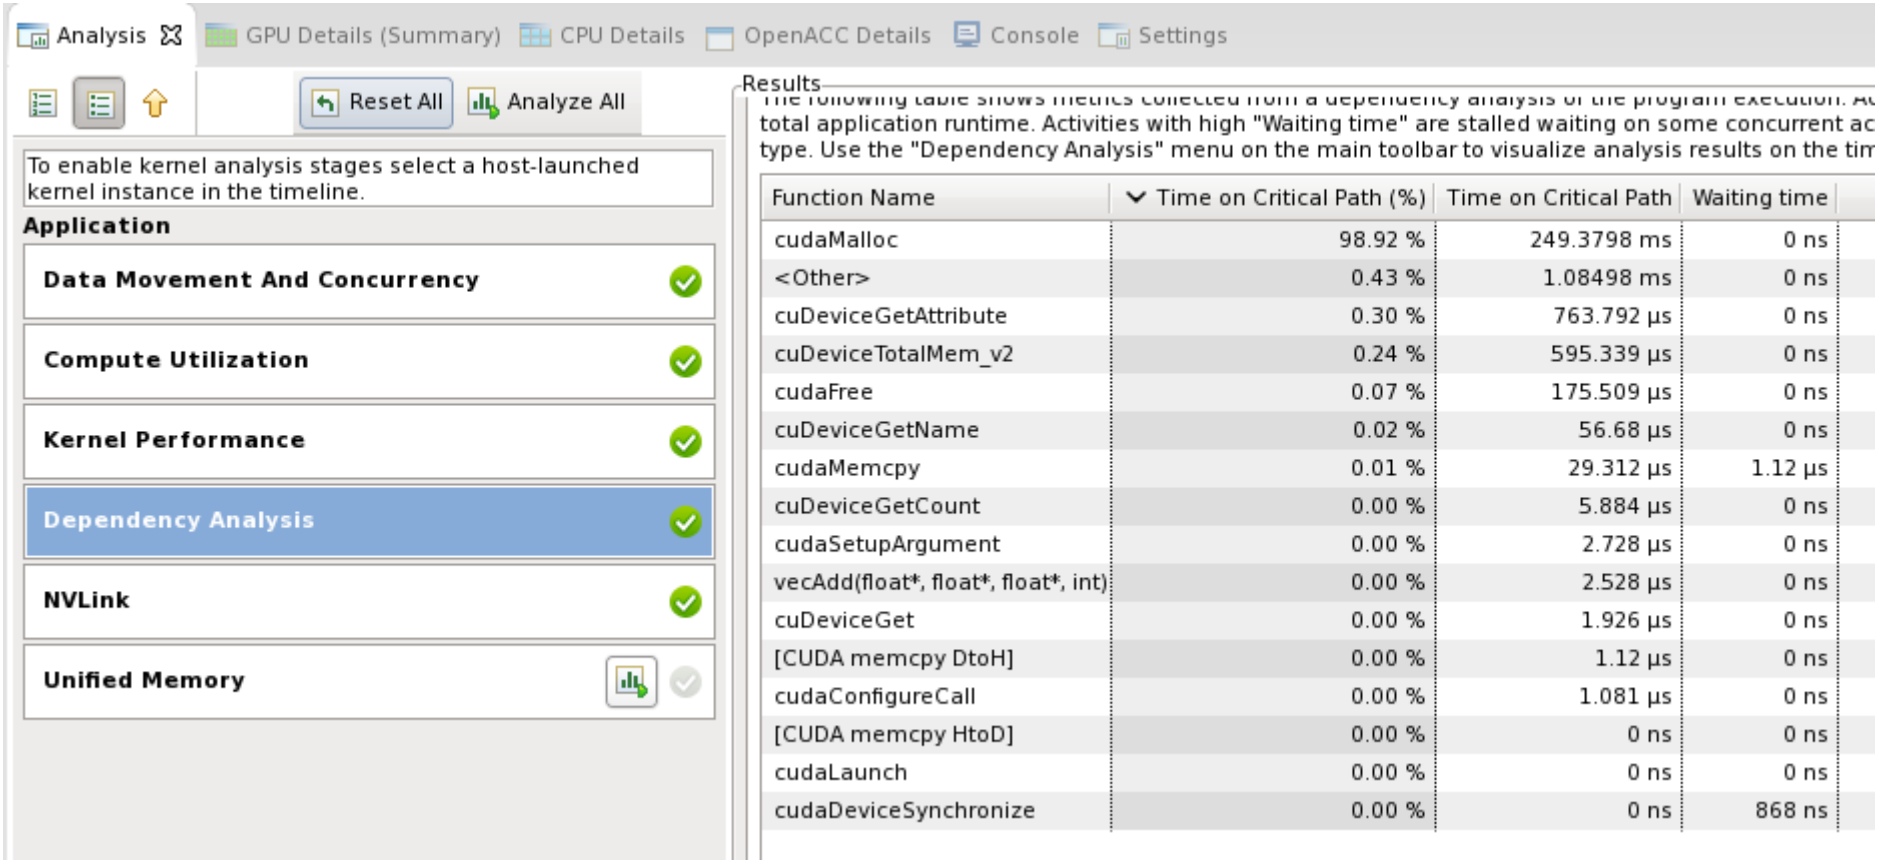

NSIGHT

- **CUDA enabled IDE**
  - Source code editor: syntax highlighting
  - Build Manger
  - Visual Debugger
  - Visual Profiler
- **Linux/Macintosh**
  - Editor = Eclipse
  - Debugger = cuda-gdb with a visual wrapper
  - Profiler = NVVP
- **Windows**
  - Integrates directly into Visual Studio
  - Profiler is NSIGHT VSE

Nsight Eclipse Edition

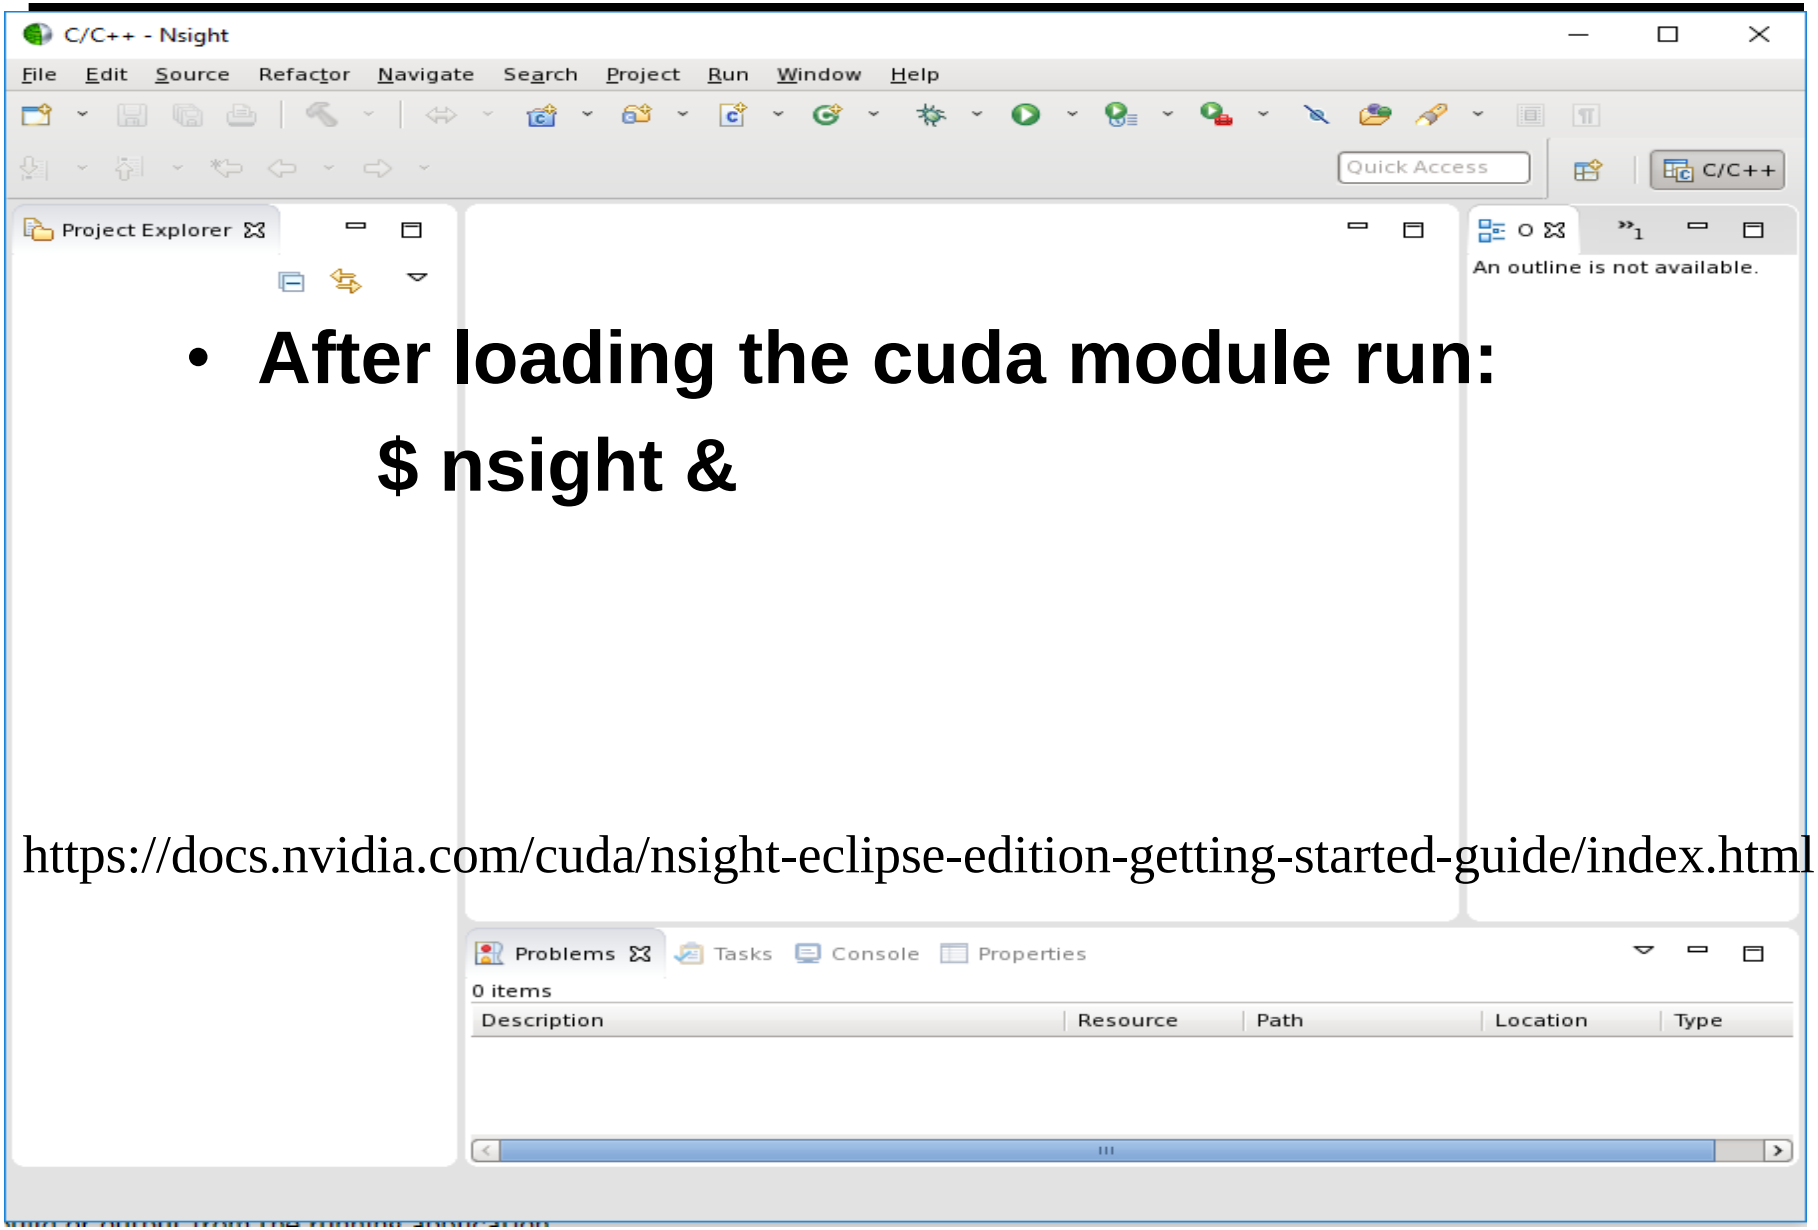

- **After loading the cuda module run:**

  $ nsight &

https://docs.nvidia.com/cuda/nsight-eclipse-edition-getting-started-guide/index.html

NSIGHT: Importing project

Instructions:

1. Run nsight
   Select default workspace
2. Click File / New / Makefile Project With Existing CodeTest
3. Enter Project Name and select the project directory
4. Click Finish
5. Right Click On Project / Properties / Run Settings / New / C++ Application
6. Browse for executable
7. In Project Explorer double click on .cu and explore source
8. Click on the build icon
9. Click on the run icon
10. Click on the profile icon

Before profiling

- Certain kinds of errors cause CUDA programs to complete, but crash under profiling
- Check your program with cuda-memcheck if code behaves incorrectly under profiling

```
cuda-memcheck ./my-program ...
```

Before profiling

- **Compile device code with optimizations**
  - Optimizations dramatically change performance
  - Remove “-G” device-debug flag from nvcc
- **Compile device code with line information**
  - Add “-lineinfo” flag to nvcc

```
nvcc –G main.cu          nvcc –lineinfo main.cu
```

Profiling, Debugging Tools Summary

- Performance programmers rely on debugging and profiling tools
  - You should setup your environment and choose the tools most convenient for you.
  - For assignments you may not need them much but for your project you will need these tools
    - quantify how well you are utilizing compute resources, memory bandwidth

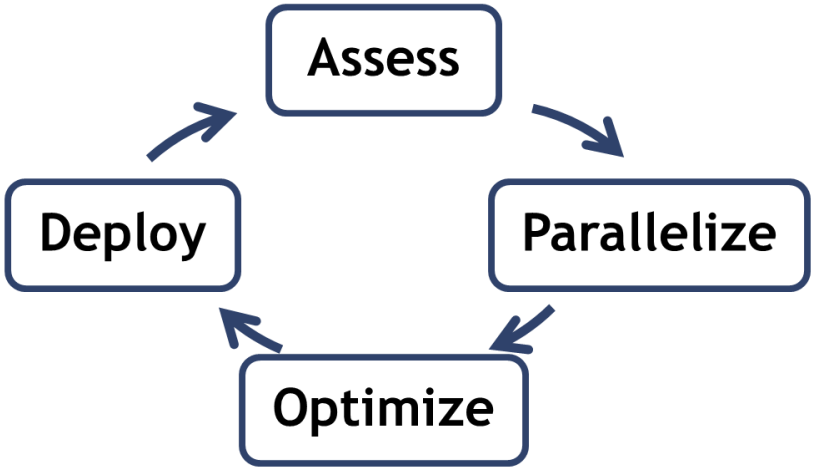

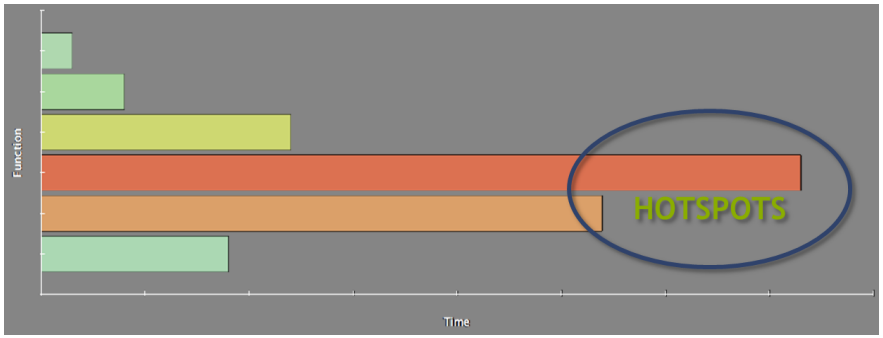

Next

- **Threads, Thread blocks, workload management**