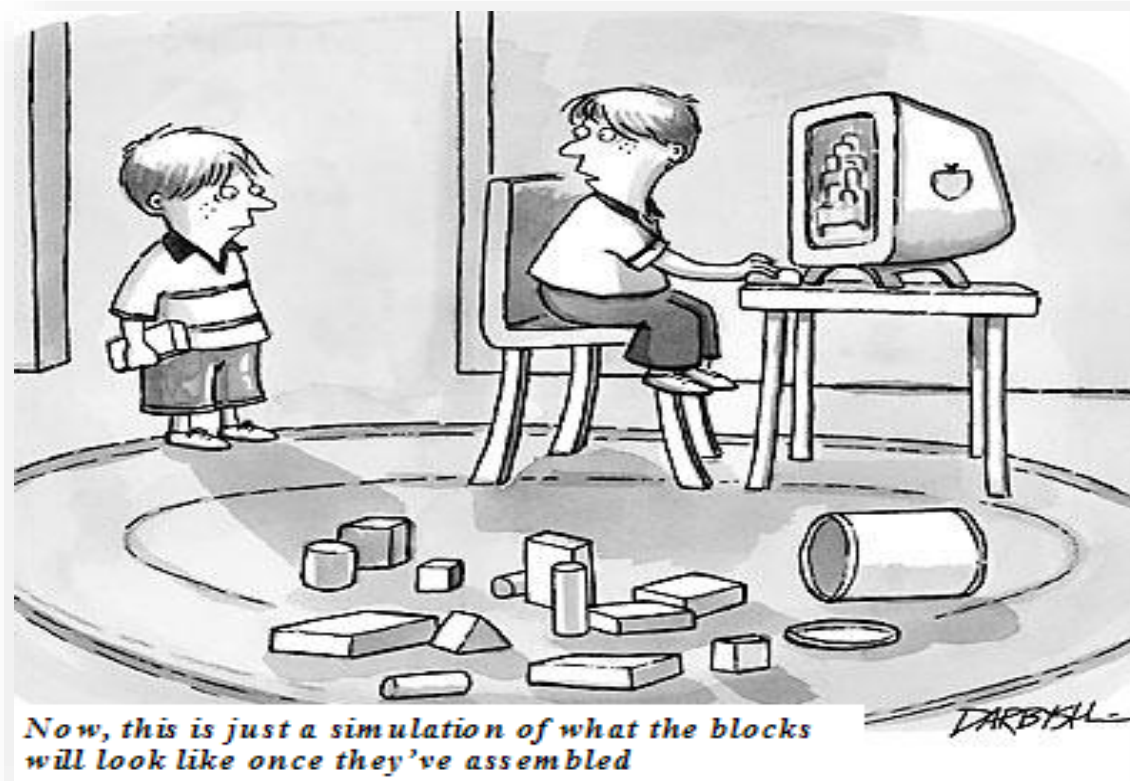


# ECE569

## Module 18

---



- Thread Synchronization

# Thread Synchronization

---

- **Fundamental concept in parallel programming**
- **Threads can access each other's results through shared and global memory**
  - They can work together
- **What if a thread reads a result before another thread writes it**
  - Need synchronization

# Barrier

---

- **Basic thread coordination mechanism**
- **When a thread calls `__syncthreads()`**
  - Forms a barrier in the thread execution path
  - Holds each thread at calling location until every thread **in the block** reaches that location
    - All threads complete a phase before moving onto next phase
      - Make sure that all threads read neighboring pixel values before those values are updated
    - Used to avoid RAW/WAR hazards when accessing shared or global memory
    - Does not synchronize threads from two different blocks



No one is left behind!

# Barrier Example: How many barriers are needed?

---

Shift elements of array to left by one element

:

```
int idx = threadIdx.x
```

```
__shared__ int array[128]
```

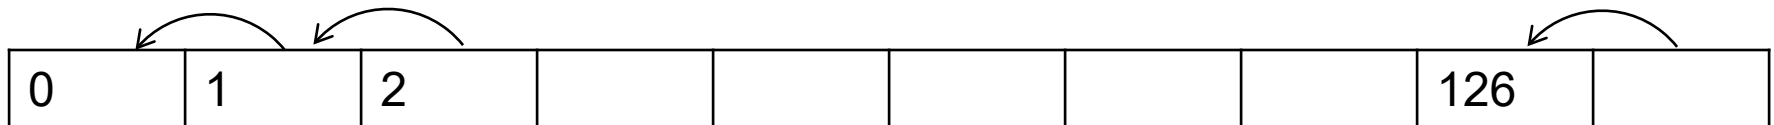
```
array[idx] = threadIdx.x
```

```
// initialize each element to thread index
```

```
If (idx < 127)
```

```
    array[idx] = array[idx+1];
```

:

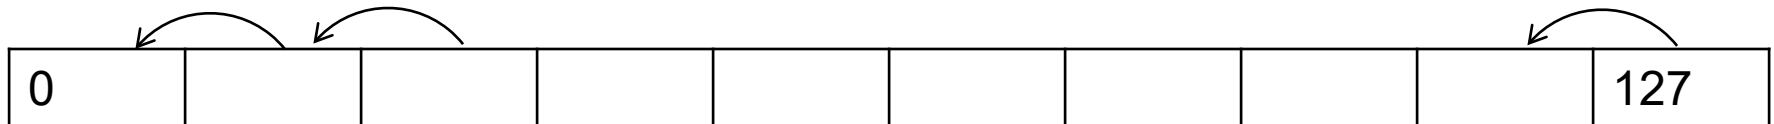


# Barrier Example: How many barriers are needed?

---

Shift elements of array to left by one element

```
:  
int idx = threadIdx.x  
__shared__ int array[128]  
array[idx] = threadIdx.x  
__syncthreads() //make sure all values written  
// initialize each element to thread index  
If (idx < 127)  
    array[idx] = array[idx+1];  
:
```



**Any other barrier(s)?**

# Barrier Example: How many barriers are needed?

---

Shift elements of array to left by one element

```
:
int idx = threadIdx.x
__shared__ int array[128]
array[idx] = threadIdx.x
__syncthreads()//make sure all values written
// initialize each element to thread index
If (idx < 127)
    array[idx] = array[idx+1];
```

**We need to complete reading [idx+1] from all entries before writing into [idx]! How do we achieve this?**

# Barrier Example: How many barriers are needed?

---

Shift elements of array to left by one element

:

```
int idx = threadIdx.x
```

```
__shared__ int array[128]
```

```
array[idx] = threadIdx.x
```

```
__syncthreads() //make sure all values written  
// initialize each element to thread index
```

```
If (idx < 127) {
```

```
    int temp = array[idx+1]; //declare local
```

```
    __syncthreads();
```

```
    array[idx] = temp;
```

```
    __syncthreads();
```

```
//last one ensures all write operations are  
completed before array is accessed again later
```

# Which one will have correct functionality without synchthreads()

```
__global__ void my_function() {  
    __shared__ int s[1024];  
    int i=threadIdx.x;  
    __synchthreads();  
    s[i]=s[i-1] // 1  
    __synchthreads();  
    if(i%2) s[i]=s[i-1]; // 2  
    __synchthreads();  
    s[i] = (s[i-1]+s[i]+s[i-1])/3.0 // 3  
    printf("s[%d]=%f\n", I, s[i]);  
    __synchthreads;
```

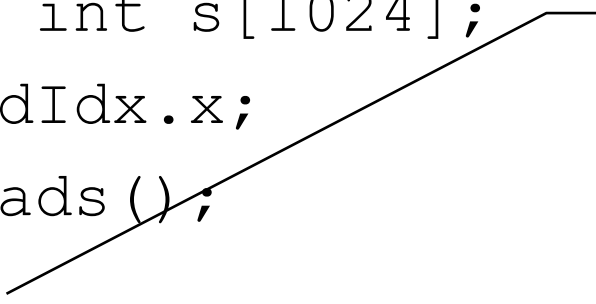


# Which one will have correct functionality without synchthreads()

---

```
__global__ void my_function() {  
    __shared__ int s[1024];  
    int i=threadIdx.x;  
    __synchthreads();  
    s[i]=s[i-1] // 1  
    __synchthreads();  
    if(i%2) s[i]=s[i-1]; // 2  
    __synchthreads();  
    s[i] = (s[i-1]+s[i]+s[i-1])/3.0 // 3  
    printf("s[%d]=%f\n", i, s[i]);  
    __synchthreads;  
}
```

```
int temp=s[i-1];  
__synchthreads();  
s[i]=temp;
```



# Which one will have correct functionality without synchthreads()

```
__global__ void my_function() {  
    __shared__ int s[1024];  
    int i=threadIdx.x;  
    __synchthreads();  
    s[i]=s[i-1] // 1  
    __synchthreads();  
    if(i%2) s[i]=s[i-1]; // 2  
    __synchthreads();  
    s[i] = (s[i-1]+s[i]+s[i-1])/3.0 // 3  
    printf("s[%d]=%f\n", i, s[i]);  
    __synchthreads;  
}
```

Only evens write  
Reads are all  
from odds.

# Which one will have correct functionality without synchthreads()

```
__global__ void my_function() {  
    __shared__ int s[1024];  
    int i=threadIdx.x;  
    __synchthreads();  
    s[i]=s[i-1]  
    __synchthreads();  
    if(i%2) s[i]=s[i-1];  
    __synchthreads();  
    s[i] = (s[i-1]+s[i]+s[i-1])/3.0  
    printf("s[%d]=%f\n", i, s[i]);  
    __synchthreads;
```

```
float temp =  
    (s[i-1]+s[i]+s[i-1])/3.0  
    __synchthreads();  
    s[i]=temp;  
    __synchthreads; // 1  
    printf();
```

```
// 2  
// 3
```

# Next

---

- **Global, Shared Memory**
  - Static vs. Dynamic Shared Memory
  - Coalesced Memory
  - Code Review