


# ECE569

## Module 34


---



- Reduction: Stride Pattern – Global and Shared Memory

# Kernel: Global Memory – Stride Pattern –Round 0

```
__global__ void global_reduce_stride(float* d_out, float* d_in) {  
    int myId = threadIdx.x + blockDim.x * blockIdx.x;  
    int tid = threadIdx.x;  
    // do reduction in global mem
```




Round0:  
Stride amount 1

```
// thread 0 writes result for this block back to global mem
```

```
}
```

# Kernel: Global Memory – Stride Pattern –Round 1

```
__global__ void global_reduce_stride(float* d_out, float* d_in) {  
    int myId = threadIdx.x + blockDim.x * blockIdx.x;  
    int tid = threadIdx.x;  
    // do reduction in global mem
```




Round1:  
Stride amount 2

```
// thread 0 writes result for this block back to global mem
```

```
}
```

# Kernel: Global Memory – Stride Pattern –Round 2

```
__global__ void global_reduce_stride(float* d_out, float* d_in) {  
    int myId = threadIdx.x + blockDim.x * blockIdx.x;  
    int tid = threadIdx.x;  
    // do reduction in global mem
```




Round2:  
Stride amount 4

```
// thread 0 writes result for this block back to global mem
```

```
}
```

# Kernel: Global Memory – Stride Pattern– Round 3

```
__global__ void global_reduce_stride(float* d_out, float* d_in) {  
    int myId = threadIdx.x + blockDim.x * blockIdx.x;  
    int tid = threadIdx.x;  
    // do reduction in global mem
```




```
// thread 0 writes result for this block back to global mem
```

```
}
```

# Kernel: Global Memory – Stride Pattern


```
__global__ void global_reduce_stride(float* d_out, float* d_in) {  
    int myId = threadIdx.x + blockDim.x * blockIdx.x;  
    int tid   = threadIdx.x;
```

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15



Block 0

16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31



Block 1

# Kernel: Global Memory – Stride Pattern

```

__global__ void global_reduce_stride(float* d_out, float* d_in) {
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int tid = threadIdx.x;
    // do reduction in global mem
    for(int stride = _____; stride < _____; stride = _____) {


        if( _____ ) {

            d_in[ _____ ] += d_in[ _____ ]; }

        }

    // thread 0 writes result for this block back to global mem
    if ( _____ ) {
        d_out[ _____ ] = d_in[ _____ ]; }
}

```



# Kernel: Global Memory – Strided Access

---

```
__global__ void global_reduce_stride(float* d_out, float* d_in) {
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int tid = threadIdx.x;
    // do reduction in global mem
    for(int stride = 1; stride < blockDim.x; stride *= 2)
    {
        __syncthreads();
        if(myId % (2*stride) == 0) {
            d_in[myId] += d_in[myId+stride]; }
    }

    // thread 0 writes result for this block back to global mem
    if (tid == 0) {
        d_out[blockIdx.x] = d_in[myId]; }
}
```




# Reduction - Tesla P100;compute v6.0;

---

n: 1<<20

Version		Time (ms)	19.9X
serial		3.27400	
global reduce stride – naïve		0.16450	



# Kernel: Shared Memory – Strided Access

```

__global__ void shared_reduce_stride(float* d_out, float* d_in){

//shared_reduce_stride<<<blocks,threads,threads*sizeof(float)>>>
int myId = threadIdx.x + blockDim.x * blockIdx.x;
int tid  = threadIdx.x;
// load shared mem from global mem
_____

// make sure entire block is loaded!
// do reduction in shared memory
for(int stride = _____; stride < _____; stride = _____){


    if( _____ ) {
        _____
    }

}

// thread 0 writes result for this block back to global mem
if ( _____ ) {
    _____
}

}

```



# Reduction - Tesla P100;compute v6.0;

n: 1<<20

Version	Time (ms)
serial	3.27400
global reduce stride – naïve	0.16450
shared stride reduce	0.15835

20.7X



Unusual we  
should have had  
better  
performance!

# Kernel: Shared Memory – Strided Access

```
__global__ void shared_reduce_stride(float* d_out, float* d_in){
extern __shared__ float sdata[];


// shared_reduce<<<blocks,threads,threads*sizeof(float)>>>
int myId = threadIdx.x + blockDim.x * blockIdx.x;
int tid  = threadIdx.x;

// load shared mem from global mem
sdata[tid] = d_in[myId];


// make sure entire block is loaded!

// do reduction in shared memory
for(int stride = 1; stride < blockDim.x; stride *= 2)  {
    __syncthreads();
    if(myId % (2*stride) == 0) {
        sdata[tid] += sdata[tid+stride]; }
    }

// thread 0 writes result for this block back to global mem
if (tid == 0) {
    d_out[blockIdx.x] = sdata[tid]; }
}
```




# Observations on the Stride Pattern



- Global
  - Memory access pattern

- Each thread
  - responsible for an even-index location of the partial sum vector (location of responsibility)
- In each step
  - second input comes from an increasing distance away
- After each step
  - half of the threads are no longer needed

# Observations on the Stride Pattern




In each iteration, two paths will be sequentially traversed for each warp

- Threads that perform addition and threads that do not

- Half or fewer of threads will be executing after the first step
  - All odd-index threads are disabled after first step

# Kernel: Shared Memory – Stride Pattern



**Next:**

**How to reorganize workload assignment to avoid divergence?**