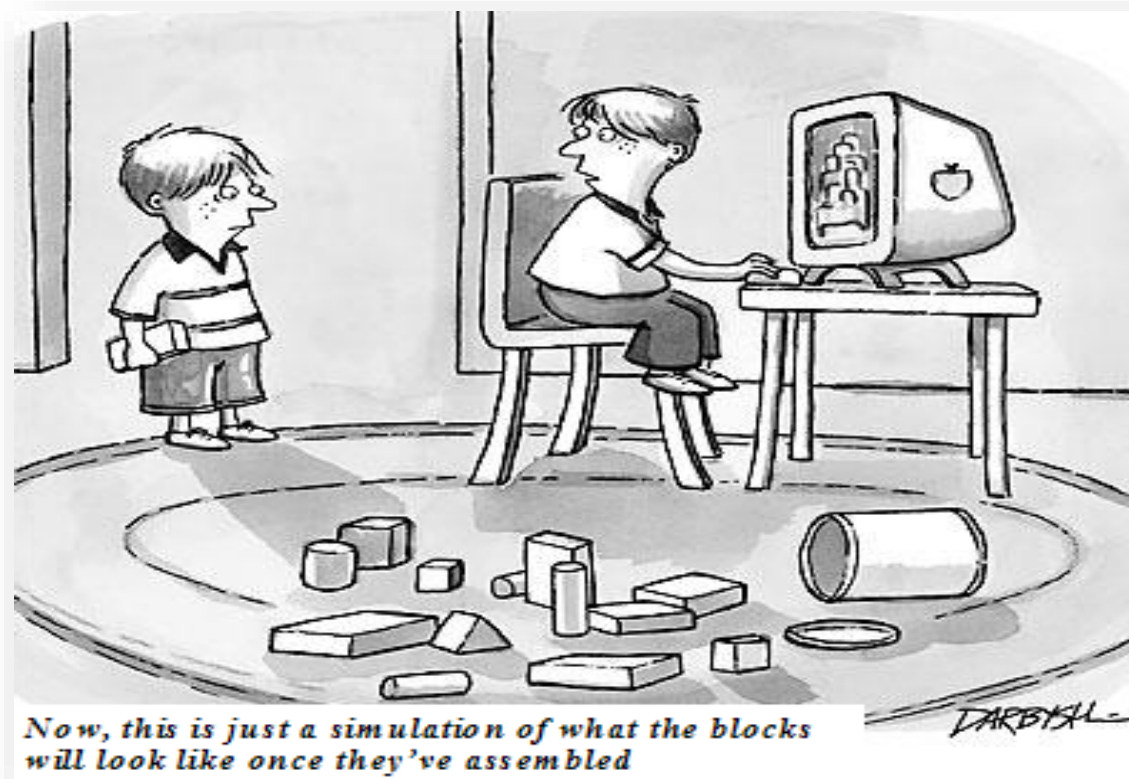


ECE569

Module 38



- Scan – Version 2 Implementation and Analysis


Implementation

```
__global__ void my_kernel(float *X, float *Y, int InputSize) {
    __shared__ float XY[SECTION_SIZE];
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int tid = threadIdx.x;
    // Read from global to shared memory
    If ( _____ )
        XY[ _____ ] = X[ _____ ]
    // Write the for loop
    for (int stride= _____; stride _____; stride= _____) {

    }
    // Write back to global memory
    If ( _____ )
        Y[ _____ ] = XY{ _____ };
}
```


Analysis of parallel version-2

- **Running Sum of N elements**
 - What is the step complexity?



Analysis of parallel version-2

- **Running Sum of N elements**
 - What is the work complexity?



Evaluation

	Step Complexity	Work Complexity
Version 1 (Naïve)	N	$O(N^2)$
Version 2	logN	NlogN

Wait a minute: my serial code has better work complexity

```
out[0] = x[0];  
for (i=1; i<n; i++)  
    out[i] = out[i-1] + x[i];
```

$O(N)$!

A factor of $\log(n)$ can hurt: 10x for 1024 elements!

- A parallel algorithm can be slower than a sequential one when execution resources are saturated from low work efficiency

Next: We can improve the work efficiency

- As we know, the fastest parallel way to produce sum values for a set of values is a reduction tree.
- With sufficient execution units, a reduction tree can generate the sum for N values in $\log N$ time units.
- Solution:
 - What if we do a reduction to find the sum (REDUCE)
 - And then use intermediate sums to use as critical piece of information to find the other running sums (CONSTRUCT)

Reading

- **Mark Harris, Parallel Prefix Sum with CUDA**
 - https://developer.nvidia.com/gpugems/GPUGems3/gpugems3_ch39.html