ECE569
Module 35

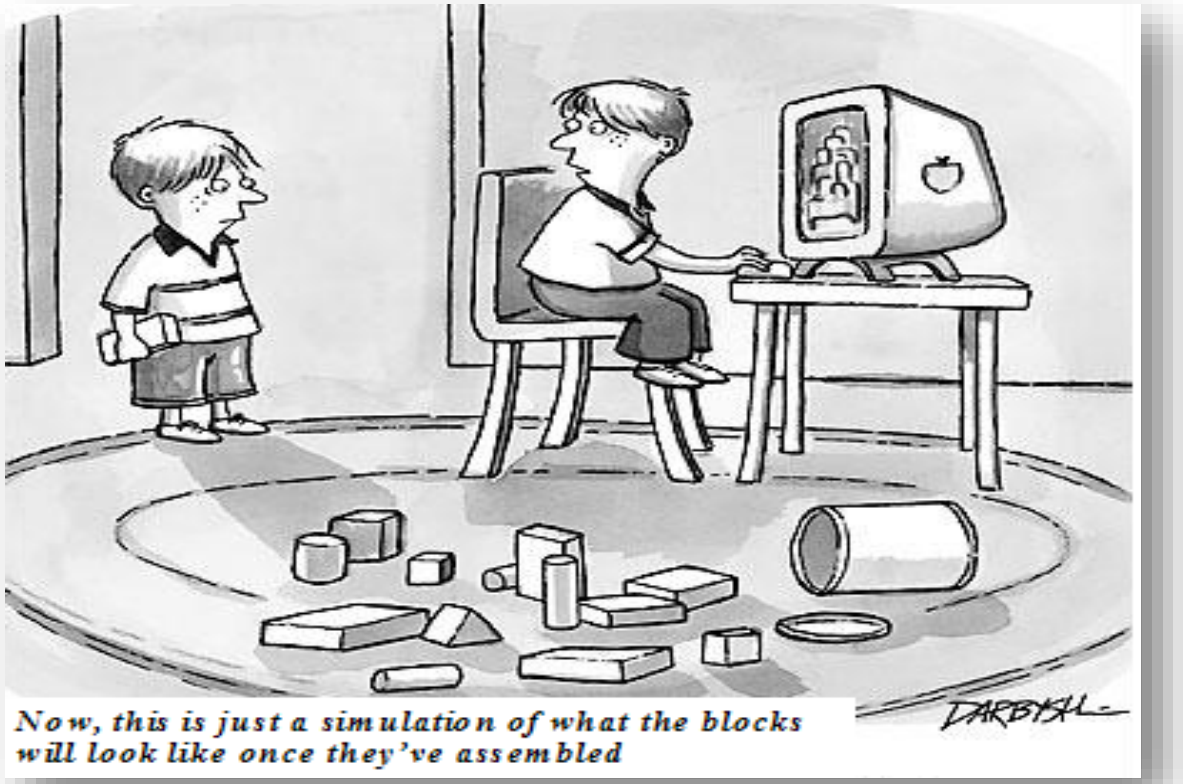

- Reduction – Stride Pattern – Shared Memory and Branch Divergence

Reduction - Tesla P100;compute v6.0;

Version	Time (ms)
serial	3.27400
global reduce stride – naïve	0.16450
shared stride reduce	0.15835

n: 1<<20

20.7X

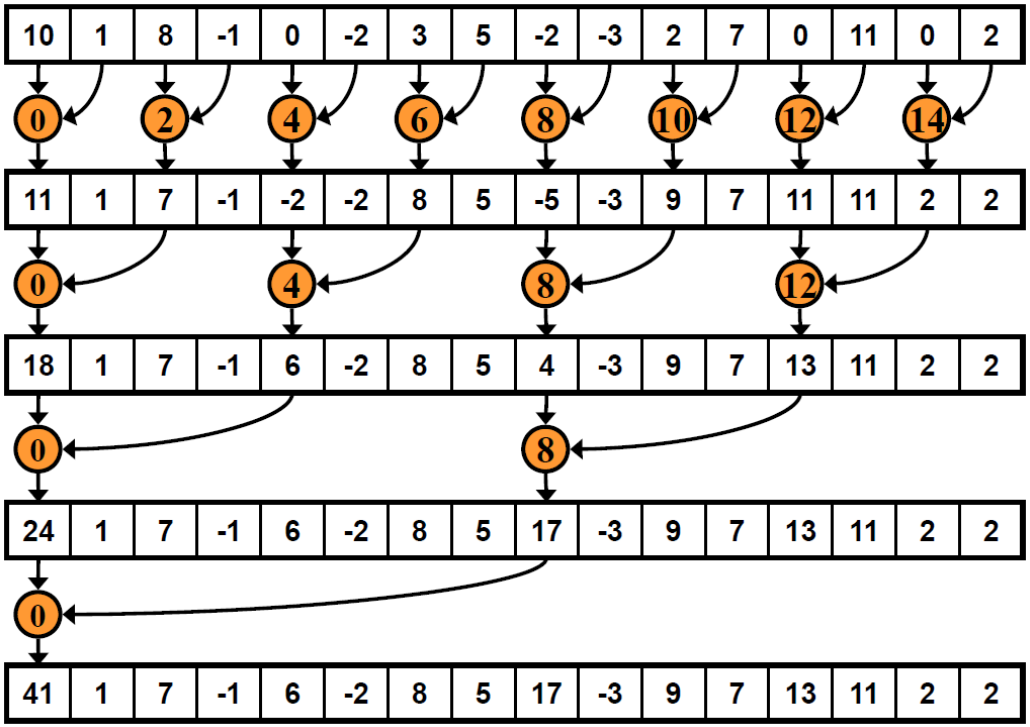

Unusual we should have had better performance!

How to reorganize workload assignment to avoid divergence?

Kernel: Shared Memory – Stride Pattern

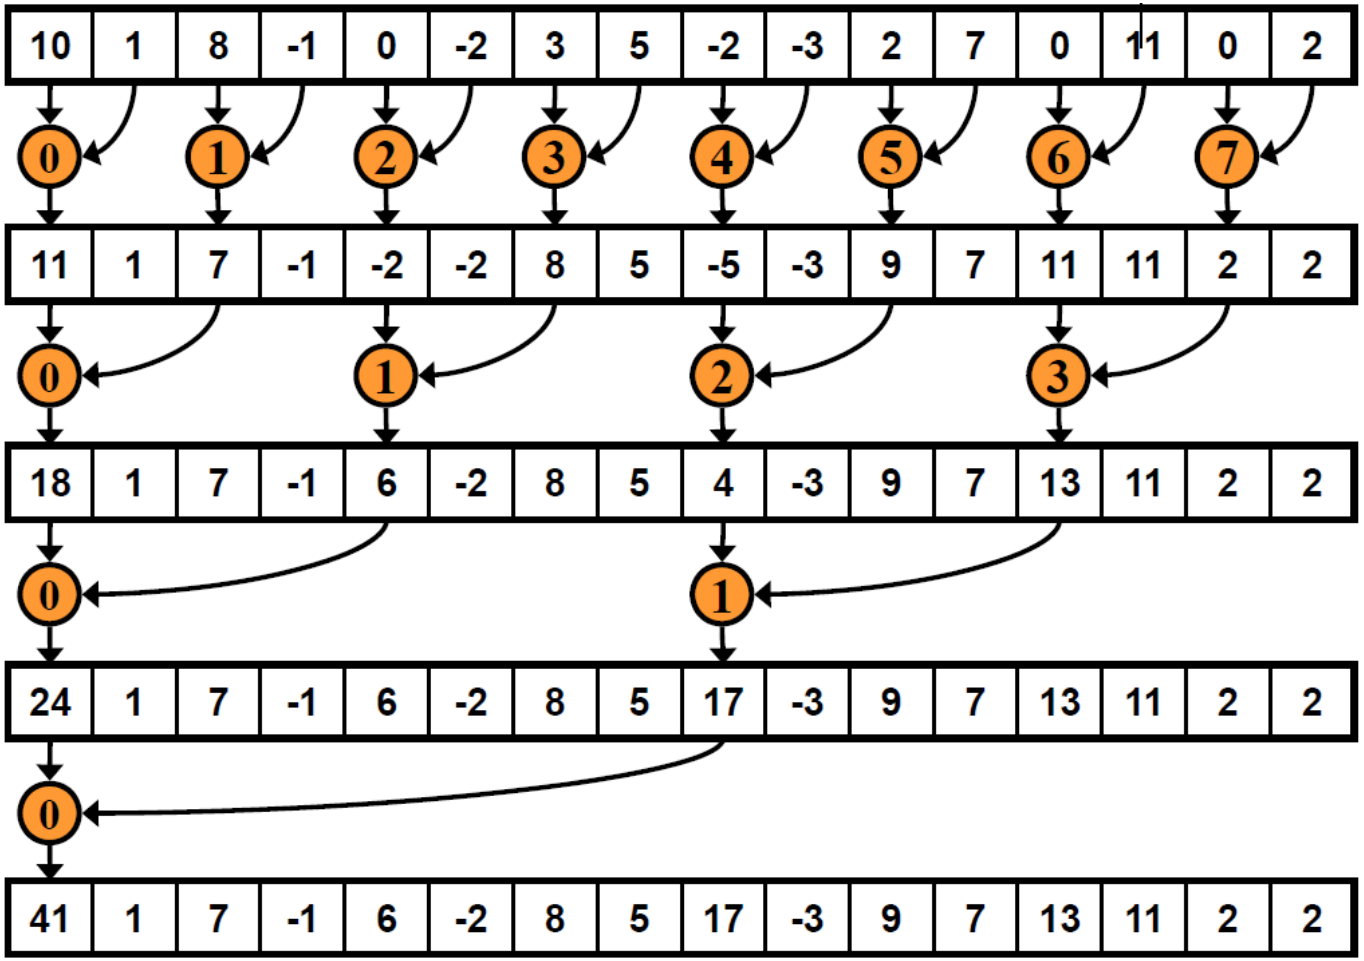

Kernel: Shared Memory – Stride Pattern

```
__global__ void shared_reduce_stride(float* d_out, float* d_in){
  extern __shared__ float sdata[];
  // shared_reduce<<<blocks,threads,threads*sizeof(float)>>>
  int myId = threadIdx.x + blockDim.x * blockIdx.x;
  int tid  = threadIdx.x;
  // load shared mem from global mem
  sdata[tid] = d_in[myId];
  // make sure entire block is loaded!
  // do reduction in shared memory
  for(int stride = 1; stride < blockDim.x; stride *= 2)  {
    __syncthreads();
    if(myId % (2*stride) == 0) {
      sdata[tid] += sdata[tid+stride]; }
    }
// thread 0 writes result for this block back to global mem
 if (tid == 0) {
        d_out[blockIdx.x] = sdata[tid]; }
}
```

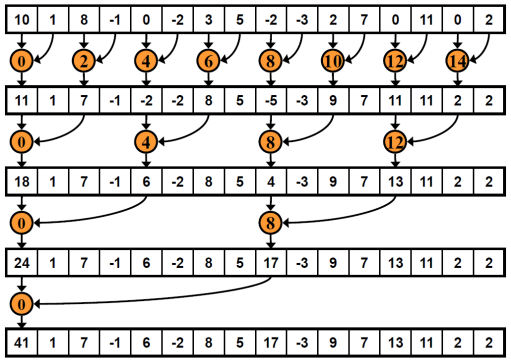

Kernel: Shared Memory – No Diverge – Stride Pattern

```
__global__ void shared_reduce_stride_nodiverge (float* d_out,
float* d_in){
  extern __shared__ float sdata[];
  // shared_reduce<<<blocks,threads,threads*sizeof(float)>>>
  int myId = threadIdx.x + blockDim.x * blockIdx.x;
  int tid  = threadIdx.x;
  // load shared mem from global mem
  sdata[tid] = d_in[myId];
  // make sure entire block is loaded!
  // do reduction in shared memory
  for(int stride = 1; stride < blockDim.x; stride *= 2)  {
    __syncthreads();
```

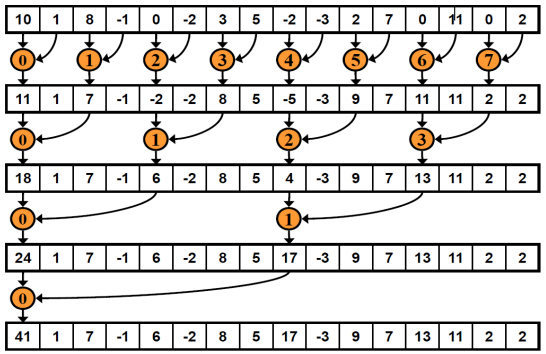


```
 }
 // thread 0 writes result for this block back to global mem
 if (tid == 0) {
        d_out[blockIdx.x] = sdata[tid]; }
}
```

Kernel: Shared Memory – No Diverge – Stride Pattern

```
__global__ void shared_reduce_stride_nodiverge (float* d_out,
float* d_in){
  extern __shared__ float sdata[];
  // shared_reduce<<<blocks,threads,threads*sizeof(float)>>>
  int myId = threadIdx.x + blockDim.x * blockIdx.x;
  int tid  = threadIdx.x;
  // load shared mem from global mem
  sdata[tid] = d_in[myId];
  // make sure entire block is loaded!
  // do reduction in shared memory
  for(int stride = 1; stride < blockDim.x; stride *= 2)  {
    __syncthreads();
    int index = 2*stride*tid;
    if( index < blockDim.x ) {
      sdata[index] += sdata[stride+index]; }
   }
// thread 0 writes result for this block back to global mem
 if (tid == 0) {
        d_out[blockIdx.x] = sdata[tid]; }
}
```

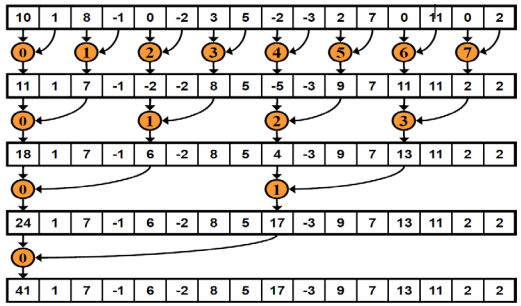

Bonus: No more expensive % operator

Reduction - Tesla P100;compute v6.0;

Version	Time (ms)
serial	3.27400
global reduce stride – naïve	0.16450
shared stride reduce	0.15835
shared_reduce_stride_nodiverge	0.09081

n: 1<<20

36.1X

1.7X!

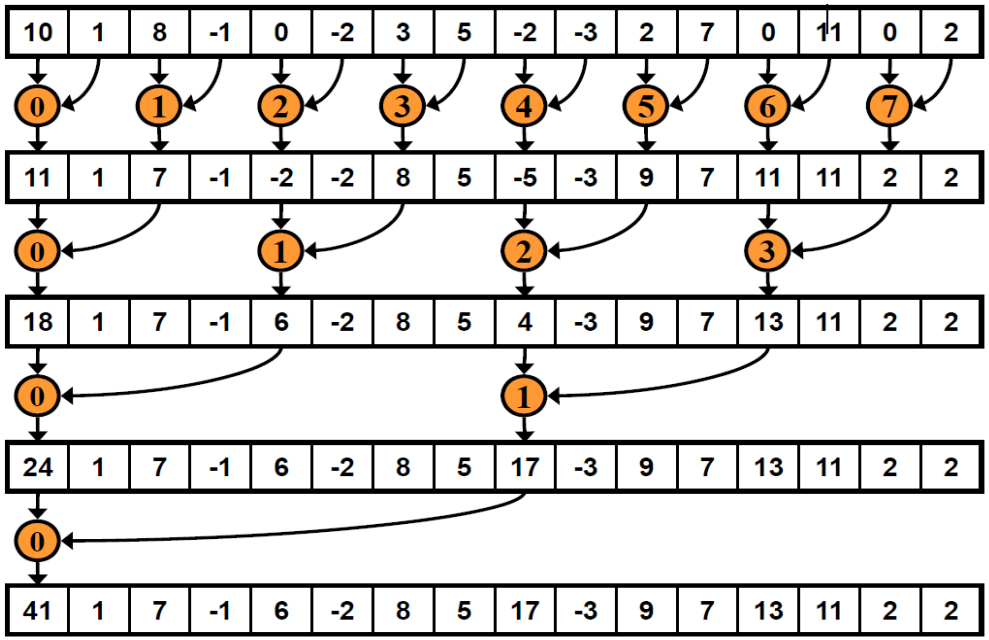

Observation on Stride Pattern-No divergence

- **Divergence free**
- **New problem:**
  - **Shared memory bank conflicts**
  - Will come back to reduction, first bank conflicts!

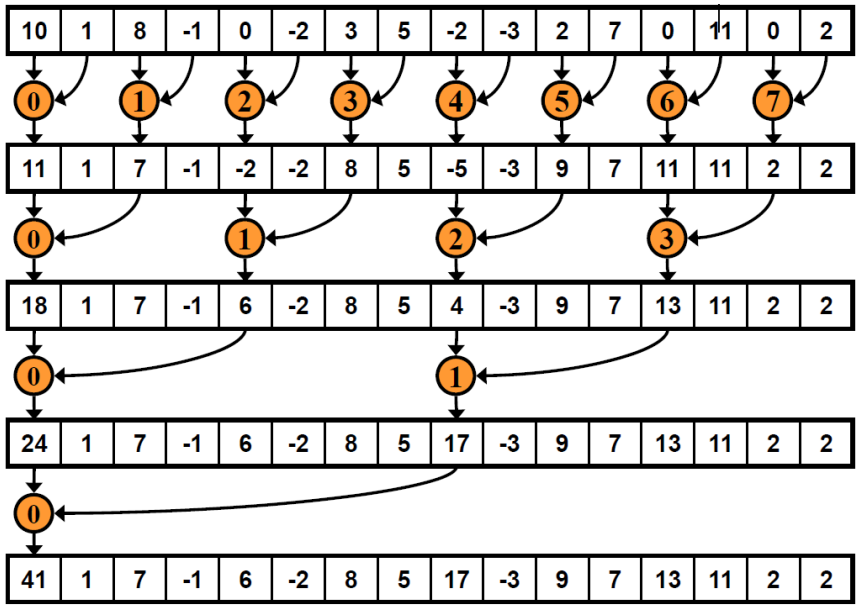

Shared Memory and Bank Conflicts

- Many threads access memory at the same time
  - To service more than one thread, memory is divided into independent banks
  - This layout essential to achieve high bandwidth
- Each SM has Shared Memory organized in 32 Memory banks

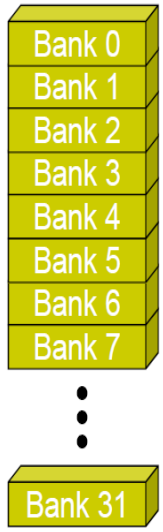

Shared Memory Architecture

- The 32 banks of the Shared Memory are organized like benches in a movie theater
  - You have multiple rows of benches
  - Each row has 32 benches which are separated and grouped in long columns
  - In each bench you can “seat” a family of four bytes (32bits total)
  - Note that a bank represents a column of benches in the movie theater, which is perpendicular to the screen

Organization of Memory Banks

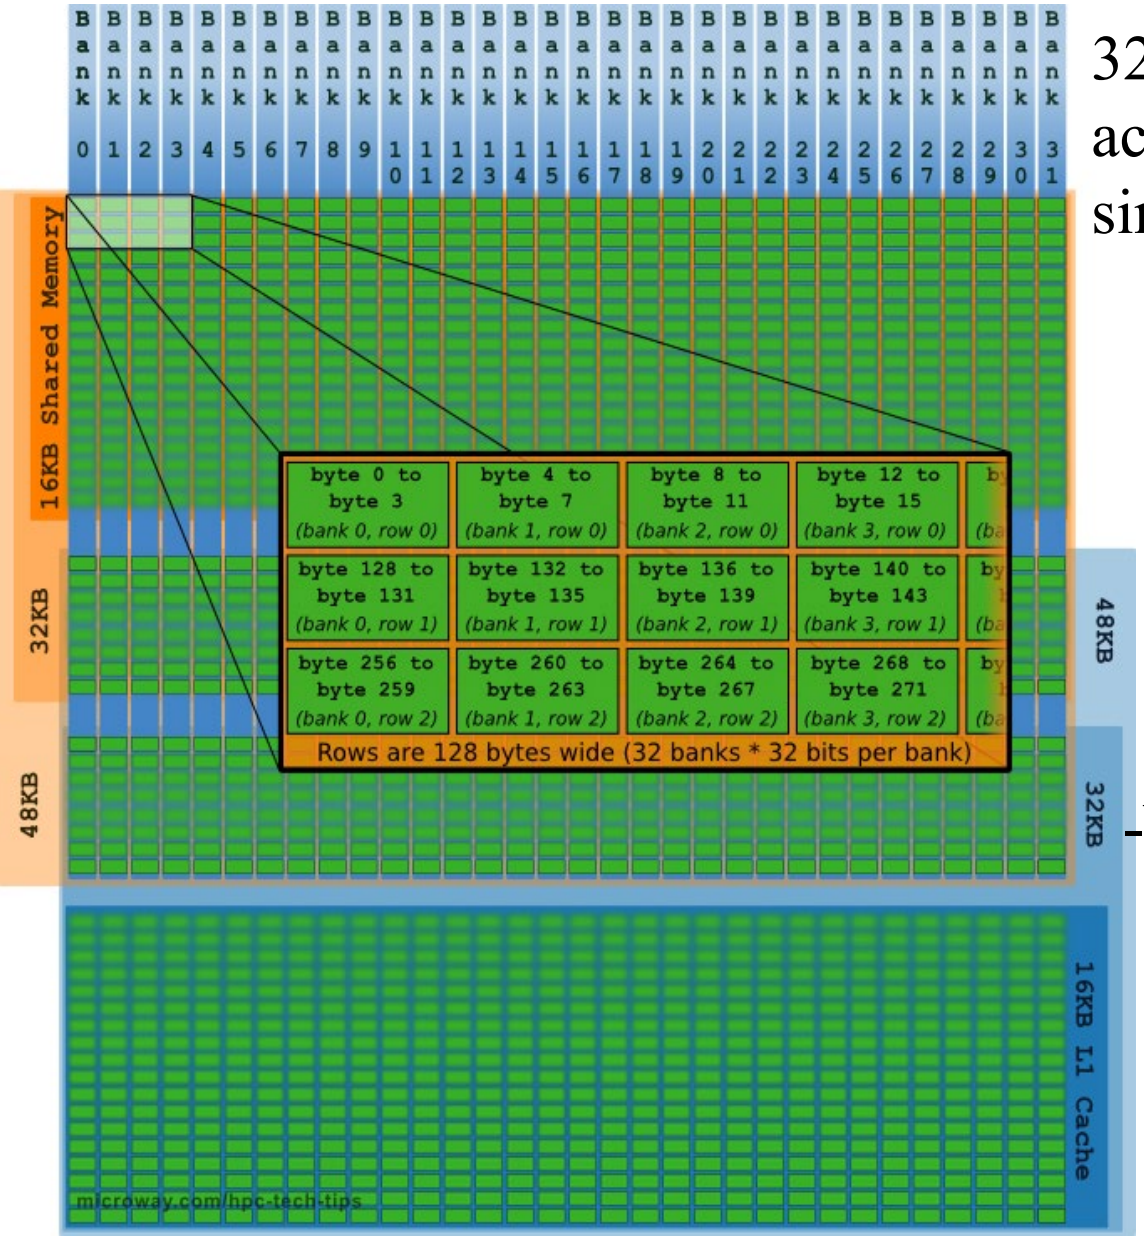

32 threads in a warp attempt to access shared memory simultaneously

Bank conflict: the scenario where two different threads access *different* words in the same bank

-

-No “bank conflicts” between threads of different warps

Bank Conflict Example

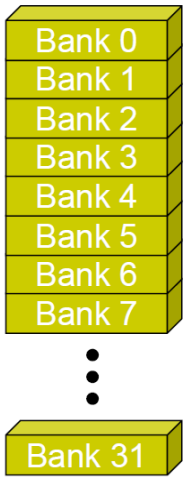

- Bank = (address of offset) % 32
- Example:
- 1D shared mem array, myShMemVar, of 1024 floats
  - myShMemVar[4]: accesses bank #4 (the fifth one – first row)
  - myShMemVar[31]: accesses bank #31 (the last one – first row)
  - **myShMemVar[50]:** access bank #18 (the 19th one – second row)
  - myShMemVar[128]: access bank #0 (the first one – fifth row)
  - **myShMemVar[178]:** access bank #18 (the 19th one – sixth row)
- NOTE: If, for instance, the third thread in a warp accesses myShMemVar[50] and the eight thread in the warp access myShMemVar[178], then you have a two-way bank conflict and the two transactions get serialized

Bank Conflict Examples(1): 4 Byte Words

Bank = (address of offset) % 32

- No Bank Conflicts
  - Linear addressing stride == 1

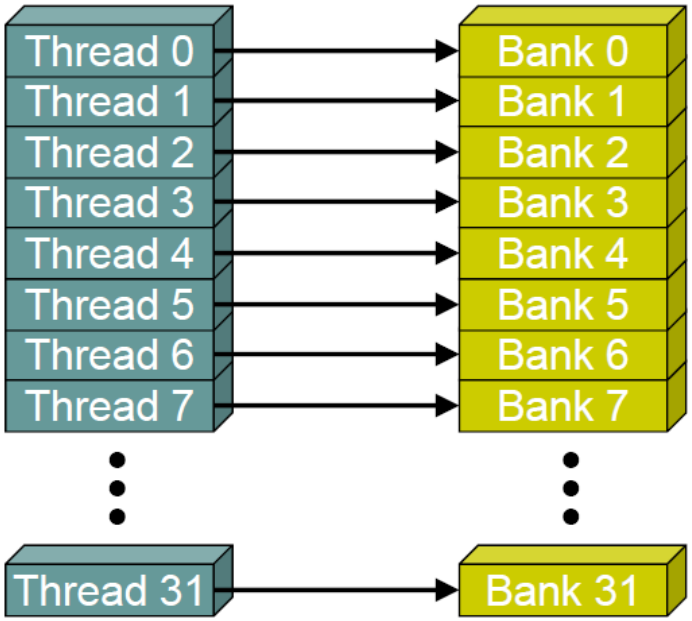

- No Bank Conflicts
  - Random 1:1 Permutation

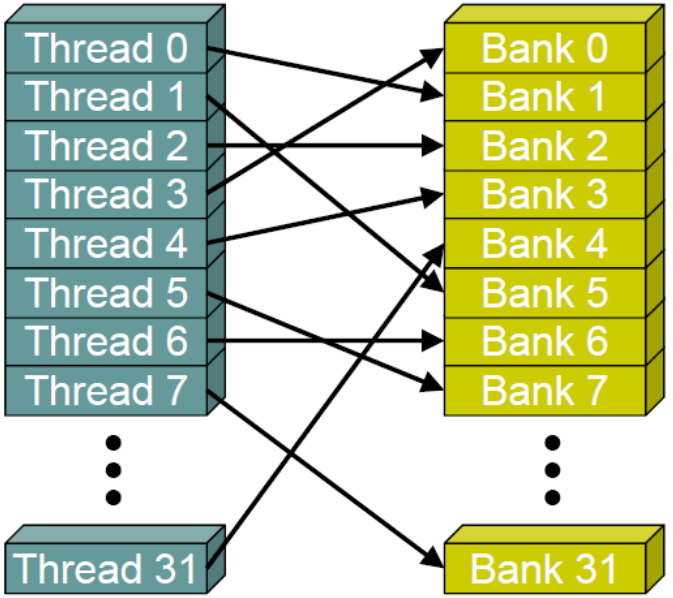

Bank Conflict Examples(2): 4 Byte Words

Bank = (address of offset) % 32

- sdata[threadIdx.x*2]++;

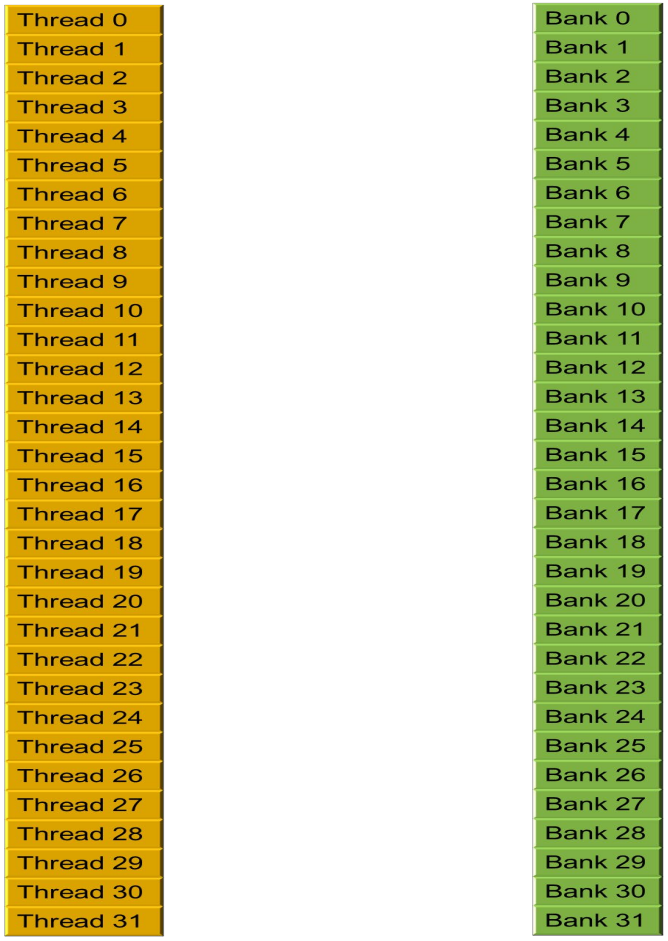

- sdata[threadIdx.x*8]++;

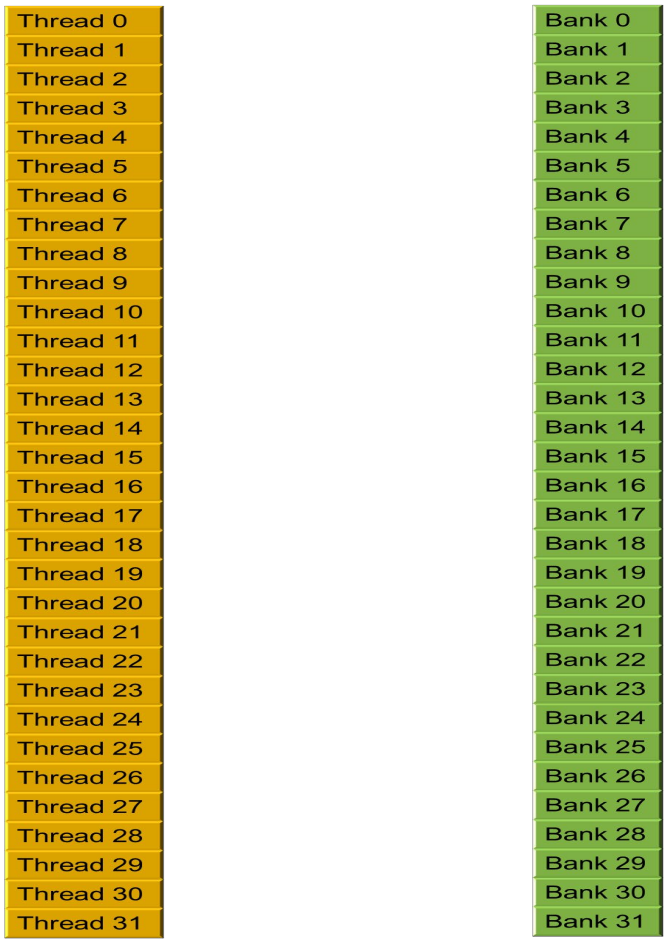

Which stride amount results with bank conflicts for all threads in a warp?

Bank Conflict Examples(3): 4 Byte Words

Which accesses have bank conflicts?

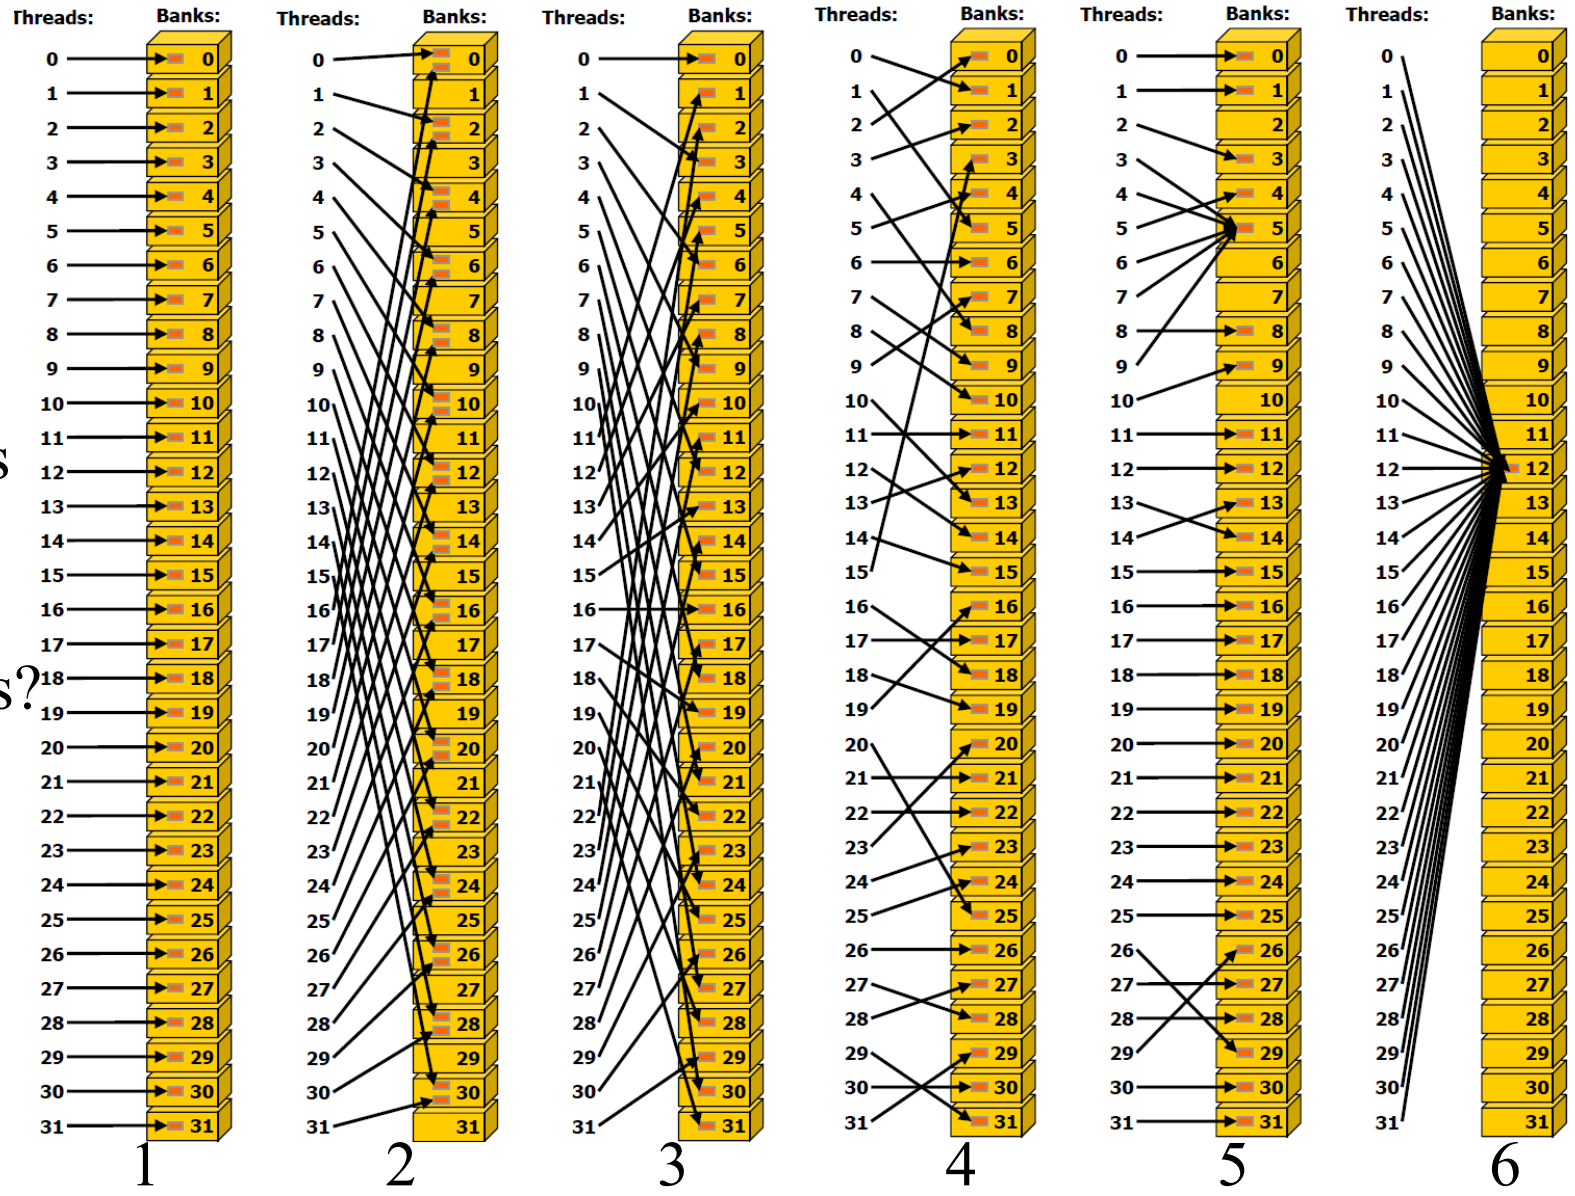

Bank Conflicts – Demo

```
__global__ void mykernel1(unsigned long long* time){
__shared__ float shared[1024];
// clock returns clock ticks  unsigned long long
startTime = clock();
//all threads accessing the same location (broadcast)
shared[0]++;
unsigned long long finishTime = clock();  *time =
(finishTime-startTime);
}
```

Replace **shared[0]++ with**

```
Kernel 2: shared[threadIdx.x]++;
Kernel 3: shared[threadIdx.x*4]++;
Kernel 4: shared[threadIdx.x*8]++;
Kernel 5: shared[threadIdx.x*32]++;
```

Refer to D2L→Demo→8.BankConflict

Kernel: Shared Memory – Stride Pattern

Next: How to eliminate bank conflicts?