



Module 44

CUDA Streams

GPU based Streaming Enabled Spectral
Correlation Density Function (SCD)

- Software Defined Radio (SDR)
 - Radio components in software
 - Rapid reconfiguration
 - higher spectral efficiency based on environmental conditions.
 - signal parameters such as transmission frequency or modulation scheme



Task of reliable and real-time signal classification involving non-cooperative communication is difficult

- Consider 5G Systems
 - Massive growth in number of edge devices, scaling to billions, in the wireless spectrum.
- Challenges:
 - Recognizing signals of a specific modulation type in real-time
 - Share the spectrum effectively without hampering the other transmissions at various priority levels

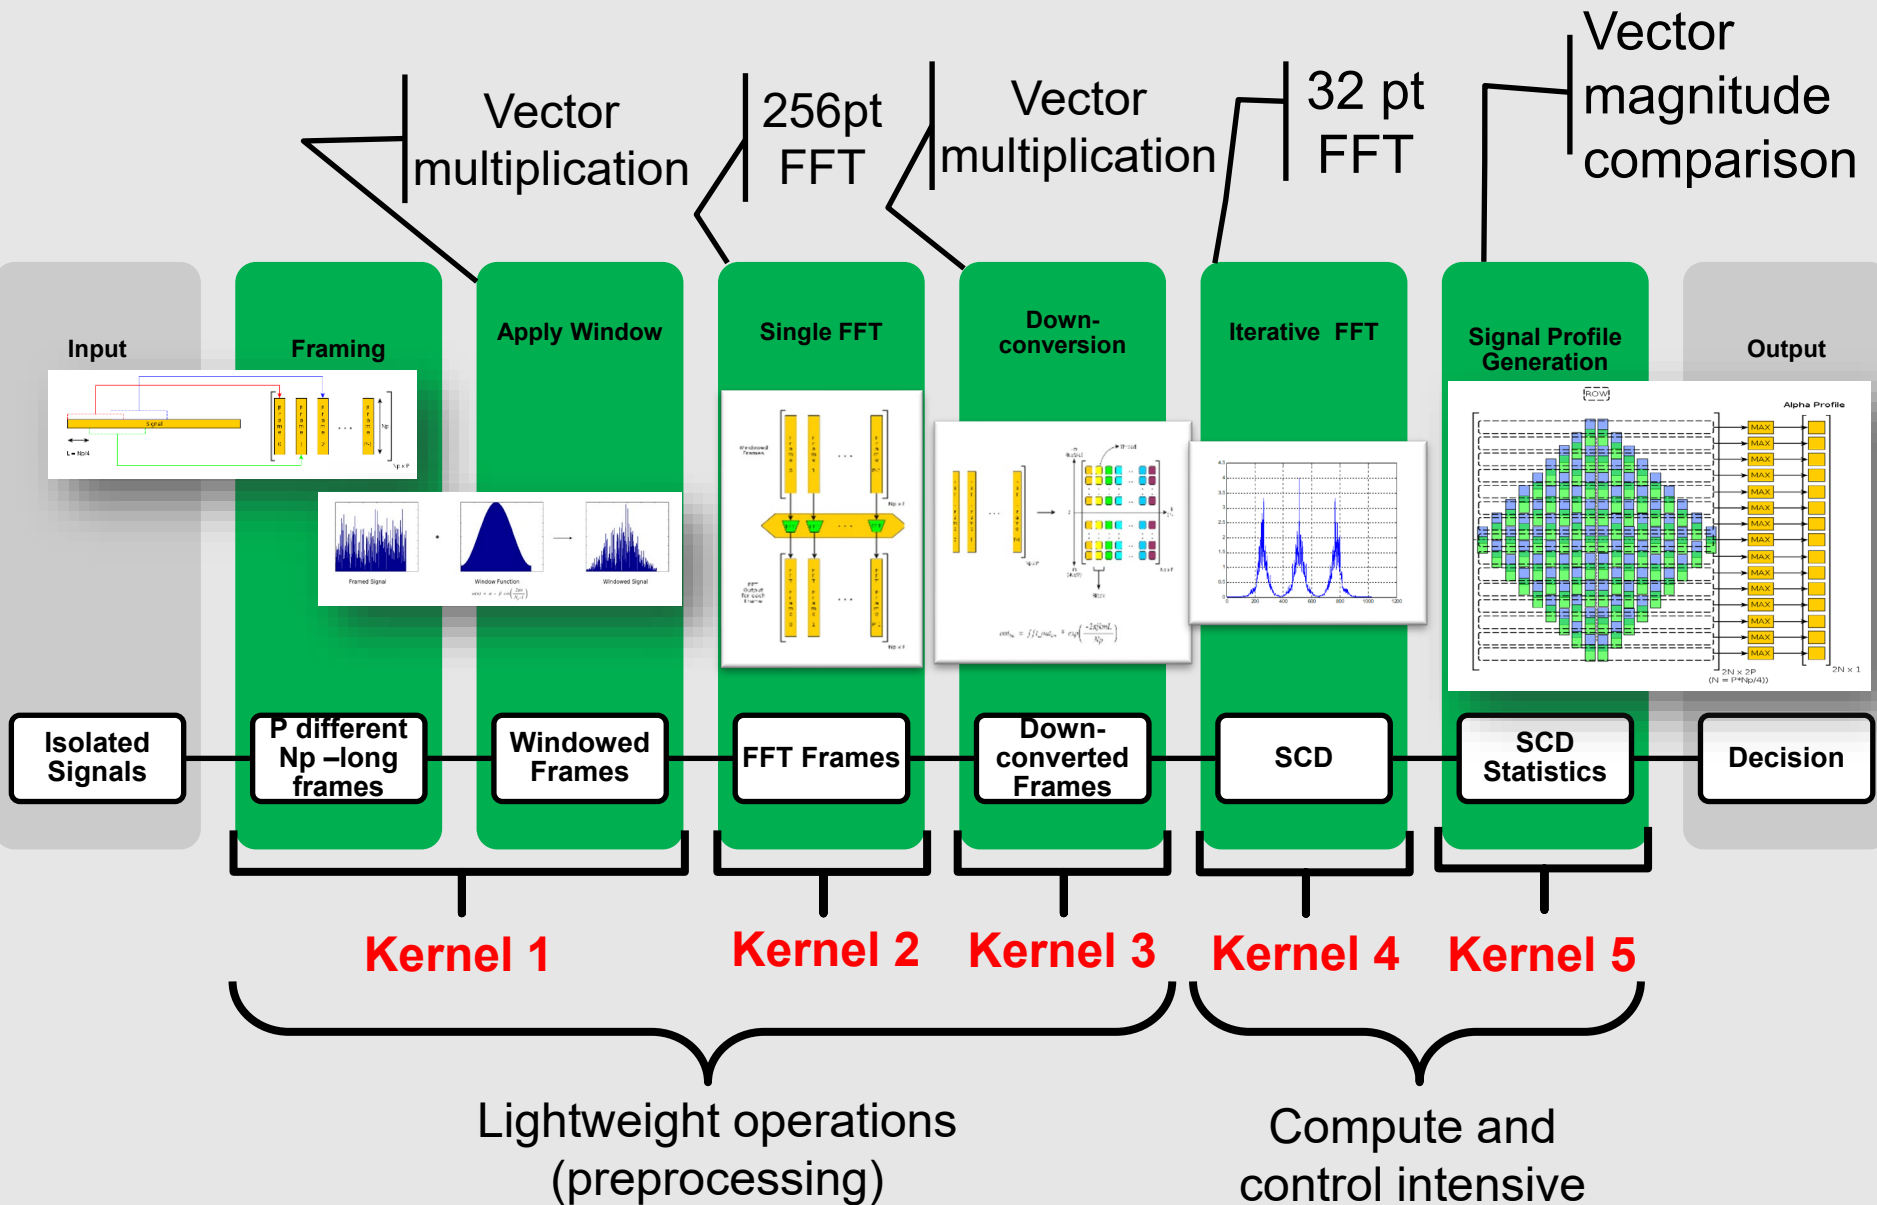
Accurate and fast modulation classification is needed!

- Spectral Correlation Density (SCD)
 - Signals that can be classified: BPSK, QPSK, 8PSK, 16QAM, GMSK, CPFSK, AM, FM, OFDM.
 - Relatively immune to noise
- Challenge:
 - Computational complexity


- High throughput SCD analysis based on GPU
 - Explore parallelization opportunities
 - Algorithmic
 - Full SCD
 - Quarter SCD (QSCD)
 - Architecture Specific
 - Kernel launch overhead
 - » Process multiple signals concurrently
 - » Data streaming
- Throughput Improvement
 - State of the art GPU based Full SCD 120 signals/second [1]
 - Our implementation [2] achieves 3300 signals/second

[1] [Nilangshu Bidyanta, Garrett Vanhoy, Mohammed Hirzallah, Ali Akoglu and Bo Ryu, “GPU and FPGA Based Architecture Design for Real-time Signal Classification,” Proc. 2015 Wireless Innovation Forum Conference on Wireless Communications Technologies and Software Defined Radio (WInnComm’15), March 24-26, 2015, San Diego, CA, pp. 70-79.


[2] Scott Marshall, Garrett Vanhoy, Ali Akoglu, Tamal Bose and Bo Ryu, “GPGPU based Parallel Implementation of Spectral Correlation Density Function,” *Journal of Signal Processing Systems*, vol. 92, pp.71-93, April 2019. DOI doi.org/10.1007/s11265-019-01448-7



- Step-1: Framing the signal
 - An arbitrary length signal split into P parts (32) of length N_p (256) each.
 - Each part overlaps with its predecessor.
 - The offset between the beginning of two consecutive parts is set to L where $L = N_p/4$.
 - Frames arranged column-wise - $N_p \times P$ matrix




- Step-2: Hamming Window
 - Steep cut-offs at both ends of the frame (step-1) introduce high frequency components
 - "raised cosine" function of length N_p
 - Window function remains the same for a given set of input parameters.




Framed Signal

*



Window Function

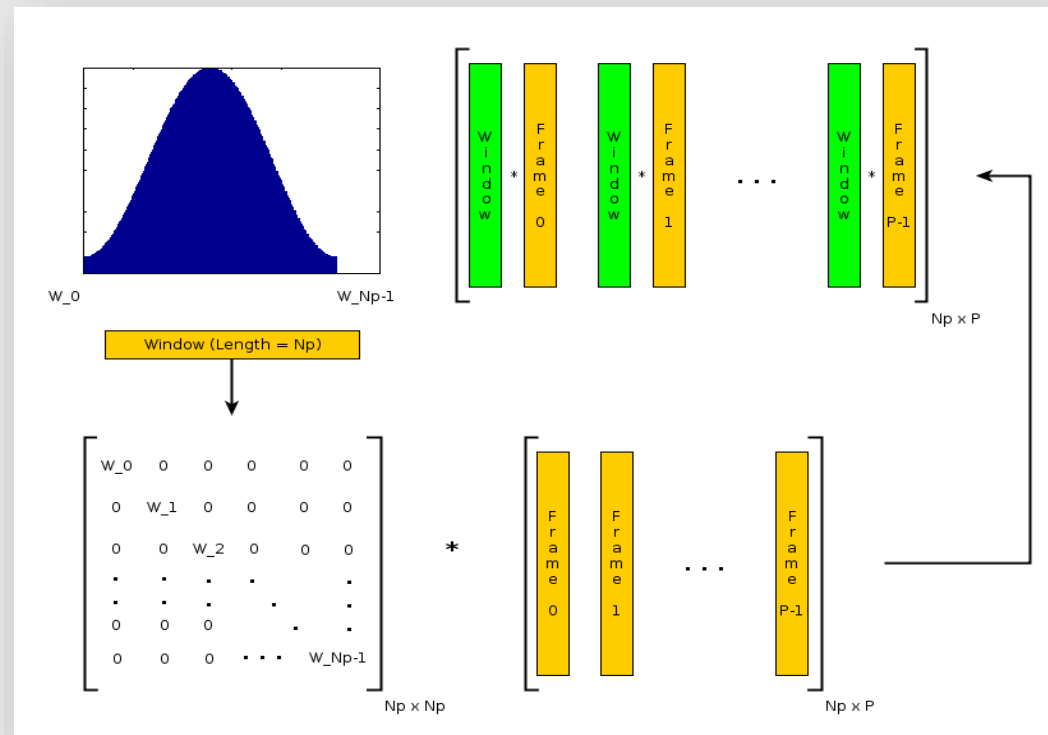


Windowed Signal

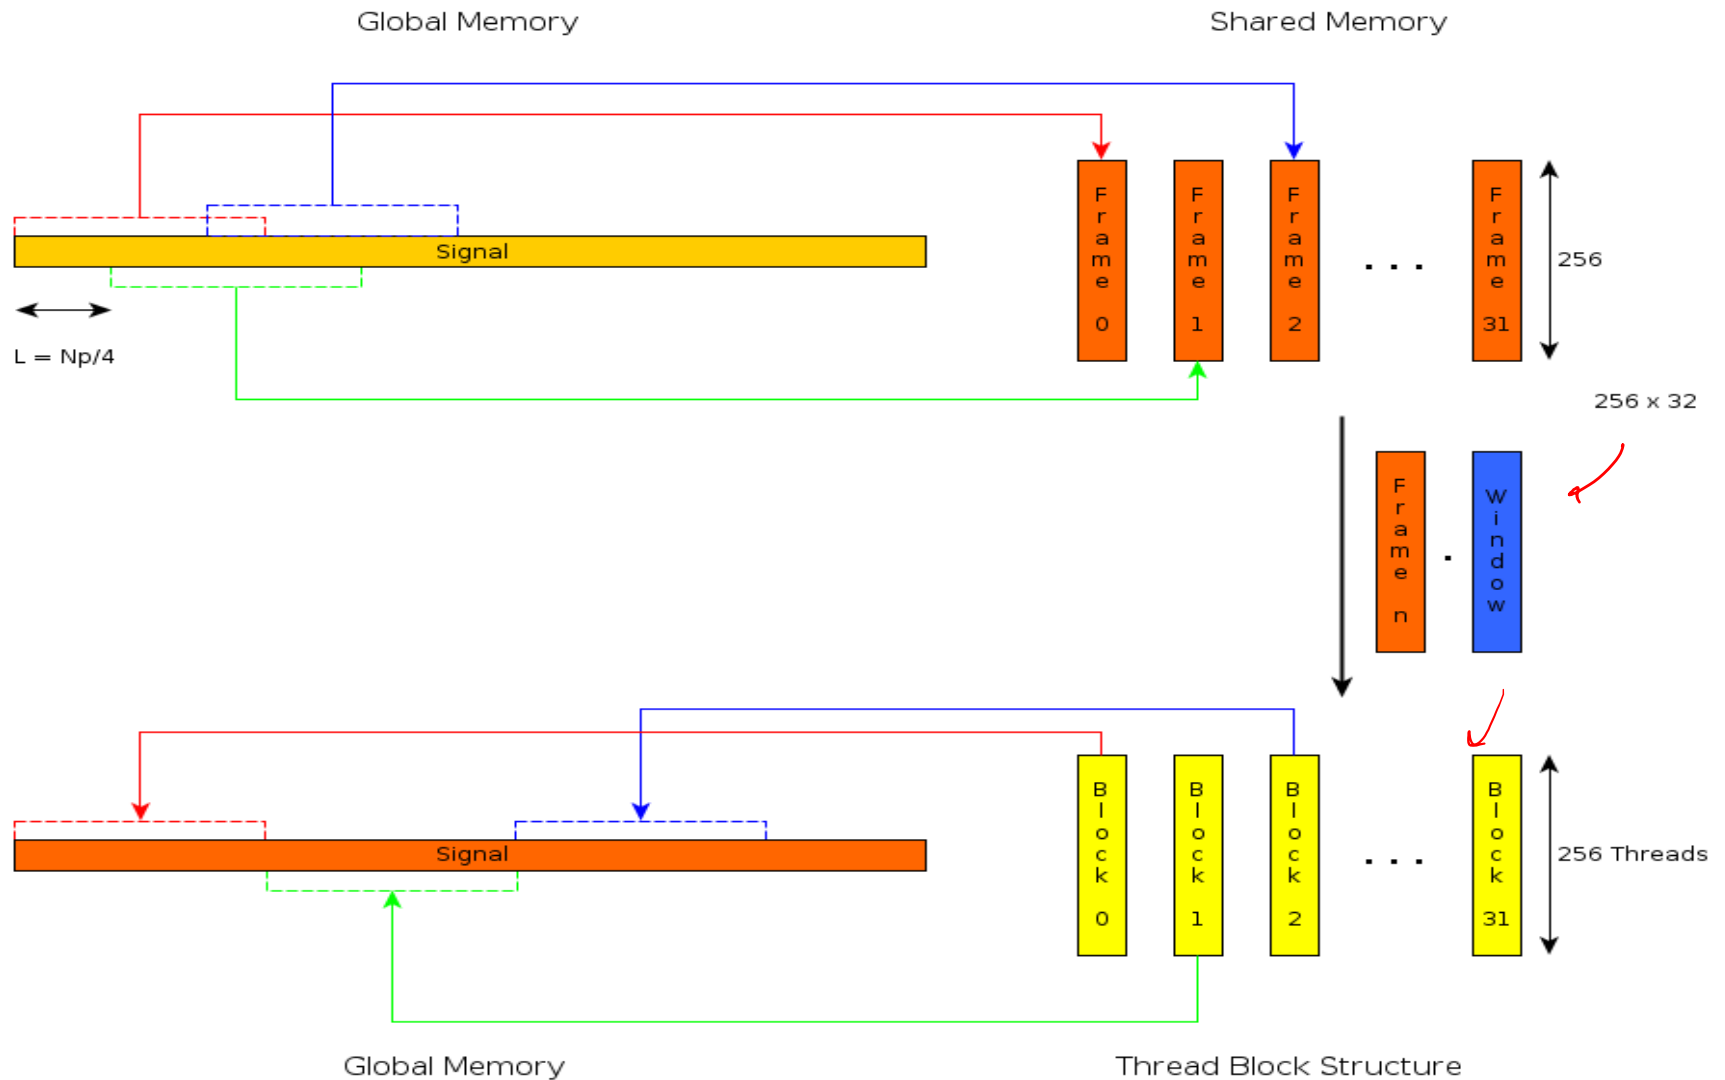
$$\omega(n) = \alpha - \beta \cos\left(\frac{2\pi n}{N_p - 1}\right)$$

- **Kernel1**: Framing and Windowing
 - Change the sparse matrix multiplication to P number of vector multiplications
 - The number of scalar multiplications is $N_p * P$
 - Factoring in the number of signals: $N_p * P * \text{signal count}$


- ✓ Typical N_p , P and signal_count:
 - 256, 32 and 30
- ✓ # of scalar multiplications:
 - 245,760



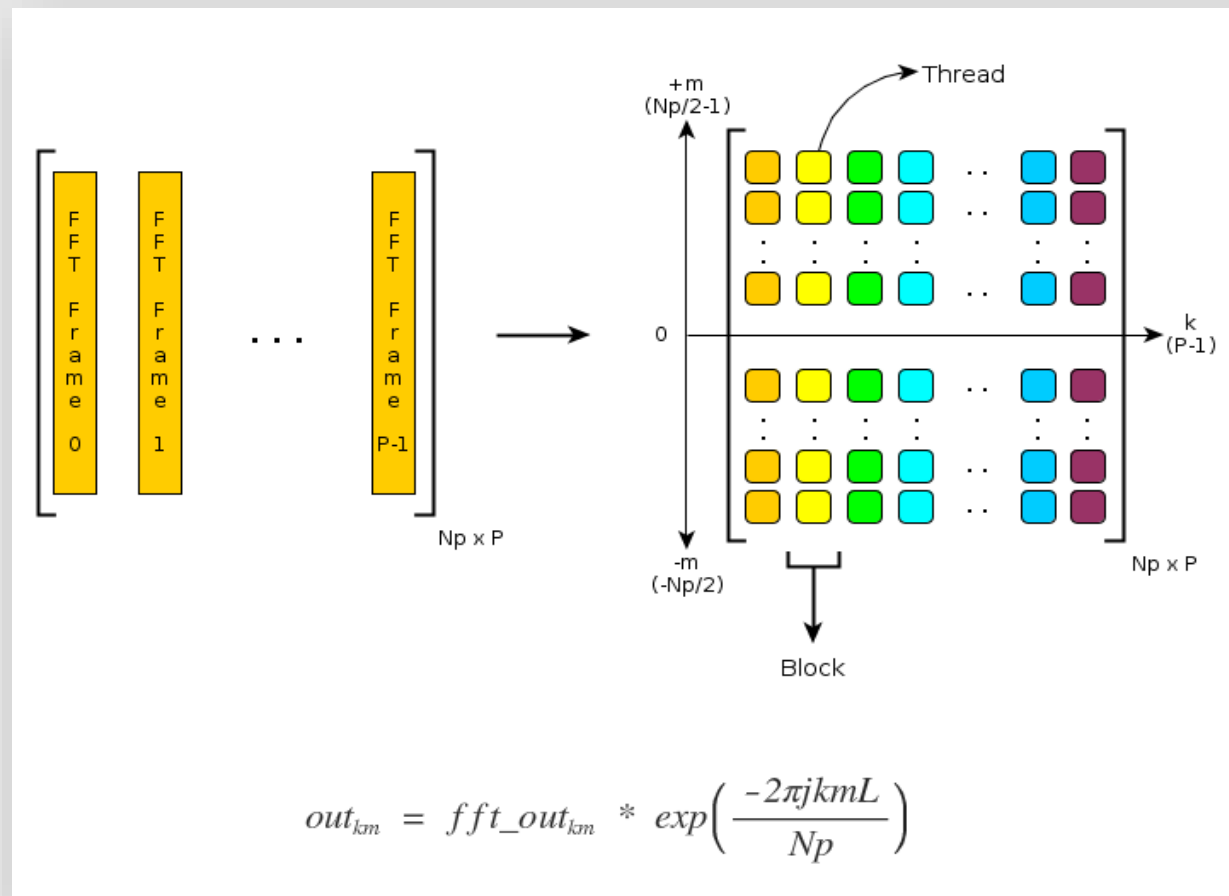
- Framing and Windowing (**Kernel 1**)



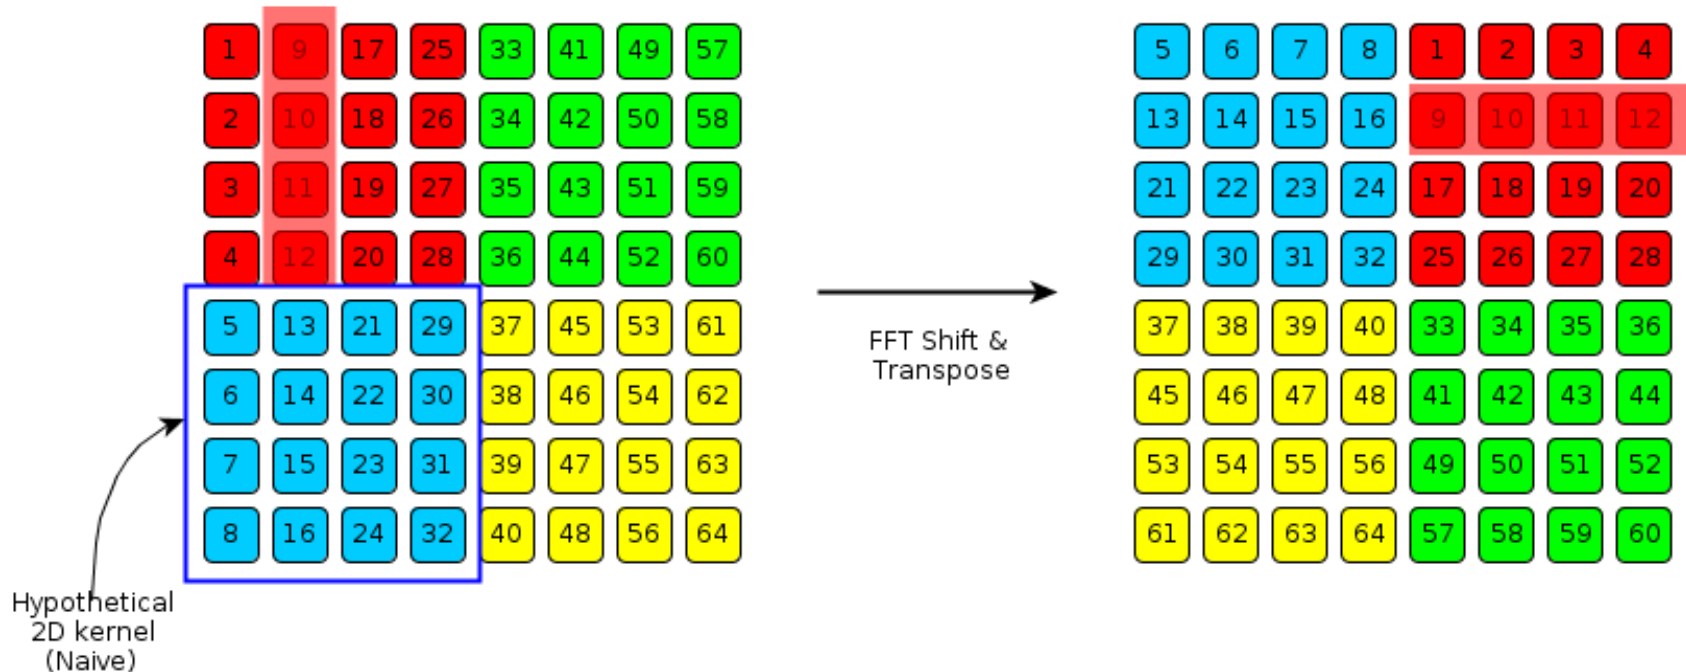
- Step-3: 256pt FFT (Kernel 2)
 - FFT applied in parallel to windowed frames

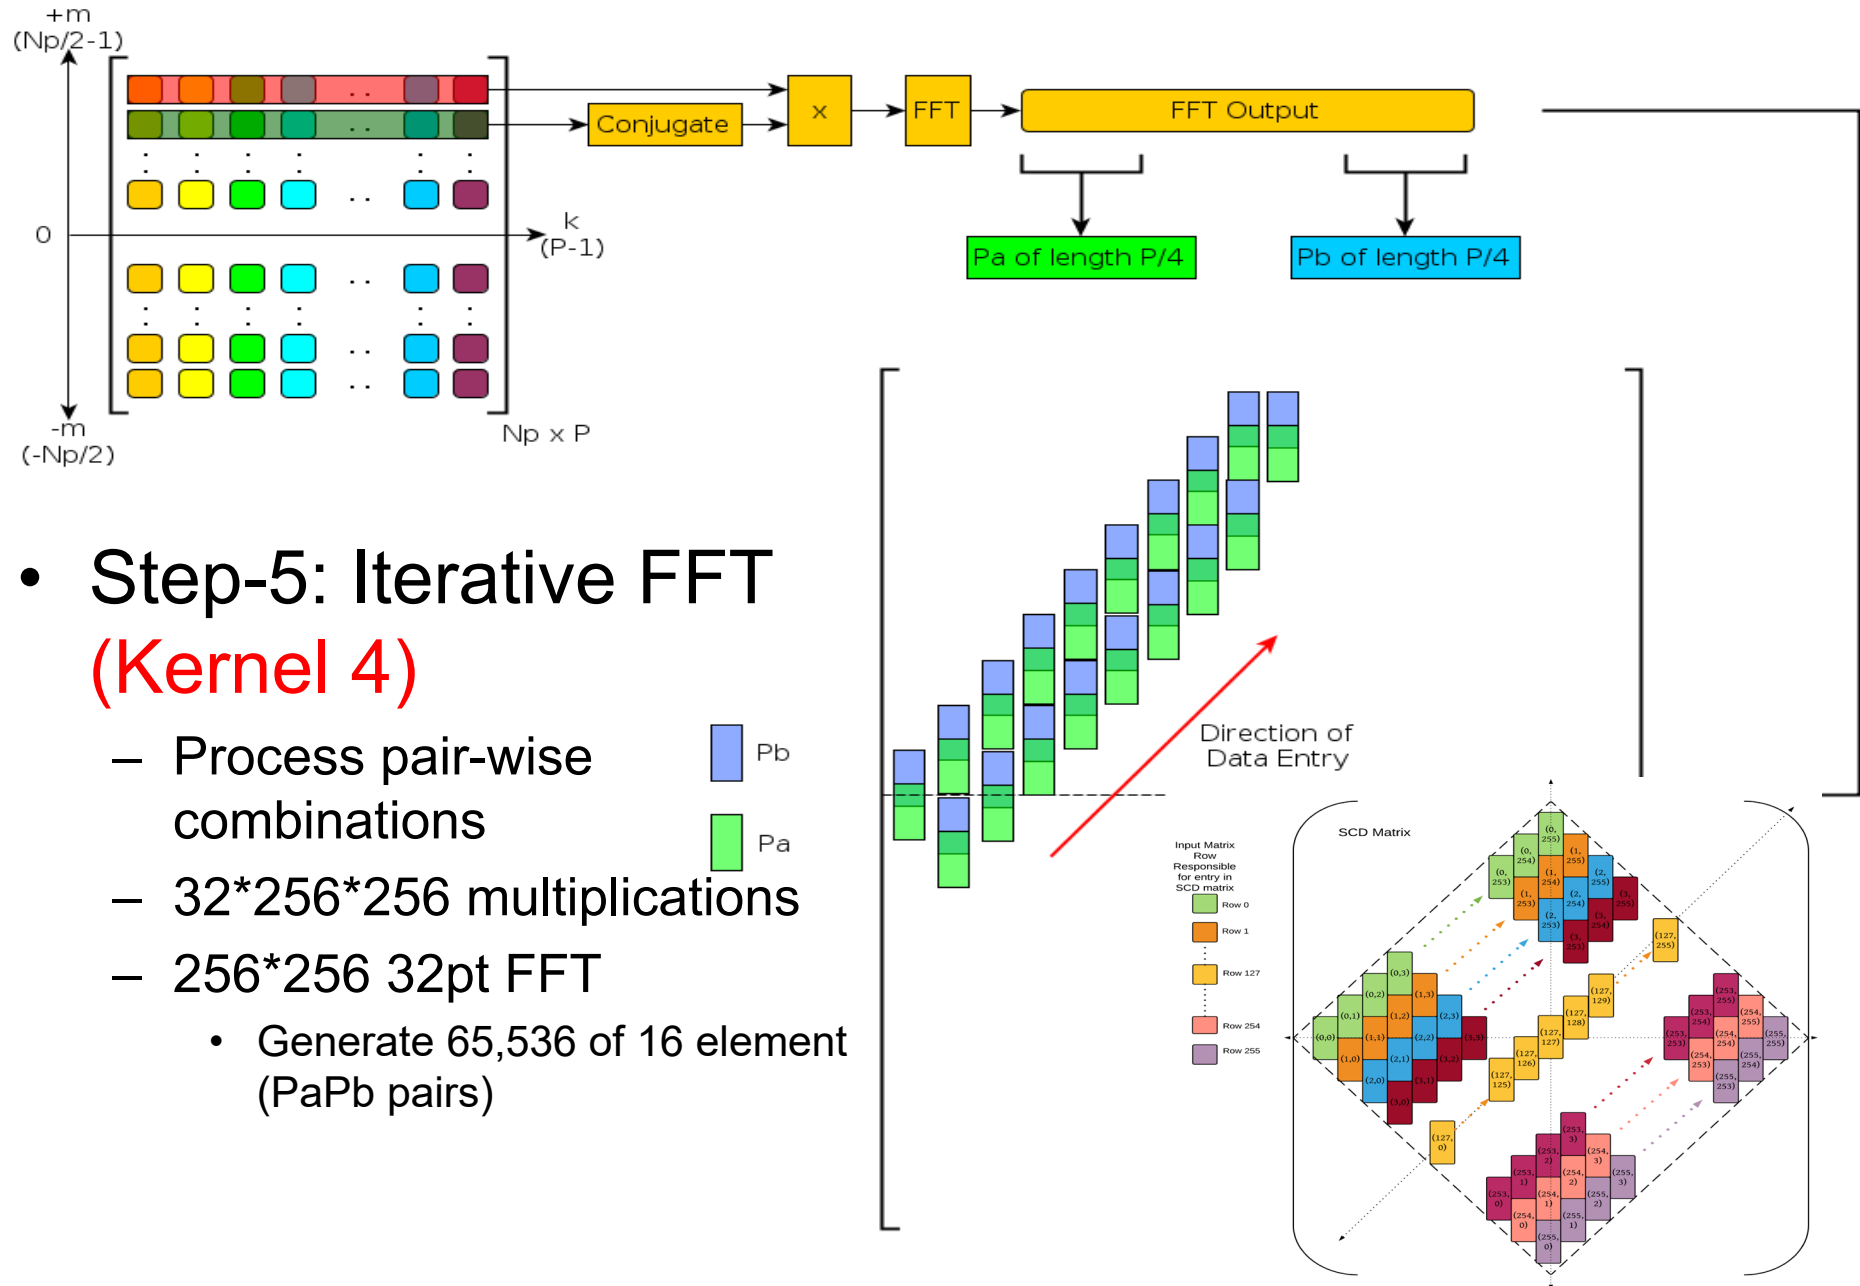


- Step-4a: Down Conversion (**Kernel 3**)
 - Apply the window to multiple frames
 - Each element is handled by a thread.

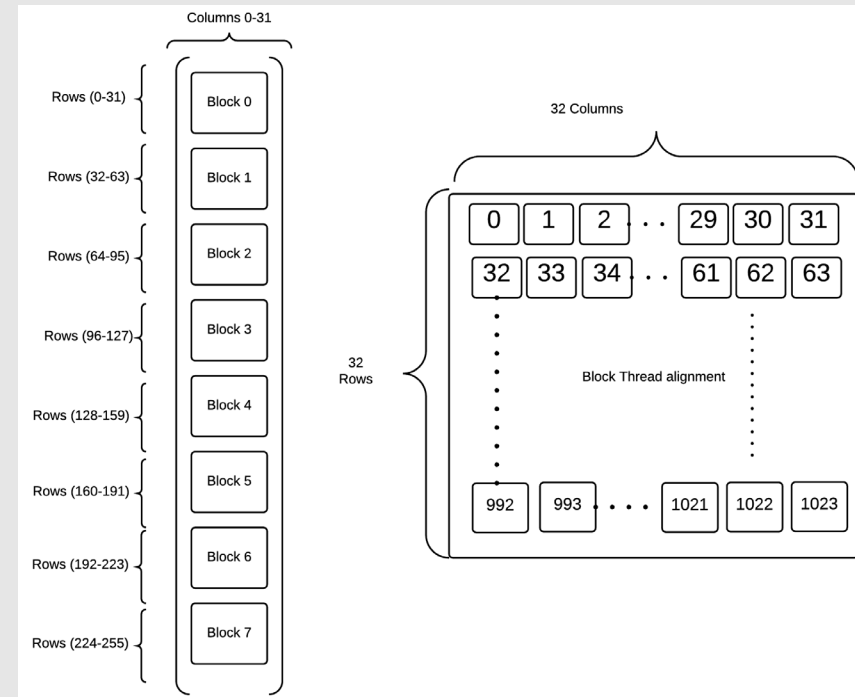



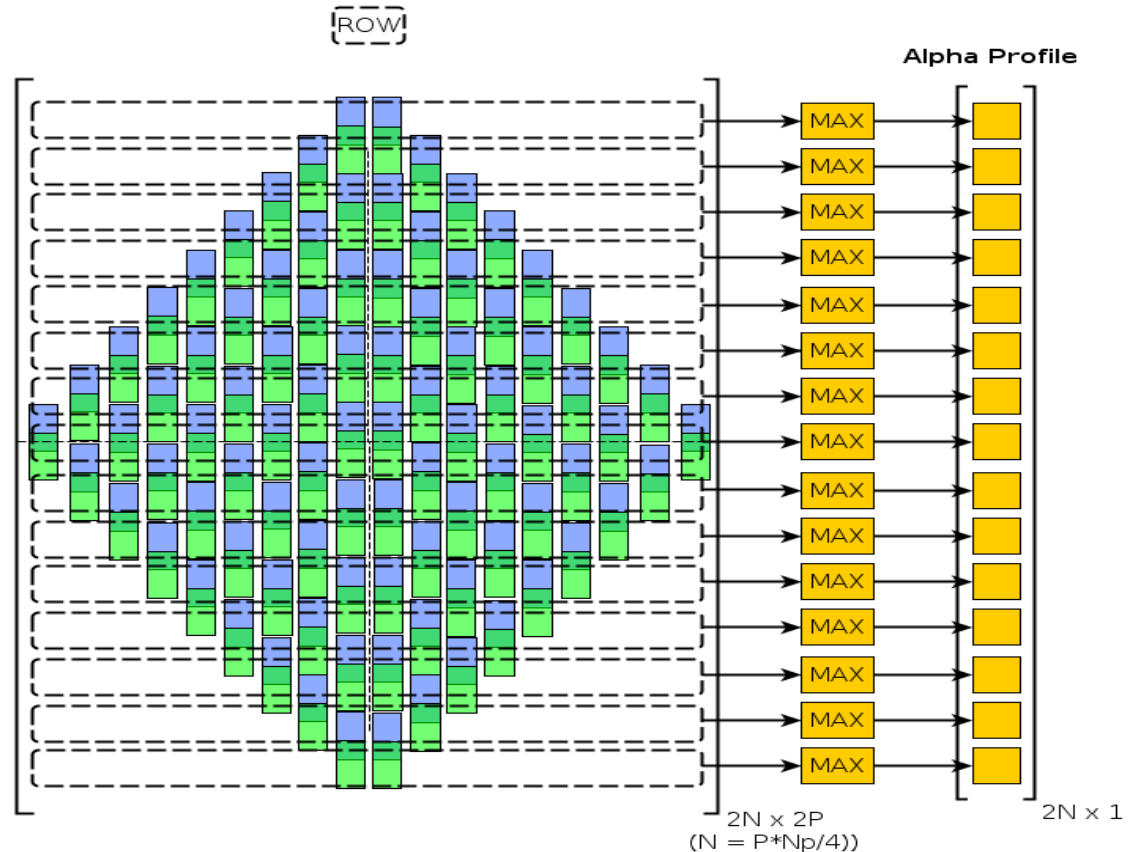
- Step-4b: FFT Shift and Transpose (**Kernel 3**)
 - Matrix is arranged in column major order
 - Setting up for coalesced memory access for FFT
 - Color coding shows how the data move before and after the transformation takes place.





- threads in blocks of 1024 threads to manage
- 32 element multiplications in groups of 32 threads (warps)





- Earlier work
 - Mapped the entire Full SCD process onto NVIDIA GPU

GPU	Matlab Intel I7 3.3GHz 32GB RAM (Reference)	CUDA K20 GPU 706MHz 2496 cores (Earlier work)*	CUDA K40 GPU 745MHz 2880 cores (This study)
Execution Time/Signal (ms) (includes data transfer)	3502.29	8.96	8.33
Speedup		390X	420X
Throughput (Signals/sec)		111	120

- Results validated against the Matlab implementation
 - minimum error of 0.0041% and a maximum error of 0.0051%.
- Execution time is based on 4096 points digital signal.

*N. Bidyanta, G. Vanhoy, M. Hirzallah, A. Akoglu, and B. Ryu, "GPU and FPGA Based Architecture Design for Real-time Signal Classification," In Proceedings of the 2015 Wireless Innovation Forum Conference on Wireless Communications Technologies and Software Defined Radio (WInnComm'15), March 24-26, 2015, San Diego, CA, pp. 70-79.

Algorithm 1 SCD Matrix and Alpha Profile Calculation

```

1: procedure SCD(data[256][32], Alpha_Profile[4096])
2:   for (i = 0; i < 256; i++) { do
3:     for (j = 0; j < 256; j++) { do
4:       Step 4.1:  $x \leftarrow data[i] * Conjugate(data[j])$ 
5:       Step 4.2:  $y \leftarrow FFT\_32(x)$ 
6:        $SCD\_Matrix \leftarrow \{y[31:24]\}\{y[7:0]\}$ 
7:       Step 4.3.1: Calculate partial Alpha_Profile
8:       Step 4.3.2: Merge partial Alpha_Profiles
9:     }# end of inner loop
10:    Step 4.3.3: Update Alpha_Profile
11:  }# end of outer loop
12:  return Alpha_Profile
  
```

Inner
loop
unrolled

Key Features	K20x	K40
Compute Capability	3.5	3.5
Memory BW (GB/sec)	250	288
Clock (MHz)	732	745
Number of MPs	14	15
Number of Cores	2688	2880
Execution Time (ms)	8.9	8.3
Throughput	111	120

Kernel Name	% Total
1 - Framing	0.13
2 - 256 point FFT	0.26
3 - FFTShift	0.11
4.1 - Conjugate	18.79
4.2 - 32 point FFT	32.93
4.3.1 - Partial	24.69
4.3.2 - Merge	9.67
4.3.3 - Update	13.44

- 4.56 micro seconds on average for a kernel launch (256*5 kernel launches)
- This corresponds to 5.84 ms of time spent on launching kernels out of the total 8.3 ms execution time,
 - which comprises 70% of the total computation time.

What if we merge
kernels?

Algorithm 1 SCD Matrix and Alpha Profile Calculation

```
1: procedure SCD(data[256][32], Alpha_Profile[4096])
2:   for (i = 0; i < 256; i++) { do
3:     for (j = 0; j < 256; j++) { do
4:       Step 4.1:  $x \leftarrow \text{data}[i] * \text{Conjugate}(\text{data}[j])$ 
5:       Step 4.2:  $y \leftarrow \text{FFT}_{32}(x)$ 
6:       SCD_Matrix  $\leftarrow \{y[31 : 24]\} \{y[7 : 0]\}$ 
7:       Step 4.3.1: Calculate partial Alpha_Profile
8:       Step 4.3.2: Merge partial Alpha_Profiles
9:     }# end of inner loop
10:    Step 4.3.3: Update Alpha_Profile
11:  }# end of outer loop
12:  return Alpha_Profile
```

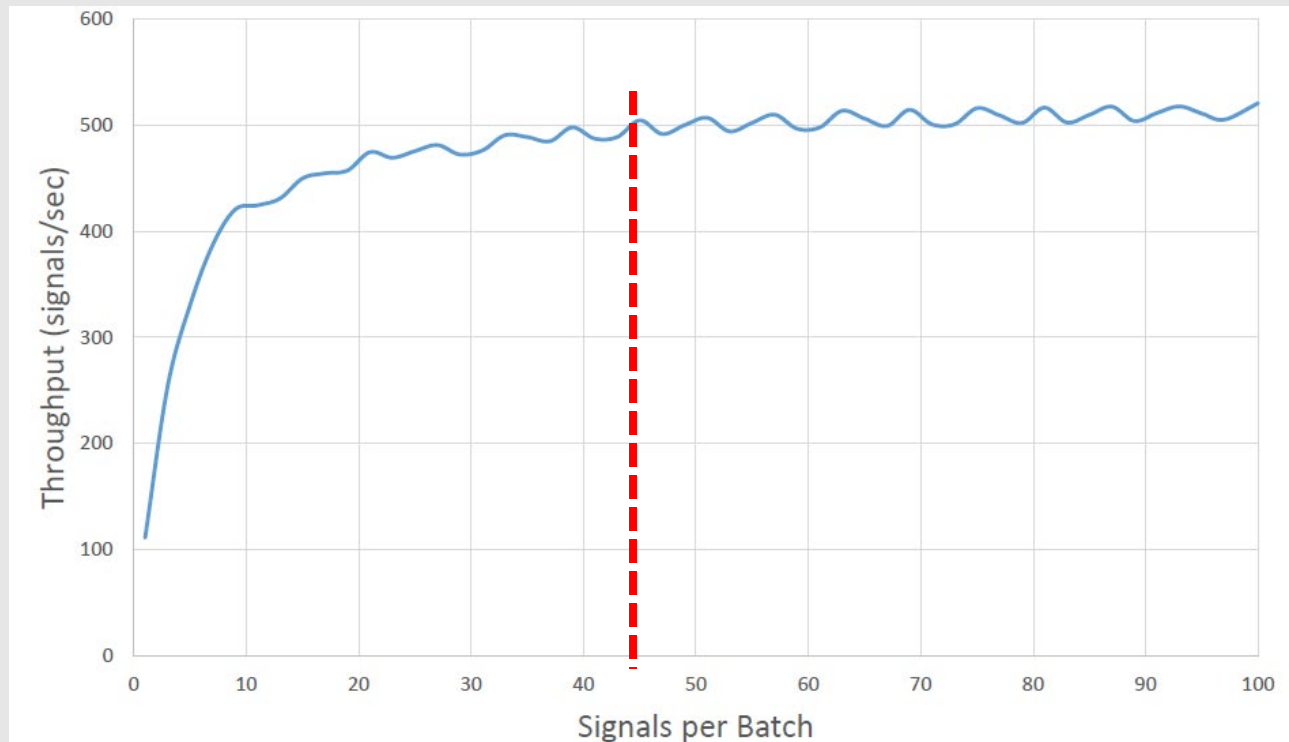
- **Batching**

- Group multiple signals together in so-called “signal-batches,”
- Increase number of calculations performed per kernel launch
- Decrease total number of kernel launches
 - *Original: 5 kernels per signal*
 - *Revised: Batch of “n” signals still uses 5 kernel launches*

Question to answer: *What is the optimal batch size?*

- **Batching – Finding Optimal Batch Size**

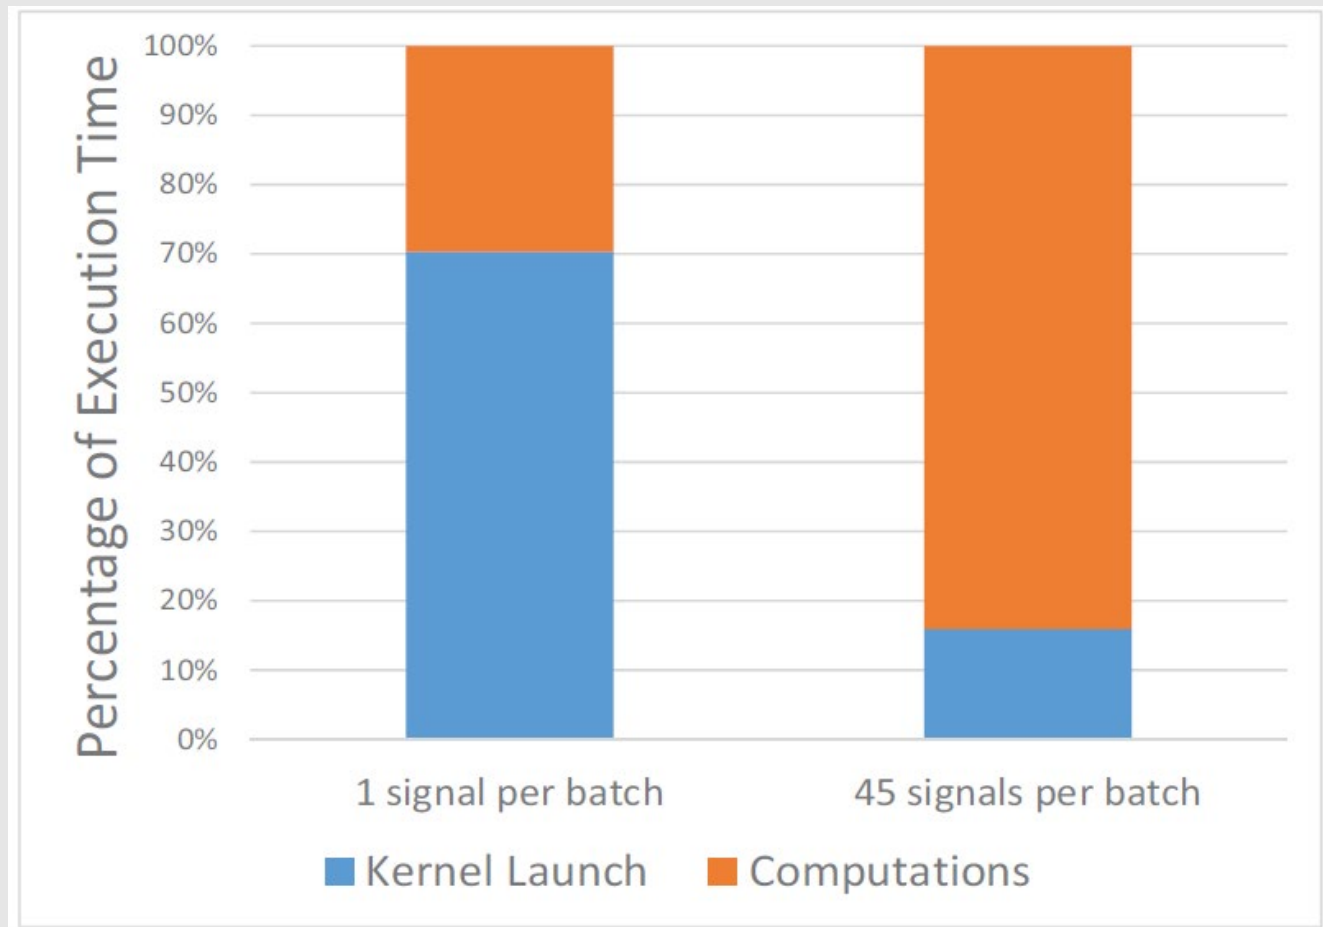
- Size 45 improves throughput from 120 to 503
 - For batch size 100: 1039 signals/second peak throughput



*Saturation trend beyond batch size 40 since all threads are occupied.
Execution time increases linearly as we add more signals to a batch*

- Complicates the utilization of the low-latency shared memory available on the GPU primarily due to memory bank conflicts.
- Therefore during each kernel launch, we ensure that individual signals of a given batch are assigned to distinct thread blocks to avoid bank conflicts.


- 70% vs 16%



- While batching reduces the kernel launch overhead, it doesn't help improve the hardware utilization
 - True or False

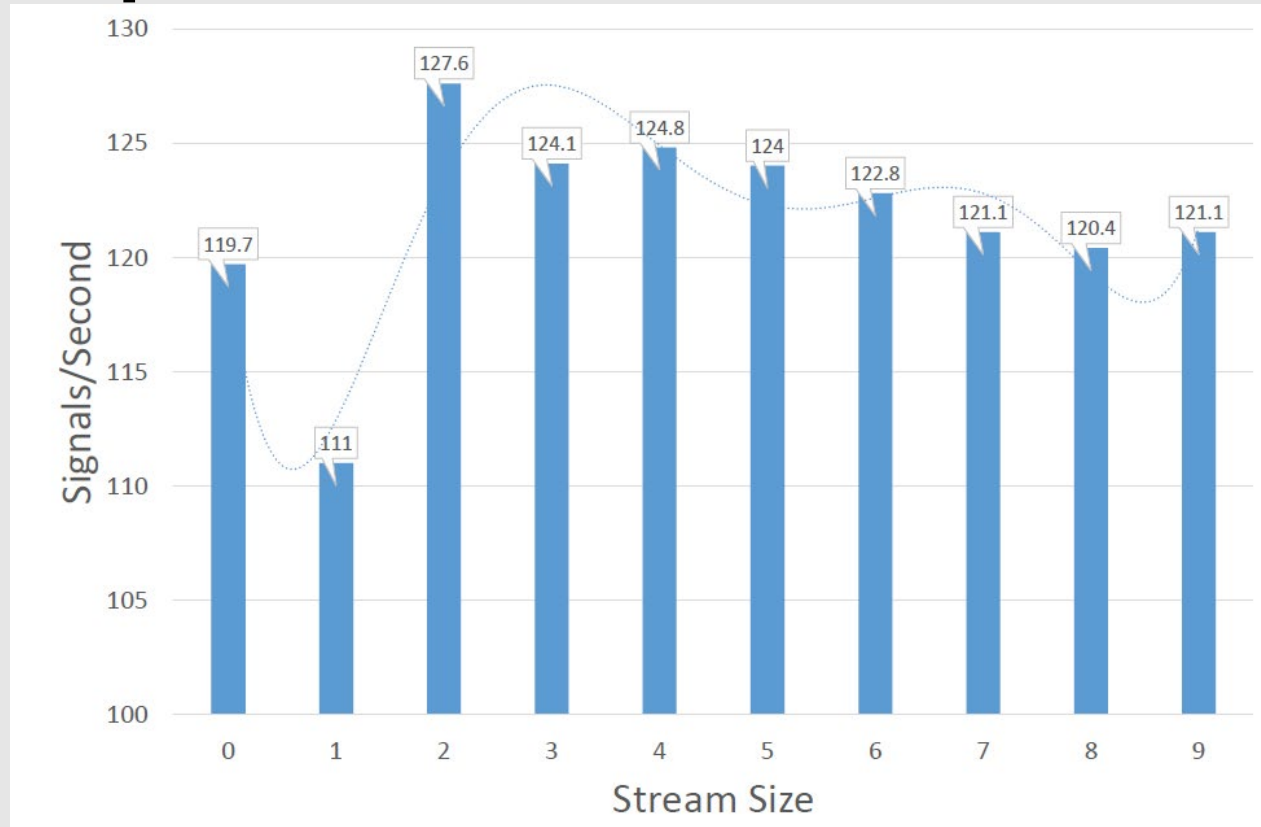
- **Batching with multiple streams**

- Different signal kernels synchronously enqueued on to *asynchronous streams*.
- symmetric multi processors (SMPs)
 - K40 has 15 SMPs
 - Allows streaming multiple signals to be executed concurrently



Question to answer: *What is the optimal stream size?*

- **Batching with multiple streams– Finding Optimal Stream Size when batch size is 1**

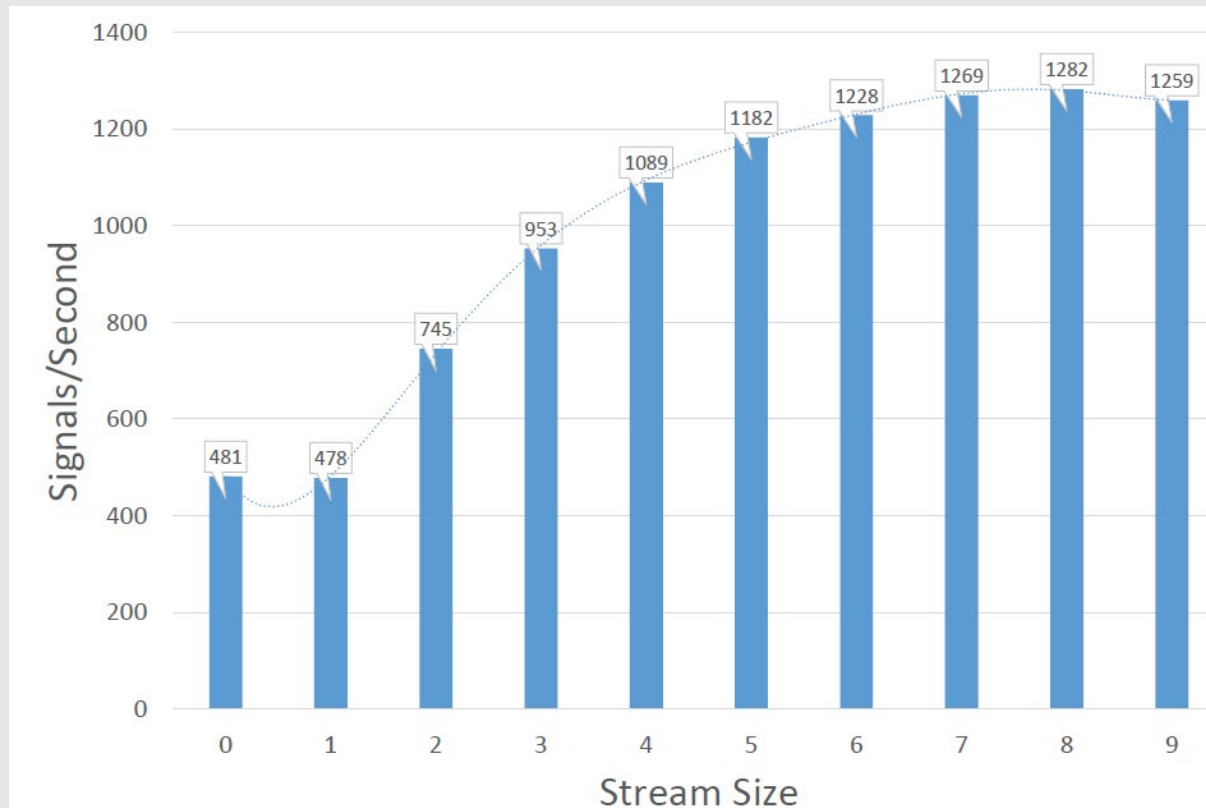


Stream size 0: No streaming, batch of 1

Stream size 1: Stream setup overhead reduces throughput

Why is streaming not helping?

- **Batching with multiple streams– Finding Optimal Stream Size when batch size is 45**

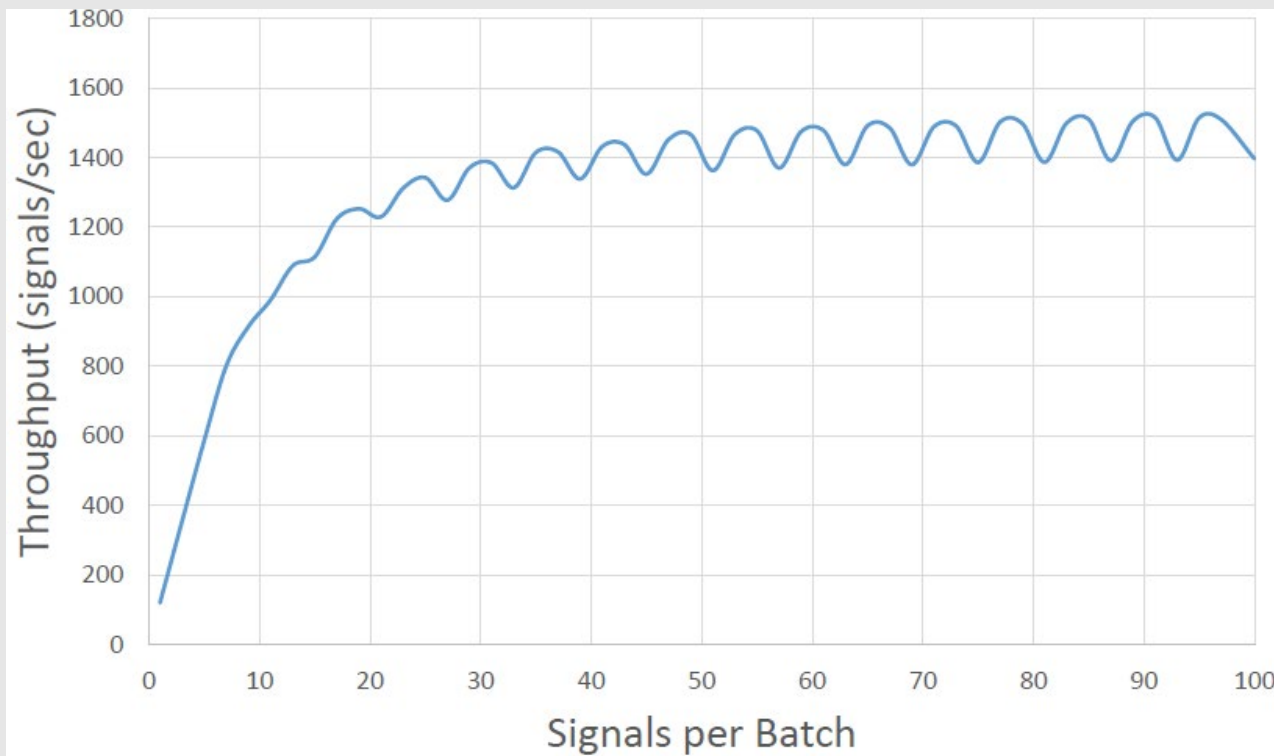


**8 streams
increases
throughput
1282**

Stream size 0: No streaming, batch of 45

Stream size 1: Stream setup overhead reduces throughput of batch size 45

- **Streaming with multiple batches– Finding Optimal Batch Size when stream size is 8**



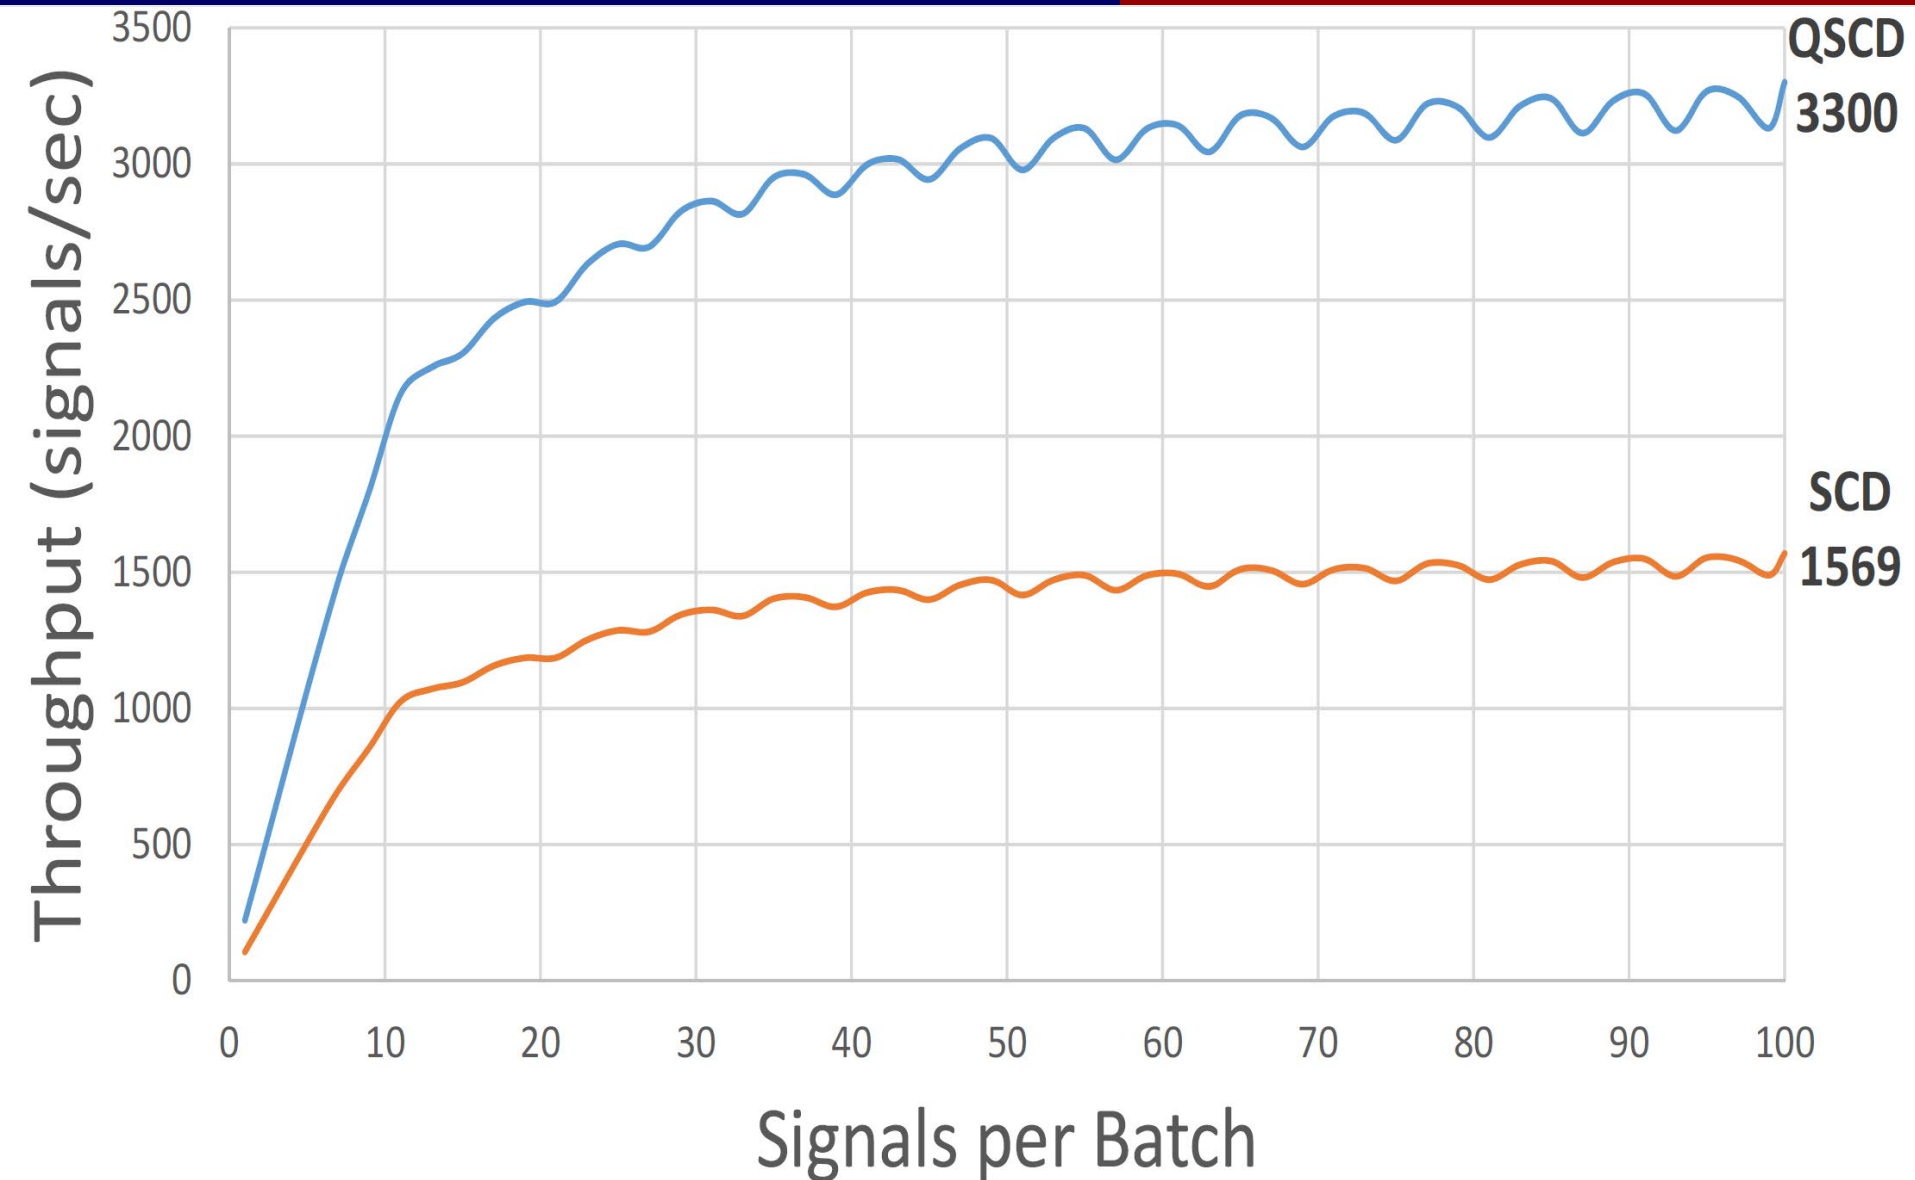
**Peak
throughput
1569 with
batch size
91**

Stream size 0: No streaming, batch of 45

Stream size 1: Stream setup overhead reduces throughput of batch size 45

- Take advantage of the symmetry in the SCD estimate

-



- Algorithmic and architecture specific optimizations
 - Quarter SCD
 - Batching
 - Streaming
- Peak throughput performance improvement over reference full SCD on GPU
 - 13.7x with batch size 45, stream size 9
 - 2719/signals/second
 - 27.5x with batch size 100, stream size 9
 - 3300 signals/second