


# ECE569

## Module 23

---



- Tiling for Matrix Multiplication

# Outline of Tiling Technique

---


- Identify a tile of global memory contents that are accessed by multiple threads
- Load the tile from global memory into on-chip memory
- Use barrier synchronization to make sure that all threads are ready to start the phase
- Have the multiple threads to access their data from the on-chip memory
- Use barrier synchronization to make sure that all threads have completed the current phase
- Move on to the next tile

# Matrix Multiplication

---


- **Data access pattern**

- Each thread - a row of M and a column of N
- Each thread block – a strip of M and a strip of N



# Tiled Matrix-Multiplication

- Divide the M and N matrices into smaller tiles.
  - Breaks up the execution of each thread into phases focusing on one tile of M and one tile of N
  - Threads collaboratively load subsets of the M and N elements into the shared memory
    - The tile is of BLOCK\_SIZE elements in each dimension
  - Threads use the elements in their dot product calculation.
- Size of the shared memory limited
  - Do not exceed the capacity
  - Choose tile size carefully

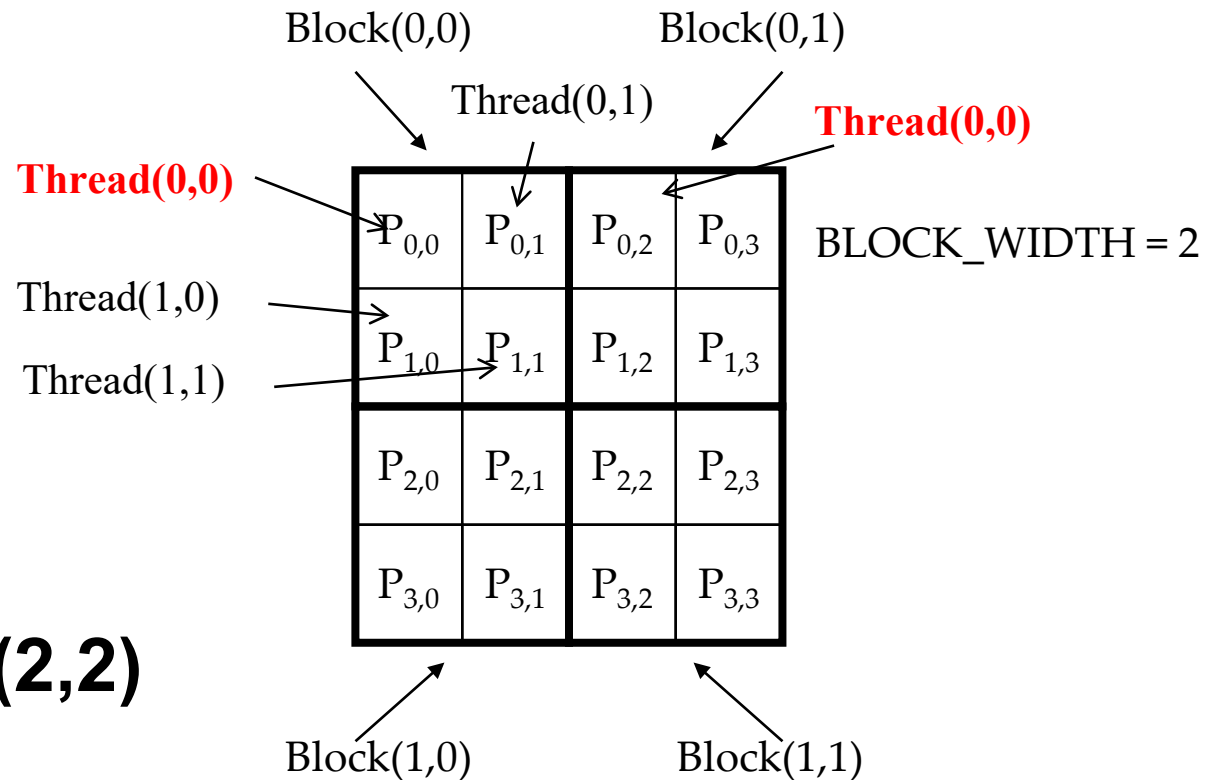


# A Toy Example: Thread to P[4x4] Data Mapping

$N_{0,0}$	$N_{0,1}$	$N_{0,2}$	$N_{0,3}$
$N_{1,0}$	$N_{1,1}$	$N_{1,2}$	$N_{1,3}$
$N_{2,0}$	$N_{2,1}$	$N_{2,2}$	$N_{2,3}$
$N_{3,0}$	$N_{3,1}$	$N_{3,2}$	$N_{3,3}$

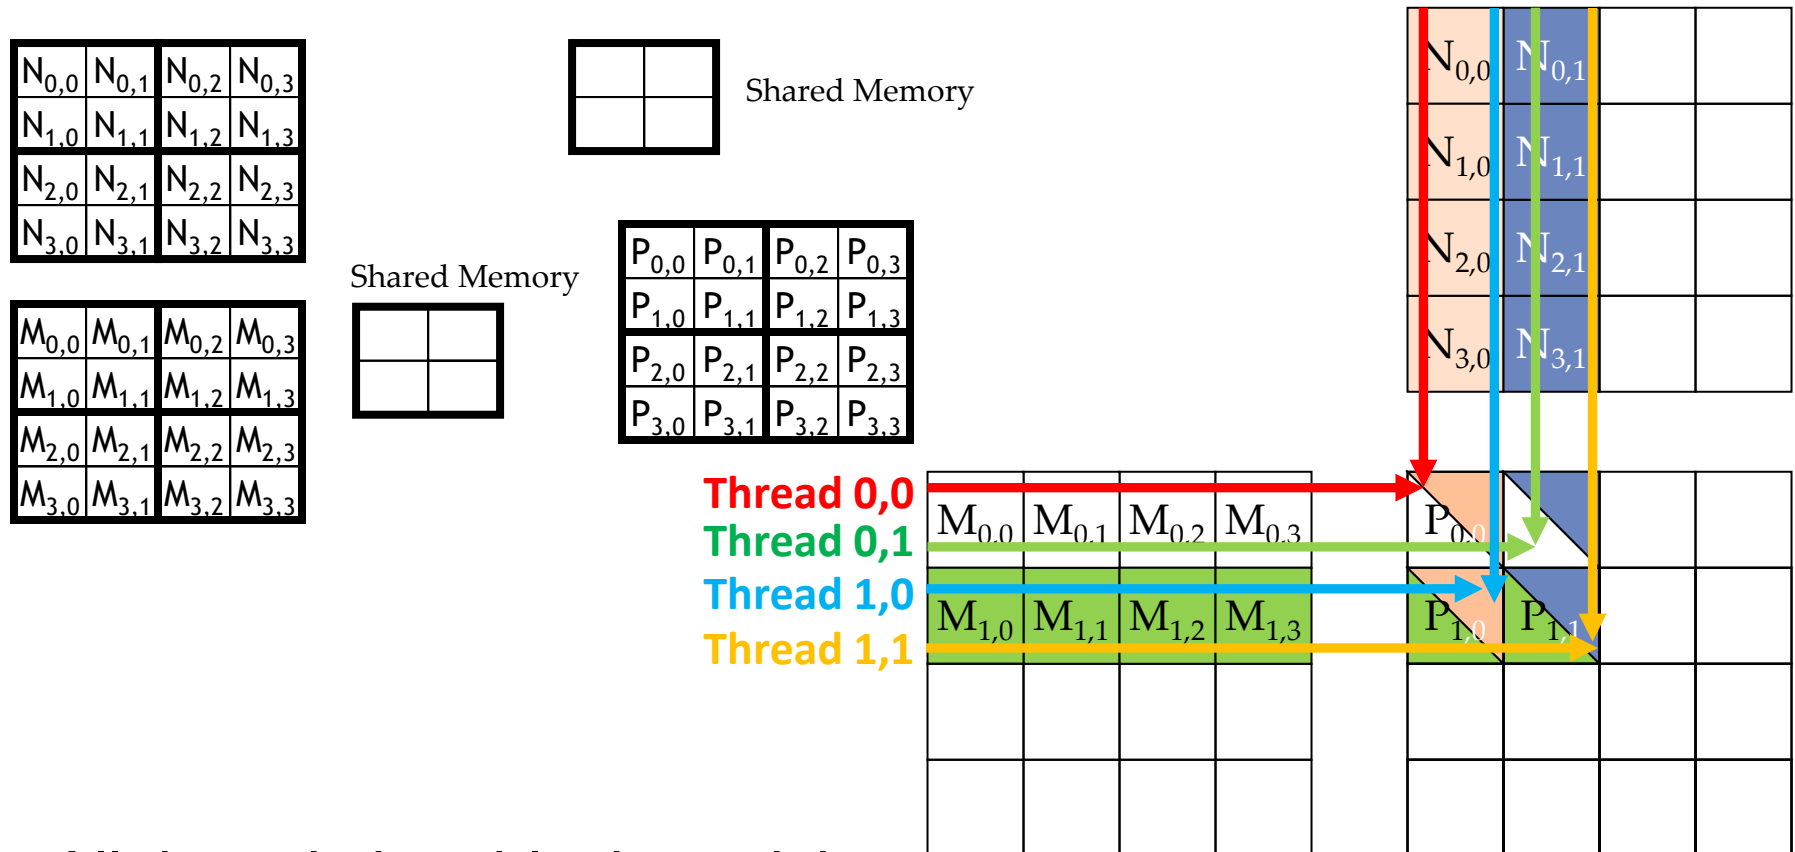
$P_{0,0}$	$P_{0,1}$	$P_{0,2}$	$P_{0,3}$
$P_{1,0}$	$P_{1,1}$	$P_{1,2}$	$P_{1,3}$
$P_{2,0}$	$P_{2,1}$	$P_{2,2}$	$P_{2,3}$
$P_{3,0}$	$P_{3,1}$	$P_{3,2}$	$P_{3,3}$

$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	$M_{0,3}$
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	$M_{1,3}$
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	$M_{2,3}$
$M_{3,0}$	$M_{3,1}$	$M_{3,2}$	$M_{3,3}$



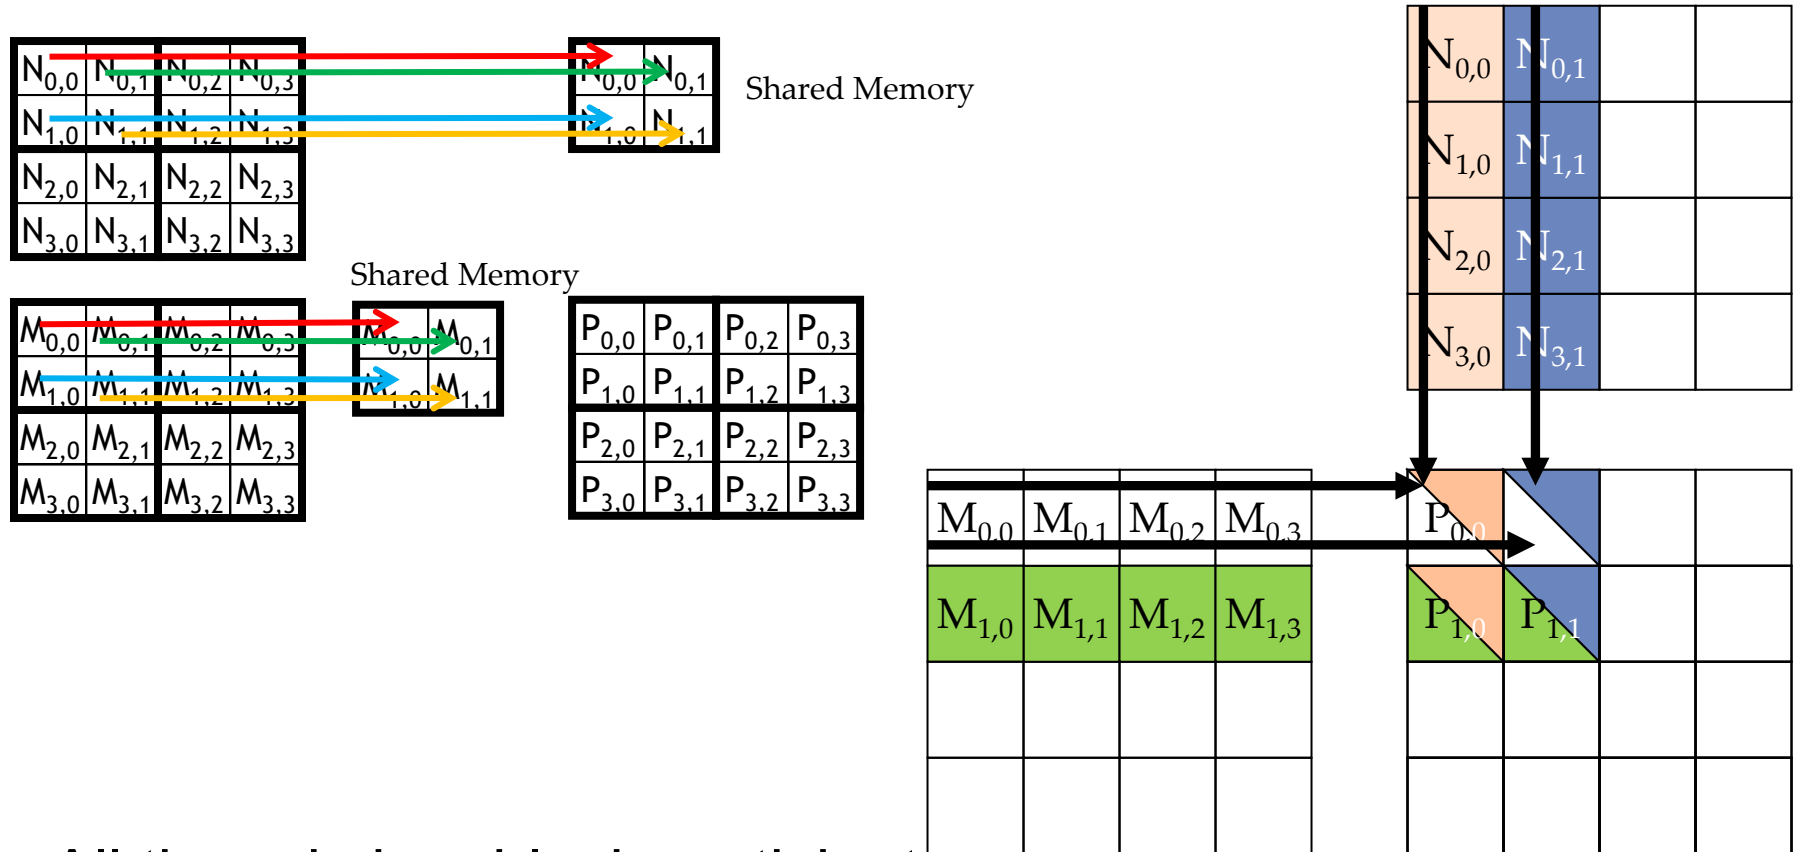
- **Grid(2,2), Block(2,2)**

# Phase 0 Load for Block (0,0)



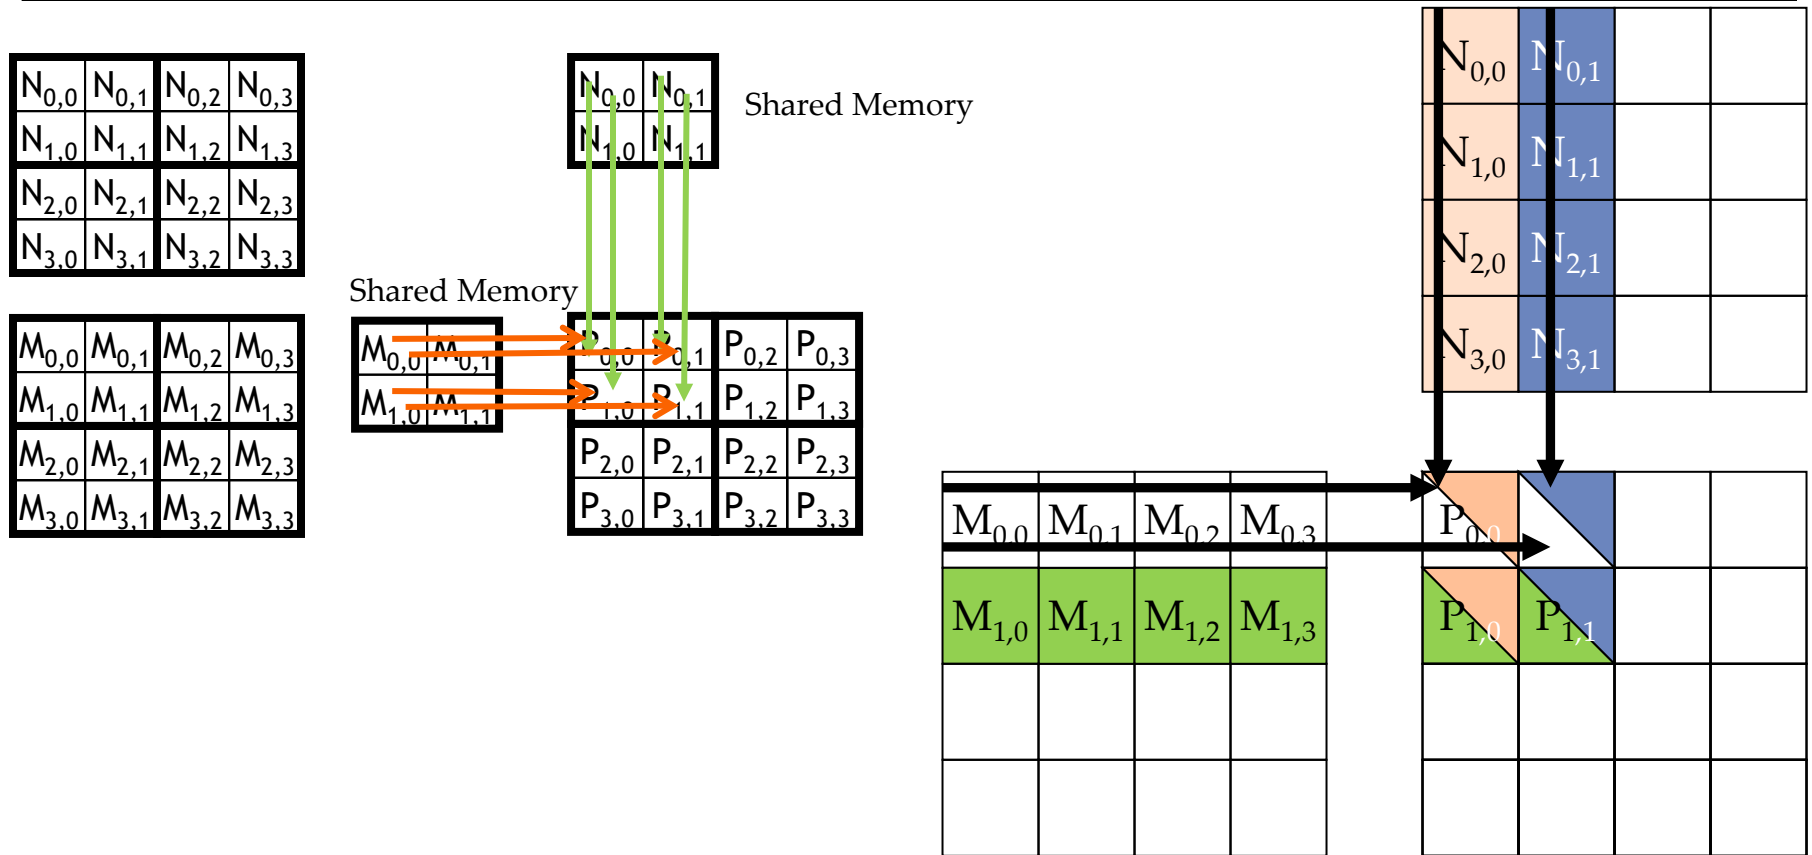
- All threads in a block participate
- Each thread loads one  $M$  element and one  $N$  element in tiled code
  - four threads of block 0,0 collaboratively load a tile of  $M$  and  $N$  into shared memory

# Phase 0 Load for Block (0,0)



- All threads in a block participate
- Each thread loads one M element and one N element in tiled code
  - four threads of block0,0 collaboratively load a tile of M and N into shared memory


# Phase 0 Use for Block (0,0) (iteration 0)



- After tiles of  $M$  and  $N$  are in the shared memory, these elements are used in the calculation of the dot product
  - Two iterations
- Note that each value in the shared memory is used twice.
  - Reduces global memory access by half

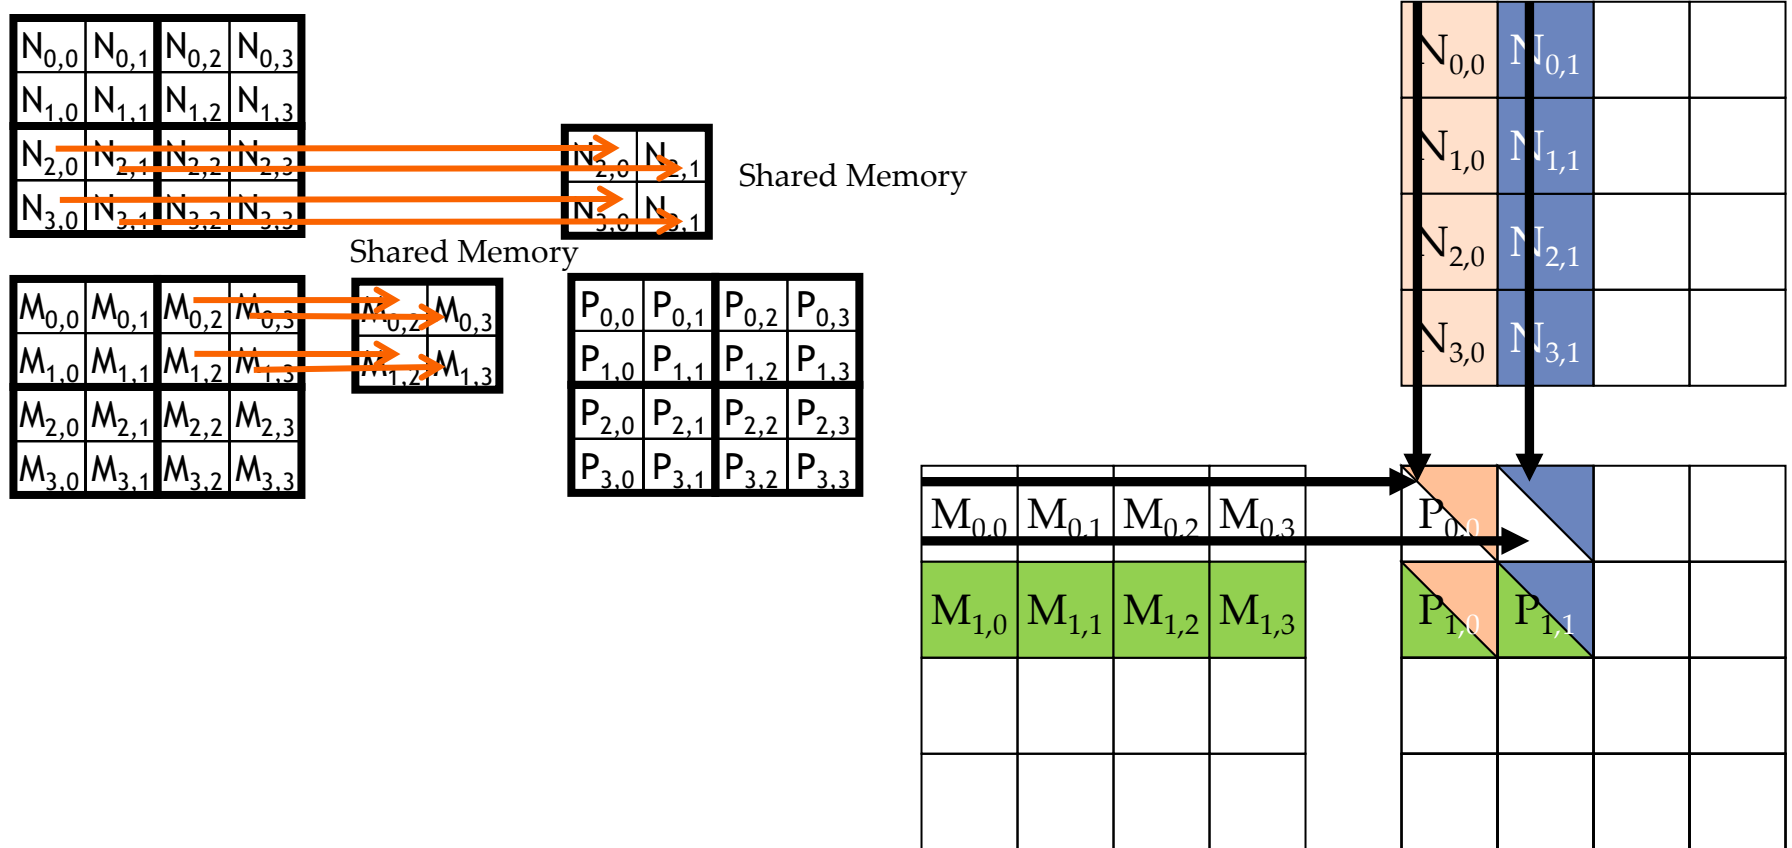


# Phase 0 Use for Block (0,0) (iteration 1)



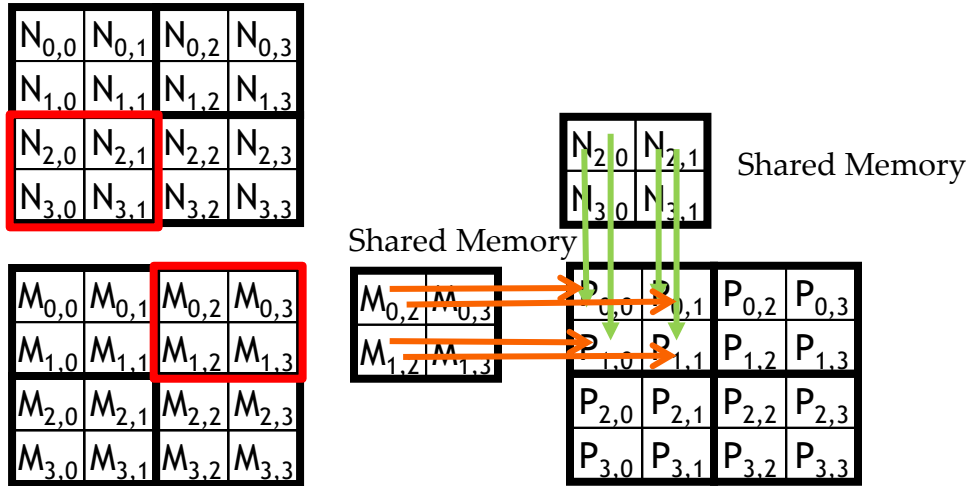
Partial accumulated dot product is a private value generated for each thread

# Phase 1 Load for Block (0,0)



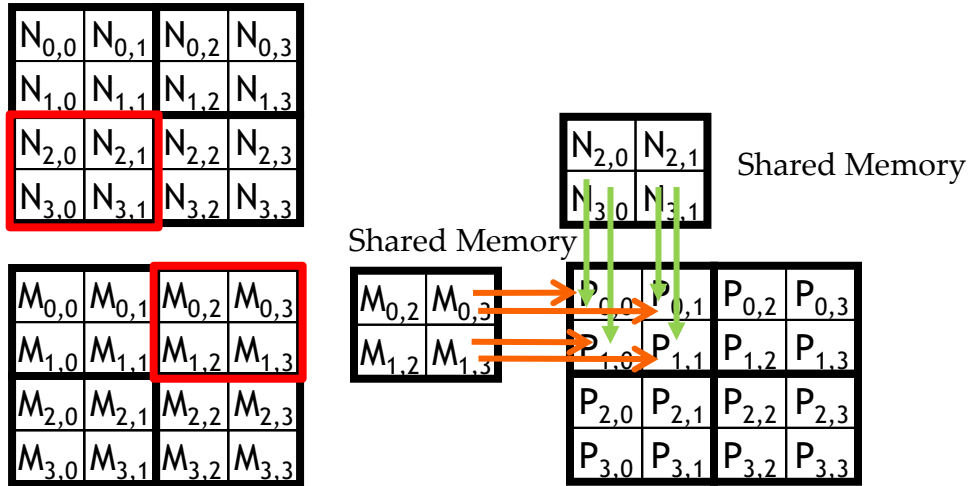
# Phase 1 Use for Block (0,0) (iteration 0)

---



# Phase 1 Use for Block (0,0) (iteration 1)

---



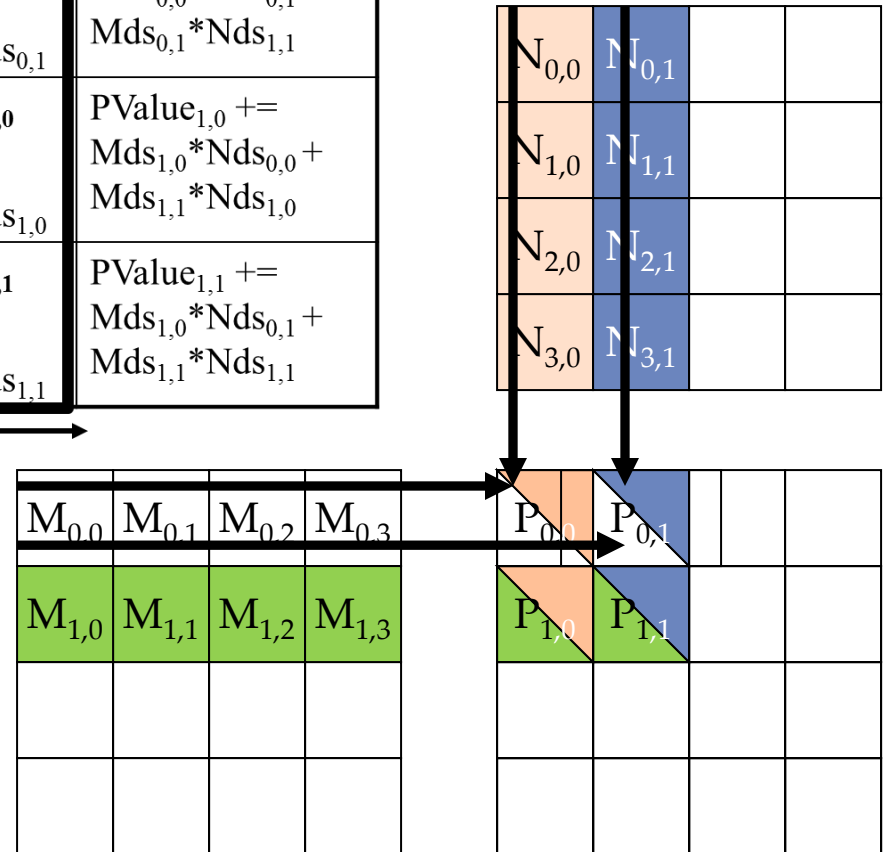
# Execution Phases of Toy Example

	Phase 0			Phase 1		
thread <sub>0,0</sub>	$\mathbf{M}_{0,0}$ ↓ Mds <sub>0,0</sub>	$\mathbf{N}_{0,0}$ ↓ Nds <sub>0,0</sub>	PValue <sub>0,0</sub> += Mds <sub>0,0</sub> *Nds <sub>0,0</sub> + Mds <sub>0,1</sub> *Nds <sub>1,0</sub>	$\mathbf{M}_{0,2}$ ↓ Mds <sub>0,0</sub>	$\mathbf{N}_{2,0}$ ↓ Nds <sub>0,0</sub>	PValue <sub>0,0</sub> += Mds <sub>0,0</sub> *Nds <sub>0,0</sub> + Mds <sub>0,1</sub> *Nds <sub>1,0</sub>
thread <sub>0,1</sub>	$\mathbf{M}_{0,1}$ ↓ Mds <sub>0,1</sub>	$\mathbf{N}_{0,1}$ ↓ Nds <sub>1,0</sub>	PValue <sub>0,1</sub> += Mds <sub>0,0</sub> *Nds <sub>0,1</sub> + Mds <sub>0,1</sub> *Nds <sub>1,1</sub>	$\mathbf{M}_{0,3}$ ↓ Mds <sub>0,1</sub>	$\mathbf{N}_{2,1}$ ↓ Nds <sub>0,1</sub>	PValue <sub>0,1</sub> += Mds <sub>0,0</sub> *Nds <sub>0,1</sub> + Mds <sub>0,1</sub> *Nds <sub>1,1</sub>
thread <sub>1,0</sub>	$\mathbf{M}_{1,0}$ ↓ Mds <sub>1,0</sub>	$\mathbf{N}_{1,0}$ ↓ Nds <sub>1,0</sub>	PValue <sub>1,0</sub> += Mds <sub>1,0</sub> *Nds <sub>0,0</sub> + Mds <sub>1,1</sub> *Nds <sub>1,0</sub>	$\mathbf{M}_{1,2}$ ↓ Mds <sub>1,0</sub>	$\mathbf{N}_{3,0}$ ↓ Nds <sub>1,0</sub>	PValue <sub>1,0</sub> += Mds <sub>1,0</sub> *Nds <sub>0,0</sub> + Mds <sub>1,1</sub> *Nds <sub>1,0</sub>
thread <sub>1,1</sub>	$\mathbf{M}_{1,1}$ ↓ Mds <sub>1,1</sub>	$\mathbf{N}_{1,1}$ ↓ Nds <sub>1,1</sub>	PValue <sub>1,1</sub> += Mds <sub>1,0</sub> *Nds <sub>0,1</sub> + Mds <sub>1,1</sub> *Nds <sub>1,1</sub>	$\mathbf{M}_{1,3}$ ↓ Mds <sub>1,1</sub>	$\mathbf{N}_{3,1}$ ↓ Nds <sub>1,1</sub>	PValue <sub>1,1</sub> += Mds <sub>1,0</sub> *Nds <sub>0,1</sub> + Mds <sub>1,1</sub> *Nds <sub>1,1</sub>

time →

Mds/Nds: shared memory array  
for the M/N elements.

**Shared memory allows each  
value to be accessed by  
multiple threads**



# Memory Bandwidth Utilization

---

What is the throughput achieved when tile size is set to 16x16 for a square matrix multiplication assuming that memory GPU device delivers 720GB/sec memory bandwidth and peak throughput of 9300GFlops?

*Hint:* You need to consider number of global memory accesses required in each phase and number of floating point operations are executed per tile.

- ☐ 360 GFLOPS
- ☐ 720GFLOPS
- ☐ 1440 GFLOPS
- ☐ 2880 GFLOPS

# Next

---

- **CUDA implementation of Tiling**